

## Highlights

**Local skill, unstable portability: sign reversal in pre-fire thermal predictors across Mediterranean wildfire regions, and the diagnostics that do not order it**

Yunus Emre Cogurcu, Emrehan Metin

- Pre-fire thermal predictors add +0.06 to +0.15 AUC within all five fire regions
- The same block adds +0.004 AUC across 20 transfer directions, interval spans zero
- Highest niche-overlap pair fails both ways; lowest transfers both ways (points)
- Of 20 diagnostics, only conditional direction agreement tracks transfer (rho 0.84)
- Label-blind adaptation compresses transfer toward chance; only target labels help

# Local skill, unstable portability: sign reversal in pre-fire thermal predictors across Mediterranean wildfire regions, and the diagnostics that do not order it

Yunus Emre Cogurcu<sup>a,\*</sup>, Emrehan Metin<sup>a</sup>

<sup>a</sup>*Çukurova University, Adana, Türkiye*

---

## Abstract

Pre-fire thermal dryness separates a fire year from a normal year. Models built on it are rarely tested outside the region they were fitted in. Five Mediterranean wildfire regions were analysed here, on about 500 m cells, with MCD64A1 labels and spatially blocked validation. Six pre-fire thermal predictors were added to a terrain, fuel and greenness baseline.

**The local skill does not travel, and a region boundary is not what stops it.** Within-region ROC-AUC rose by +0.056 to +0.153 in every region. Across twenty ordered transfer directions the same block contributed +0.004 [−0.028, +0.036]. That is not distinguishable from zero, and its sign changes from pair to pair. The static baseline transferred no better, at 0.537 against 0.541. A within-region half-split, applied with no refit through the same code path, returned 0.574, and skill is near chance from about 10 km outward, so most of the loss is incurred before any boundary is crossed. No further loss attributable to region crossing is detectable on top of it. Separation does not order the matrix, however: geographic distance gives  $\rho = -0.32$  with an

---

\*Corresponding author.

*Email address:* `ycogurcu@cu.edu.tr` (Yunus Emre Cogurcu)

interval spanning zero, the nearest pair is among the worst, and six directions are below chance with interval support, which extrapolation alone does not produce. Two predictors reverse their association between regions: elevation and the LST anomaly. Removing them cost  $-0.081$  of within-region skill, supported in every region and mostly due to elevation. It changed transfer by  $+0.014$   $[-0.017, +0.045]$ . Both figures are post-selection estimates. A local cost is measured. No compensating transfer gain is.

**The failure is invisible to the diagnostics that can be run before deployment.** Twenty candidate diagnostics were rank-correlated against observed transfer. Only two had intervals excluding zero, and both were conditional. The stronger one measures agreement in the sign of each predictor’s association (Spearman  $\rho = 0.84$  over sixteen directions). No marginal measure was shown to order the matrix, including predictor-space dissimilarity. At the point estimates, the pair with the highest burned-niche overlap failed in both directions, while the lowest transferred in both. Signed associations need labels on both sides. The family that works is therefore the one a practitioner does not have.

**The mechanism is a reversal of sign.** Predictors do not weaken across regions. They reverse the direction of their association with burning. A distance in predictor space cannot see that. The same reversal appears inside one study area, between two fires eleven months apart on an identical grid. That arm rests on one fire and eleven positive-carrying blocks, so it corroborates rather than establishes.

Label-free alignment pushed fourteen of twenty directions towards chance rather than repairing them. Transfer skill has to be measured, not inferred

from similarity. The price is target labels: thirty-two labelled 5 km blocks recovered 85 to 89 % of the target’s matched ceiling in three of six directions.

*Keywords:* wildfire, model transferability, area of applicability, land surface temperature, domain adaptation, concept shift

---

## 1. Introduction

Wildfire is a defining disturbance of Mediterranean-basin landscapes, and a changing climate is reshaping where and how it burns [27]. The pressure to anticipate *where* fire will occur has produced a large and still-growing literature on fire susceptibility mapping. It has recently been synthesised from both a methodological and an application perspective [39, 13]. The dominant pattern is well established. Geospatial predictors are assembled over a study region and paired with a historical record of burned area. A supervised classifier is then fitted, most often a random forest [4]. The resulting susceptibility surface is published with a cross-validated AUC in the 0.85 to 0.95 range.

This paper is about a property that the pattern does not report. Predictors that describe the *state* of a surface in a particular season are more informative about that season than predictors that describe a place. They are also, as we show, not portable: the relationship they carry is reparameterised locally, existing everywhere but pointing in different directions in different places. The dynamic block is worth +0.056 to +0.153 ROC-AUC inside every region and contributes +0.004 between them, an estimate whose interval spans zero, with a sign that is a property of the source-target pair rather than of the block.

The obvious reading of that pair of numbers is a trade-off, and our own earlier framing took it. The data do not support it. The static baseline transfers at a mean of 0.537 against the thermal model’s 0.541, so the predictor class that ought to travel does not travel either, and removing the reversing  
25 predictors returns +0.014 of transfer on an interval that also spans zero. Two nulls on the portability axis are not an exchange. What is measured is a local gain that does not export, and a local cost when the reversing predictors are removed, with no compensating transfer gain on either arm.

### *1.1. Portability goes unmeasured*

30 A susceptibility model’s reported skill is almost always an estimate of *within-region* performance. Held-out folds come from the same study area and season, often from the same fire event, and where folds are drawn at random over cells, spatial autocorrelation inflates the estimate further. The problem is documented across ecological modelling [31, 29] and is addressed  
35 by spatially blocked cross-validation [37, 23]. Blocking makes the within-region estimate honest. It says nothing about whether the fitted relationship holds elsewhere.

Because only that side of the ledger is reported, portability is never entered at all. A predictor block is adopted on the strength of the increment  
40 it delivers inside its training footprint, and whether that increment survives a change of region is not asked, even though any regional product built from locally trained models implicitly promises generalisation beyond that footprint. For fire specifically the transfer question is asked far less often than in adjacent fields. Meteorological fire-danger indices are known not  
45 to port cleanly between fire environments [30], and the two recent studies

that test model transfer systematically, across Mediterranean countries [7] and US counties [15], both report that it largely succeeds between similar regions. Both transfer models whose dominant predictors are *spatially stationary*: terrain, human modification, night-time lights, population density, long-term climatologies. Each describes a place rather than a season, which is the class our thesis expects to travel well. Species distribution modelling has examined transferability far more systematically [43], and there geographic and environmental similarity do *not* reliably predict transfer success [38, 32].

### 1.2. Pre-fire thermal dryness is the natural test case

The dynamic class used here is pre-fire thermal dryness, and it is the class most plausibly *expected* to transfer. Standard susceptibility predictors are terrain, fuel type, greenness, long-term climatology and proximity to roads and settlements, and most are static or near-static over the timescale at which fire danger varies. Such a set explains poorly why a particular summer burned and the preceding one did not. What changes between those summers is the state of the surface: fuel dryness, moisture stress, and how anomalously hot the land surface has become relative to its own baseline. Satellite thermal observation gives direct access to part of that state through fuel moisture content [44], and pairing surface temperature with a vegetation index was established as a live-fuel-moisture estimator for fire-danger rating by Chuvieco et al. [6]. The Temperature-Vegetation Dryness Index [33] formalises that feature space into an internally normalised measure. In principle it should be less exposed to absolute-temperature offsets between regions than raw land surface temperature. That theoretical portability advantage is tested here, and it is not found: the two internally normalised

channels transfer no better than the four absolute ones, at means of 0.544 and 0.548 over the twenty directions (Appendix A(f)).

That pre-fire thermal state carries genuine information about subsequent fire is established [17, 19, 18, 8]. What has not been established is whether  
75 the predictive *skill* of a classifier trained on it in one fire region survives strict, label-free application to another. The physics linking moisture stress to combustion is universal, so portability should be most expected for this class, which is what makes it diagnostic: a loss here cannot be dismissed as a peculiarity of a locally defined covariate.

### 80 1.3. *Why the loss is invisible to the diagnostics in use*

A deficit on transfer can arise two ways. Under **covariate shift** the marginal distribution of the predictors differs while the predictor-response relationship is preserved, which is in principle correctable without target labels, by per-region standardisation or covariance alignment such as CORAL  
85 [35], an approach with an established remote-sensing literature [36]. Under **concept shift** the conditional relationship itself changes: a predictor positively associated with burning in one region is negatively associated with it in another. No label-free realignment of input distributions can repair a sign reversal, because the information needed to detect it is exactly the  
90 information being withheld. The taxonomy is canonical [25].

This distinction is what makes the failure hard to anticipate. The diagnostics the field relies on to decide whether a model may be applied to new ground, above all area-of-applicability measures built on dissimilarity in predictor space [21, 22, 16], are marginal by construction. A target region  
95 can sit well inside the training data’s predictor envelope while the relation-

ship between those predictors and burning points the other way. This paper measures both sides on the same pairs.

#### 1.4. Contributions

Three findings carry this paper. Each is stated at the strength its interval  
100 supports, and the nearest prior work is named inside the claim.

**Contribution 1. Local skill and portability, measured separately and per direction.** The thermal block is worth +0.056 to +0.153 ROC-AUC inside every one of five regions, with every bootstrap interval above zero and the result holding in both analysis populations. Across the twenty  
105 ordered transfer directions it contributes +0.004  $[-0.028, +0.036]$ . That interval spans zero under all four between-direction resampling units we computed, and the sign varies by pair: paired deltas run from  $-0.148$  to  $+0.132$ , twelve positive and eight negative. Two controls bound the reading. The baseline arm transfers at a mean of 0.537, against the thermal model’s 0.541,  
110 so the static predictor class is not the portable one either. And a within-region half-split, fitted on one contiguous half and applied to the other with no refit through the same code path, returns 0.574, with skill near chance from about 10 km outward. Most of the loss is therefore incurred **without leaving the region**, and **no additional loss attributable to region crossing is detectable on top of it**. That is a null rather than a positive  
115 attribution to distance: the within-region curve is at the chance floor before the cross-region separations begin, so the comparison could not have come out otherwise, and the half-split confounds separation with training-set composition. What separation demonstrably does *not* do is order the matrix.  
120 Geographic distance gives  $\rho = -0.32$  with an interval spanning zero, the

nearest pair is among the worst, and six directions are below chance with interval support, which extrapolation alone does not produce. The variation around the floor is what Contributions 2 and 3 address. Feature removal measures the local cost. Dropping the two reversing predictors, elevation and the LST anomaly, costs  $-0.081$  of mean within-region AUC, supported  
125 in every region. Roughly three quarters of that is elevation, a *baseline* terrain variable. Mean transfer changes by  $+0.014$ , whose interval also spans zero. Both are post-selection estimates (Section 4.6b). A local cost is measured. No compensating transfer gain is. The within-region half is not novel,  
130 comparable results existing for these landscapes [1, 12]; the paired contrast against portability is. Dimarco et al. [7] transfer a predominantly stationary predictor set successfully across a comparable Mediterranean design, though their response variable is ignition rather than burned area (Section 5.7).

**Contribution 2. The loss is invisible to the diagnostics that  
135 could be run before deployment, and visible to one that cannot.**  
Twenty candidate diagnostics from four families are evaluated against observed transfer under one bootstrap framework. Area-of-applicability-style predictor-space dissimilarity [21, 22, 16], climatic and geographic distance, learned domain separability, niche-overlap statistics and burn-pattern regime  
140 distances were none of them shown to order the matrix; the domain classifier is at ceiling, with  $\text{AUC} \geq 0.96$  for every pair. Only two diagnostics have intervals excluding zero. Both measure conditional direction agreement, that is, whether each predictor’s signed association with burning points the same way in both regions (Spearman  $\rho = +0.84$  and  $+0.81$  over sixteen directions  
145 from eight pairs). Both are computed on a data-selected subset of predic-

tors, and their all-nine-feature counterparts span zero, so the result rests on that selection. Equalising the families onto a common twelve directions leaves the ordering unchanged, so it is not an artefact of unequal samples. Signed associations need burned labels on both sides, so this is a mechanism  
150 diagnosis rather than a pre-deployment screen, and the marginal family, the only one runnable before deployment, is the one that fails. The same point survives without any ranking, at the point estimates: the pair with the highest burned-niche overlap fails in both directions while the lowest transfers in both. That contradicts an expectation the fire literature carries [7, 15]  
155 and agrees with species distribution modelling [38, 32]. Shift decomposition in applied remote sensing is not itself new [11]; the fire application and the head-to-head of four diagnostic families on the same pairs are.

**Contribution 3. The mechanism is a reversal of sign, and it survives holding geography fixed.** Predictors do not merely weaken across  
160 regions, they reverse the direction of their association with burning, which is why a distance in predictor space cannot see the failure. The sharpest reversal is elevation, whose association points opposite ways in Manavgat and in Bejís and Muğla, with each region’s own interval excluding chance and the paired difference excluding zero (Appendix B, Table B3). The same reversal  
165 appears inside a single study area, between two fires eleven months apart on an identical grid, where season and population also differ but place does not (Section 4.8). That arm rests on one fire and eleven positive-carrying 5 km blocks, below the floor this design sets for itself, so it corroborates the mechanism rather than establishing it.

170 Two consequences follow. Both are reported as supporting results, not

leading ones. First, label-free alignment by standardisation and covariance alignment [35] does not repair transfer. It compresses transfer towards chance in fourteen of twenty directions. We believe this is the first application of covariance alignment to fire susceptibility. Second, the residual gap is  
175 conditional, so the resource that closes it is target labels. That price is measurable. Thirty-two labelled 5 km blocks recover 85 to 89 % of the target’s matched ceiling in three of six directions, 51 to 57 % in two more, and 30 % in the sixth.

Alongside these, a leakage-audited, spatially blocked evaluation and trans-  
180 fer protocol is released with code, configuration and frozen outputs, so that most of this result can be re-run rather than taken on trust. One component is an exception and is named in the data-and-code availability statement: the reproduction-check code is not in the released repository at the commit of record. The few-shot export, previously also listed as an exception, was  
185 regenerated at commit `6d7a6a71` and its validator now passes on every check. This matters given evidence that wildfire transfer conclusions are sensitive to evaluation design and task formulation [42].

A companion paper treats the observational layer beneath this one, where preprocessing decisions taken before any model is fitted are shown to carry  
190 their own uncertainty budget.

Section 2 reviews the relevant literature. Section 3 describes the study regions, the data, and the modelling, cross-validation, transfer and adaptation protocol. Section 4 reports results. Section 5 discusses them, including limitations. Section 6 concludes.

## 195 2. Related work

### *2.1. Fire susceptibility mapping, and what it reports*

Fire susceptibility mapping has been synthesised recently from both methodological and application perspectives [39, 13]. The dominant pattern is stable: geospatial predictors are assembled over a study region, paired with a historical burned-area record, and fitted with a supervised classifier, most  
200 often a random forest [4], and the resulting surface is published with a cross-validated AUC in the 0.85 to 0.95 range. That figure is a within-region estimate. Comparable within-region results already exist for the landscapes studied here [1, 12].

### 205 *2.2. Pre-fire thermal dryness*

Satellite thermal observation gives repeated access to surface state through fuel moisture content [44]. Pairing land surface temperature with a vegetation index was established as a live-fuel-moisture estimator for fire-danger rating by Chuvieco et al. [6], and the Temperature-Vegetation Dryness Index [33] formalises that feature space into an internally normalised measure,  
210 which is why it is often expected to travel between regions better than raw temperature. That pre-fire thermal state carries information about subsequent fire is established [17, 19, 18]. Gelabert et al. [8] found dead fine fuel moisture and its anomalies to be the most influential predictor of human-caused ignition likelihood across Europe, testing generalisation by pooled  
215 fitting with per-site evaluation. What has not been tested is whether a classifier built on this class survives strict, label-free application to a region it never saw.

### 2.3. *Spatial validation and transferability*

220 Random cross-validation over spatially autocorrelated cells inflates skill estimates [31, 29], and spatially blocked designs are the standard remedy [37, 23]. The remedy is not uncontested: blocked schemes have been argued to introduce pessimistic bias of their own [40, 24, 5]. This paper uses spatial blocking throughout and reports every result at three block sizes, so that a  
225 reader can see which verdicts depend on the choice.

Blocking makes a within-region estimate honest but says nothing about another region. For that, the field’s instrument is the area of applicability and related predictor-space dissimilarity measures [21, 22, 16], which ask whether a target’s predictor values fall inside the training data’s envelope.  
230 Species distribution modelling has examined transferability far more systematically than fire has [43], and two studies there report that geographic and environmental similarity do not reliably predict transfer success [38, 32], and the transferability of satellite-derived empirical models is an active question in fire remote sensing itself [20].

### 235 2.4. *Dataset shift and label-free adaptation*

The canonical taxonomy separates covariate shift, where the predictor distribution moves but the predictor-response relationship holds, from concept shift, where the relationship itself changes [25]. Covariate shift is in principle correctable without target labels, by per-region standardisation or  
240 by covariance alignment such as CORAL [35], and domain adaptation has an established remote-sensing literature [36, 28]. Concept shift is not: no realignment of input distributions can repair a reversal in the sign of an association, because detecting it requires the labels being withheld. Shift

decomposition in applied remote sensing is not itself new [11]. We are not  
245 aware of a prior application of covariance alignment to fire susceptibility, fire  
occurrence or burned-area prediction.

### 2.5. Cross-region generalisation of fire models

Few studies test fire-model transfer directly. Podschwit et al. [30] give  
the long-standing reference point that meteorologically derived danger in-  
250 dices do not port cleanly between fire environments. WildfireGenome [15]  
runs a leave-one-county-out matrix across seven US counties and reports  
strong within-county performance with highly variable off-diagonal transfer,  
similar pairs transferring well and dissimilar pairs collapsing; its label, how-  
ever, is a principal-component composite of hazard *indicators* rather than  
255 observed burned area, and it applies no adaptation. Xu et al. [42] argue that  
wildfire transfer conclusions depend strongly on evaluation design, a caution  
we address by fixing the protocol in a project log before the diagnostics were  
computed and by reporting every sensitivity axis; that log is not a formal pre-  
registration, and Section 3.11 states what was fixed and when. Kondylatos  
260 et al. [14] provide Mesogeos, a 1 km Mediterranean datacube for data-driven  
wildfire modelling.

#### **The nearest neighbour, and the contrast the present paper draws.**

Dimarco et al. [7] is the closest Mediterranean analogue: 500 m predic-  
tors harmonised across four countries, tree ensembles under spatial cross-  
265 validation, transfer tested both leave-one-country-out and as a full 4 x 4  
matrix, with every transfer exceeding AUC 0.80 and bioclimatically similar

countries scoring higher. They apply no domain adaptation.<sup>1</sup>

Much of the design is shared — Mediterranean regions, 500 m cells, MCD64A1-derived targets, tree ensembles, spatially aware validation, an explicit transfer matrix — and two things differ. **The predictor class:** their model rests on attributes of a place, all spatially stationary, ours on the state of a surface in one season. **The response variable:** theirs is human-driven ignition, ours is burned area, and those are different quantities with different dominant controls. Read together the two results suggest that the relationship between domain similarity and transfer success may be predictor-class dependent, with Vesk et al. [38] and Rousseau and Betts [32] aligning from species distribution modelling while Dimarco et al. and WildfireGenome run against it. We present that as a live disagreement, and Section 5.7 states why our own data do not settle it.

Methods, where no land surface temperature or TVDI variable appears and the only temperature predictor is a static ERA5-Land seasonal climatology [26]. We follow their Methods, which we take to be authoritative, and note the discrepancy so that the comparison is transparent. Every statement about their work was checked against the full text, read 2026-08-08.

### 3. Data and methods

#### 3.1. Study regions and temporal windows

Five Mediterranean wildfire regions are analysed: Manavgat 2021 and Muğla 2021 in Türkiye, Bejís 2022 in Spain, North Evia 2021 in Greece

---

<sup>1</sup>Their Results text refers to "LST anomalies", which is inconsistent with their own

and Montiferru 2021 in Sardinia. Each is a place-based rectangular area of  
 290 interest in EPSG:4326, defined from place coverage rather than from a fire  
 perimeter, and not tuned on burned prevalence, gate outcome or any model  
 metric. One AOI choice was label-informed and is stated as such. The North  
 Evia box was extended after the legacy box was found to carry an atypically  
 high burned prevalence. The extended geometry was then defined from place  
 295 anchors, and the legacy variant is kept as a sensitivity arm (Appendix A).  
 The AOIs are deliberately not clipped to fire perimeters, so that unburned  
 cells around each fire form the negative class rather than being excluded by  
 construction. A sixth region, Kozan 2023, is carried as a negative control  
 and is excluded from all modelling by the gate of Section 3.3.

300 Each region has two non-overlapping windows. The **predictor window**  
 closes the day before the **label window** opens, so no predictor observation  
 can post-date the first labelled burning. Window lengths are matched within  
 a region and vary between them with the event: 56 to 60 days for predic-  
 tors and 35 to 49 days for labels. A four-year baseline of window-symmetric  
 305 composites precedes each predictor window and supplies the climatological  
 reference for the anomaly channels. Region identifiers, bounding boxes, win-  
 dow dates and baseline years are registered in `core/regions.py` and are  
 reproduced in Table 1.

### *3.2. Burned-area label and the $\sim 500$ m analysis grid*

310 Labels come from MODIS MCD64A1 Collection 6.1 [9] retrieved through  
 Google Earth Engine [10], whose omission and commission characteristics  
 [3] bound every model here. The analysis grid is **reconstructed rather**  
**than native**: the pipeline’s 30 m EPSG:4326 reference grid is partitioned

into 17 x 17 blocks, giving a nominal 510 m cell that approximates rather  
315 than reproduces the MODIS cell, and is square in degrees but not on the  
ground. A cell's burn date is the mode of its positive sub-pixel day-of-year  
values, tested against the region's label window; the label never affects a cell's  
eligibility for modelling. Two safeguards are recorded rather than assumed:  
cells burning inside a region's own predictor window are removed, which ran  
320 for three of five regions, and burning in earlier years is screened for none.  
Appendix C.1 gives the full specification, including what follows from the  
grid's shape.

### *3.3. Burned-landcover admissibility gate*

Before any modelling each region passes a gate asking what fraction of  
325 its burned cells is dominated by natural vegetation, using ESA WorldCover  
classes on the same cells. A region is admitted at 0.50 with at least 30  
burned cells, and rejected as a cropland-dominated control at 0.50 cropland.  
The purpose is to separate burned area produced by natural-fuel combustion  
from post-harvest stubble burning, which MCD64A1 does not distinguish.  
330 Verdicts are in Section 4.1; the full rule is in Appendix C.1.

### *3.4. Predictor variables*

Ten predictors are used, four baseline and six thermal, all summarised  
per cell over the predictor window. All optical and thermal predictors come  
from Landsat 8 Collection 2 Level-2 (LANDSAT/LC08/C02/T1\_L2), quality-  
335 screened per pixel from the QA\_PIXEL band with the water bit deliberately  
preserved; the coarse-resolution thermal input is MODIS/061/MOD11A1.

- **NDVI**, the predictor-window median of Landsat surface reflectance, is the baseline’s one time-varying member.
- **Elevation** and **slope** come from the Copernicus DEM GLO-30.
- 340 • **Land cover** is ESA WorldCover v200 [45], entering as the dominant class code.
- **Current LST**, the predictor-window median of Landsat surface temperature in °C.
- 345 • **LST anomaly**, a z-score of the current-window LST median against the four baseline years, set to no-data where the baseline standard deviation is below 1.0 °C or observations are too few.
- 350 • **TVDI** and **TVDI difference**. The Temperature-Vegetation Dryness Index [33] is computed in the LST-NDVI feature space, with wet and dry edges taken as the 2nd and 98th LST percentiles within each of twenty NDVI bins and the index clamped to [0, 1]. The edges are percentiles of the LST values a given scene contains, so they are fitted per AOI and per window, and a TVDI of 0.5 denotes a different physical moisture state in each region. The companion paper reports what that scene dependence does and does not explain. **TVDI difference** 355 **ence** is the raw anomaly of that index against the same four window-symmetric baseline years used for the LST anomaly: `tvdi_difference = current_tvdi - mean(baseline_tvdi)`, in index units rather than standard deviations, which is why it is reported alongside the z-scored channel rather than in place of it.

360 • **Downscaled LST** and **fused LST**. A MODIS-to-Landsat downscal-  
 ing model is trained on the predictor window and applied to the full 30  
 m grid. The fused product equals observed Landsat LST wherever that  
 is valid, and the downscaled surface only where it is not. Gap-filling  
 therefore never replaces or blends a valid observation. The gap-filled  
 365 share is 0.11 % to 9.70 % by region. The downscaler’s own inputs in-  
 clude coordinates, which is the one route by which a coordinate-derived  
 surface re-enters a feature set from which Section 3.13 excludes coor-  
 dinates; the companion paper measures what the increment is worth  
 without these two channels.

### 370 3.5. *Cell aggregation, validity and analysis populations*

Cell-level values are means over the  $\sim 510$  m cell, computed from valid  
 30 m pixels only, with the valid fraction recorded. A cell is `valid_for_`  
`modeling` when it meets a 30 % valid-pixel floor and no thermal channel  
 is missing. Two populations are then defined. The **primary** population  
 375 is natural vegetation, cells whose combined tree, shrub and grass fraction  
 reaches 0.50, which excludes cropland from every burnable mask. The **sec-**  
**ondary** population is all valid cells, reported throughout as a sensitivity.  
 Every headline result is given for both.

### 3.6. *Feature sets and classifier*

380 Two nested feature sets are compared. The **baseline** set is NDVI, el-  
 evation, slope and land cover. The **thermal** set is the baseline plus the  
 six thermal channels. Land cover is one-hot encoded and numeric features

are standardised inside the fitting pipeline. The classifier is a random forest [4] with 300 trees, unlimited depth, `min_samples_leaf = 3`, balanced class weights and `random_state = 42`, identical for every region, population, feature set and transfer direction, so that no comparison in this paper is confounded by a model choice.

### 3.7. *Spatial-block cross-validation and bootstrap uncertainty*

Cross-validation is spatially blocked. Blocks are squares of the analysis grid formed as `row_500m // B` by `col_500m // B`, and folds are drawn over whole blocks, so no block is split between training and test. Five folds are used. The block size B is reported at 2, 10 and 20 cells, about 1, 5 and 10 km, for the within-region analysis.

Uncertainty is a spatial-block bootstrap: blocks are resampled with replacement, 1000 replicates, seed 42, with 2.5 and 97.5 percentiles as the interval. Because it is blocks that are resampled, what bounds an interval's reliability is the number of blocks carrying at least one burned cell, and those counts are reported alongside the intervals. Where a verdict rests on too few such blocks it is stated as indicative rather than as an interval.

### 3.8. *Cross-region transfer protocol*

For each ordered pair of regions the model is fitted on the source population and applied to the target population with **no target label used at any point**: no refitting, no threshold selection, no calibration. Both the baseline and the thermal feature set are transferred, so that the thermal block's contribution is a paired difference within a direction rather than a

comparison across directions. All twenty ordered directions among the five regions are computed.

### 3.9. Label-blind domain adaptation

Two label-free remedies are tested on every direction, in both feature sets.

410     **Region-wise z-score.** Each region’s numeric features are standardised using its own statistics, source statistics from source data and target statistics from target data, never pooled. The classifier is refitted on the z-scored source and applied to the z-scored target. This removes first- and second-order marginal offsets.

415     **CORAL after region-wise z-score.** Starting from the z-scored features, the source covariance is aligned to the target covariance by the standard whitening-recolouring map [35], with regularisation  $\lambda = 10^{-5}$  and eigenvalues floored at  $10^{-12}$ . Critically **the transform is applied to the source only**; the target is left exactly as it is. The classifier is refitted on the aligned  
420 source. Neither variant sees a target label, and both are verified label-blind at run time.

$\lambda$  sensitivity was assessed over nine values from 0 to  $10^{-1}$  on four directions in both feature families, moving transfer AUC by at most 0.014 in any direction and 0.008 within the thermal family. The canonical  $\lambda = 1$  lies  
425 outside the released sweep; the companion paper reports it.

### 3.10. Transfer-gap decomposition and the concept-shift diagnostic

For each direction the gap between the target’s own within-region skill and the raw transfer result is split in two. One part is what the best label-free adaptation recovers, and the other is what it does not. The recovered fraction

430 is defined as (adapted – raw) / (within – raw), signed and unclipped, and its interval comes from the same paired bootstrap. The recovered part bounds what covariate-level correction can achieve. The remainder is the conditional residual.

The mechanism is diagnosed by **signed univariate association**. For 435 each numeric predictor the raw ROC-AUC of that predictor against **burned** is computed in each region and never folded to  $\max(\text{AUC}, 1 - \text{AUC})$ , so a value below 0.5 is read as a direction rather than as weakness. A reversal is called bootstrap-supported only when the two regions’ intervals are disjoint, under the same 10-cell spatial-block bootstrap.

### 440 3.11. *Transferability diagnostics versus transfer*

Twenty candidate diagnostics from four families are computed for every region pair and rank-correlated (Spearman) against observed raw transfer AUC, under one bootstrap framework that resamples unordered region pairs with both of their ordered directions travelling together. The families are 445 marginal predictor-space measures  $P(x)$ , including area-of-applicability dissimilarity, climatic and geographic distance and a learned domain classifier; burned-niche overlap  $P(x|y=1)$ ; burn-pattern regime distance  $P(y)$ ; and conditional direction agreement  $P(y|x)$ . No family-wise error control is claimed, and the number of variants computed within each family is reported with 450 the result. Appendix C.2 gives the full specification of each diagnostic and of the bootstrap.

### 3.12. Interventions

**Pooled multi-region training.** Leave-one-region-out: the model is trained on the pooled primary populations of the other four regions and  
455 evaluated on the held-out region, with folds blocked as before. This asks whether pooling recovers what single-source transfer loses.

**Removal of direction-reversing features.** The two predictors whose signed association reverses between regions with bootstrap support are **elevation\_mean** (Manavgat against Bejís and against Muğla) and **lst\_anomaly\_mean**  
460 (Bejís against Evia). They are dropped and everything is refitted, within-region and across every direction, and the two are also dropped singly so the cost can be attributed. This measures what the removal costs locally and what it returns on transfer on the same footing. Note that the two features are selected by the same reversal analysis against which the result is then  
465 read, so both quantities are post-selection estimates and no correction for that selection is applied.

### 3.13. Leakage control

An explicit forbidden-column set is enforced at every model fit. Coordinates (**lon**, **lat**, **row**, **col** and their normalised forms), every burn-date and  
470 label-provenance column, and the agreement fraction are excluded from all feature sets, and the enforcement runs as an assertion rather than a convention. The natural-vegetation mask is used only to define the population, never as a predictor. Spatial blocking prevents a cell from sharing a fold with its own neighbours.

**Reproducibility.** All randomness uses seed 42 and the bootstrap uses  
475 1000 replicates throughout. The transfer and adaptation analysis runs in

an environment separate from the upstream pipeline’s. Every within-region model was therefore refitted there and compared against the frozen upstream output, and the independently implemented adaptation was compared against the pipeline’s own. The within-region comparisons agree exactly, and the twenty transfer directions agree to within  $1.6 \times 10^{-7}$ . Random-forest fits are not bit-identical across library versions. All numbers here were therefore produced under scikit-learn 1.9.0, or verified against it. If that version is not fixed, cross-region point estimates carry an implementation tolerance of about  $\pm 0.02$  to  $0.03$ . The companion paper reports the version sensitivity and the reproduction check in full.

**Sensitivity analyses.** Every headline result is repeated across two analysis populations, three spatial-block sizes, the CORAL regularisation sweep and both feature sets. Where a conclusion depends on one of those choices, the dependence is reported rather than resolved by choosing the favourable setting.

### 3.14. *Same-geography event-to-event comparison (Mugla 2021 versus 2022)*

Mugla admits a comparison in which place is held fixed and the event varies: a second fire burned inside the identical AOI, on the identical analysis grid, eleven months after the first. Signed univariate AUCs are computed for both arms under the same 10-cell spatial-block bootstrap used elsewhere. Season, year and population all differ between the arms, since the 2022 arm is defined by removing the 2021 scar, and Section 4.8 reports what that costs. Appendix C.3 gives the window dates, the population construction and the asymmetry audit.

## 4. Results

### *4.1. Study regions and the admissibility gate*

All five candidate regions pass the burned-landcover gate as wildfire candidates and the negative control fails it as intended. Kozan 2023 returns  
505 a natural-vegetation fraction of 0.017 against the 0.50 threshold and is excluded from all modelling. The separation is not marginal: the five admitted regions carry fractions of 0.723 to 0.991, so the threshold falls in an empty interval rather than between neighbouring cases. That is consistent with the gate separating burned area produced by natural-fuel combustion from post-  
510 harvest stubble burning, which MCD64A1 does not distinguish, though one negative control cannot establish it.

North Evia is analysed on an extended AOI. Relative to the legacy box, the extended box leaves the burned scar essentially unchanged while cutting overall prevalence from 0.361 to 0.122 and TSG prevalence from 0.676 to  
515 0.287. The effect of that choice on transfer is reported in Appendix A(a).

### *4.2. Within-region: the thermal increment replicates in five regions*

Adding the six thermal predictors to the baseline raises spatially blocked out-of-fold ROC-AUC in every region. The increment's bootstrap interval excludes zero in all five regions at 1 km and at 5 km blocking. Those are the  
520 two scales this design supports as intervals. At 10 km the point estimates hold, from +0.048 to +0.154. They rest on 6 to 33 positive-carrying blocks, so they are reported as indicative rather than as intervals. The table note gives the reason.

Table 1: Region summary (recap of Table 1; final numbering at assembly). Counts from each region’s Step 8A dataset statistics; gate fractions from each region’s burned-landcover gate output. TSG = the primary natural-vegetation population. The TSG columns use the canonical modelled population, **burnable\_tree\_shrub\_grass and valid\_for\_modeling == True**, which is the population every model in this paper was fitted and scored on.

| Region                                   | Total cells | Valid cells | Burned | Prevalence<br>(all<br>valid) | TSG cells | Burned in TSG | TSG prevalence | Burned<br>natural-<br>veg<br>frac-<br>tion | Gate verdict |
|------------------------------------------|-------------|-------------|--------|------------------------------|-----------|---------------|----------------|--------------------------------------------|--------------|
| Manavgat<br>2021                         | 24,150      | 24,087      | 796    | 0.033                        | 20,511    | 784           | 0.038          | 0.984                                      | pass         |
| Bejis<br>2022                            | 15,759      | 15,759      | 1,103  | 0.070                        | 15,190    | 1,100         | 0.072          | 0.991                                      | pass         |
| Muğla<br>2021                            | 73,098      | 73,045      | 3,026  | 0.041                        | 41,730    | 2,911         | 0.070          | 0.958                                      | pass         |
| North<br>Evia<br>2021<br>(ex-<br>tended) | 22,925      | 22,906      | 2,788  | 0.122                        | 9,298     | 2,664         | 0.287          | 0.945                                      | pass         |
| Monticelli<br>2021                       | 3,234       | 3,173       | 697    | 0.220                        | 2,544     | 539           | 0.212          | 0.723                                      | pass         |

Table 2: Within-region baseline versus thermal performance and block-size robustness. Primary (TSG) population; spatially blocked 5-fold CV (Section 3.7); paired spatial-block bootstrap, 1000 replicates. Block sizes 2/10/20 cells  $\approx$  1/5/10 km.

| Region                       | Block ( $\approx$ scale) | Baseline AUC | Thermal AUC | $\Delta$ AUC | $\Delta$ AUC 95% CI |
|------------------------------|--------------------------|--------------|-------------|--------------|---------------------|
| Manavgat<br>2021             | 2 ( $\sim$ 1 km)         | 0.803        | 0.870       | +0.067       | [+0.055, +0.079]    |
|                              | 10 ( $\sim$ 5 km)        | 0.748        | 0.797       | +0.050       | [+0.023, +0.077]    |
|                              | 20 ( $\sim$ 10 km)       | 0.683        | 0.731       | +0.048       | [+0.014, +0.085]    |
| Bejís<br>2022                | 2 ( $\sim$ 1 km)         | 0.862        | 0.918       | +0.056       | [+0.048, +0.065]    |
|                              | 10 ( $\sim$ 5 km)        | 0.779        | 0.825       | +0.045       | [+0.018, +0.069]    |
|                              | 20 ( $\sim$ 10 km)       | 0.739        | 0.795       | +0.057       | [+0.031, +0.090]    |
| Muğla<br>2021                | 2 ( $\sim$ 1 km)         | 0.743        | 0.859       | +0.116       | [+0.106, +0.125]    |
|                              | 10 ( $\sim$ 5 km)        | 0.698        | 0.777       | +0.079       | [+0.050, +0.105]    |
|                              | 20 ( $\sim$ 10 km)       | 0.673        | 0.733       | +0.061       | [+0.030, +0.094]    |
| North Evia<br>2021<br>(ext.) | 2 ( $\sim$ 1 km)         | 0.759        | 0.912       | +0.153       | [+0.142, +0.166]    |
|                              | 10 ( $\sim$ 5 km)        | 0.716        | 0.864       | +0.148       | [+0.119, +0.183]    |
|                              | 20 ( $\sim$ 10 km)       | 0.679        | 0.833       | +0.154       | [+0.124, +0.189]    |
| Montiferu<br>2021            | 2 ( $\sim$ 1 km)         | 0.781        | 0.883       | +0.101       | [+0.080, +0.125]    |
|                              | 10 ( $\sim$ 5 km)        | 0.620        | 0.720       | +0.100       | [+0.017, +0.186]    |
|                              | 20 ( $\sim$ 10 km)       | 0.555        | 0.681       | +0.126       | [+0.053, +0.228]    |

*Table note (resampling units).* The bootstrap resamples spatial blocks,  
525 so what bounds an interval’s reliability is the number of blocks that carry  
at least one burned cell. Those counts fall quickly as blocks coarsen. At  
2 cells they are 235 (Manavgat), 302 (Bejís), 843 (Muğla), 716 (Evia) and  
192 (Montiferru), out of 5 439, 3 967, 11 316, 2 566 and 743 blocks. At 10  
cells they are 28, 19, 70, 41 and 16, out of 237, 176, 576, 155 and 35. At  
530 20 cells they are **12, 6, 33, 15 and 6**, out of 60, 48, 167, 50 and 12. No  
bootstrap replicate was invalid at any block size, with one exception. At  
20 cells Bejís had 5 of 1000 replicates single-class. That is the symptom  
of the six positive-carrying blocks just reported, and a further reason to  
read the row as indicative. An equal-tailed percentile interval built on six  
535 positive-carrying blocks has no meaningful coverage. Montiferru at 20 cells  
also feeds only 12 groups into a 5-fold grouped split, so its models train on  
about ten blocks each. **The 20-cell row of this table should be read  
as indicative rather than as an interval.** The 10-cell row is the coarsest  
blocking this design supports properly, because every region there has 16  
540 to 70 positive-carrying blocks. The increment holds at that scale in all five  
regions. Source: `paper/referee2_numbers.md`, block C, counted from the  
frozen per-cell prediction tables.

This within-region result is not itself novel. It is reported because the  
transfer arms below are measured against it.

#### 545 4.3. Cross-region transfer, and what label-free adaptation does to it

**Raw transfer is heterogeneous and includes anti-predictive di-  
rections.** Raw target AUC spans 0.326 to 0.686. Twelve of 20 directions are  
above chance with CI support. Six are *below* chance with CI support: both

directions of Manavgat to Bejís and of Manavgat to Muğla, plus Bejís to  
 550 Evia and Evia to Bejís. Two intervals span 0.5. Even the best raw transfer,  
 Evia to Manavgat at 0.686, stays far below that target’s own within-region  
 thermal performance of 0.870. Across all directions the raw deficit against  
 the within-region reference is 0.184 to 0.592 AUC.

Those counts belong to the 2-cell blocking of Table 5. At the more conser-  
 555 vative 10-cell blocking the same points give 9 above, 4 below and 7 uncertain,  
 and no direction changes side of the chance line (Section 4.7). Four of the  
 six below-chance directions keep their support there. Bejís to Manavgat and  
 Manavgat to Muğla lose it and carry no verdict. The qualitative statement  
 is unchanged. The counts should not be read as exact.

560 **The static baseline does not transfer either.** This is the control  
 for the reading the rest of the paper invites, and it constrains that reading  
 sharply. The same twenty directions were run with the terrain, fuel and  
 greenness baseline alone. The mean target AUC is **0.537**, against **0.541**  
 for the thermal model. Static attributes of a place are the class Dimarco  
 565 et al. transfer successfully, and are the class this paper’s framing treats as  
 portable. Here that class is itself barely above chance. The paired per-  
 direction contrast is given in Section 4.6b. What matters here is that the  
 transfer failure below is not specific to the dynamic block. A baseline that  
 does not travel, plus pre-fire thermal state, gives a model that does not travel.

570 **The thermal block’s paired contribution to transfer, with its  
 interval.** The two matrices were differenced direction by direction. The  
 mean is **+0.004**. The directions are not independent, because each region  
 appears in eight of the twenty. The interval therefore depends on what is

575 treated as the resampling unit. All four units the design permits give the same answer:

Table 3:

| Resampling unit                    | n  | 95 % interval on the mean paired contribution |
|------------------------------------|----|-----------------------------------------------|
| Directions, naive                  | 20 | $[-0.027, +0.034]$                            |
| Unordered pairs, cluster bootstrap | 10 | $[-0.028, +0.036]$                            |
| Unordered pairs, t on pair means   | 10 | $[-0.034, +0.042]$                            |
| Regions, leave-one-out jackknife   | 5  | $[-0.037, +0.046]$                            |

Every interval spans zero, and the point estimate is a small fraction of each width. The per-region jackknife shows how little the mean is anchored. Holding out Manavgat, Bejís, Muğla, Evia and Montiferru in turn gives +0.0148, +0.0021, +0.0065, **-0.0081** and +0.0060. **Dropping Evia alone**  
580 **reverses the sign of the headline.** None of these four units propagates within-direction sampling variability. They resample between directions only.

**Crossing a region boundary adds nothing detectable to a loss already incurred inside the region.** Two controls establish this, and the section is careful about what they do and do not support.

585 The first asks whether this evaluation can register transfer at all. Within each region the modelled cells were cut in half by a straight line, a model fitted on one half and applied to the other with no refit, no recalibration and no threshold selection, which is the protocol the cross-region arms use. The relationship is shared by construction: same region, same season, same  
590 fire, same units. Fitted through the same code path, which reproduces the frozen cross-region AUCs to 0.000000, the eighteen usable splits give a mean

of **0.574**, ranging from 0.294 to 0.750. The harness is therefore not broken, since it returns 0.750 on the easiest split.

**That 0.574 is not a clean measure of separation.** A straight cut  
 595 does not produce two exchangeable halves. Source burned counts across the  
 eighteen splits run from 84 to 2,564, one region contributes a split with 700  
 source positives against another with 84, and two further splits were unusable  
 because one half of Manavgat contains no burned cells at all. Blocked cross-  
 validation, by contrast, trains on about four fifths of a region with positives  
 600 drawn from across the whole scar. The drop from the blocked-CV range of  
 0.859 to 0.918 down to 0.574 therefore mixes separation with training-set  
 composition and with where a single scar happens to fall, and this design  
 does not separate them.

The second control asks how skill varies with separation. Because each  
 605 cut is a straight line, a target cell’s distance to the nearest training cell is its  
 distance to that line, so one fitted model per split yields an AUC at every  
 separation:

Table 4:

| Separation from the<br>training half | Bins | Burned cells | Mean AUC | Weighted by burned<br>cells |
|--------------------------------------|------|--------------|----------|-----------------------------|
| 0 to 5 km                            | 18   | 4,504        | 0.692    | 0.674                       |
| 5 to 10 km                           | 16   | 3,369        | 0.519    | 0.574                       |
| 10 to 20 km                          | 11   | 3,740        | 0.499    | 0.494                       |
| 20 to 40 km                          | 6    | 2,215        | 0.445    | 0.495                       |
| 40 to 80 km                          | 3    | 1,357        | 0.541    | 0.555                       |
| 80 to 160 km                         | 1    | 27           | 0.421    | 0.421                       |

Skill declines with separation and is near chance from about 10 km outward. Three cautions attach. The near bins are inflated by autocorrelation, which is what blocked validation exists to remove, so 0.692 is an upper bound rather than an estimate. The far bins are thin, one of them resting on 27 burned cells, and the profile is not monotonic. And the two aggregations disagree by up to 0.05, so the level of the plateau is sensitive to a choice the data do not settle.

**What this supports, and what it does not.** Placing the twenty transfer directions at their AOI-centroid separations of 306 to 2,802 km gives a mean of 0.541, which is where the within-region curve already sits by 10 to 20 km. The defensible statement is therefore a **null**: no additional loss attributable to region crossing is detectable on top of a loss already largely incurred inside a region. It is not evidence that distance *causes* the cross-region deficit. Once the within-region curve is at the chance floor, any cross-region mean near 0.5 lies on its continuation, so this comparison could not have come out otherwise, and the two ranges do not overlap: within-region separations stop at 86 km and cross-region ones begin at 306.

**Separation also does not order the matrix, and something else must.** Every cross-region direction is past the saturation, so separation is effectively constant across the matrix while transfer runs from 0.326 to 0.686. Tested directly, geographic distance does not order transfer, at Spearman  $\rho = -0.32$  with an interval spanning zero, and the two nearest directions, Manavgat and Muğla at 306 km, are among the worst at 0.470 and 0.401 while directions eight times more distant transfer above chance. **Six directions are below chance with interval support**, the sharpest at 0.326 [0.305,

0.349]. Extrapolating an uninformative model produces about 0.5, not a reliably reversed ranking, so those six require a mechanism that acts on the direction of the relationship rather than on its strength. That is what Sections 4.4 to 4.8 are about. Full detail on both controls is in `paper/positive_control.md` and `paper/distance_curve.md`.

**Label-blind adaptation compresses the matrix toward chance rather than repairing it.** Under region-wise z-scoring the twenty directions span 0.431 to 0.630 and under CORAL 0.443 to 0.624, roughly half the raw spread, with no adapted direction exceeding 0.631 against within-region references of 0.859 to 0.918. Adaptation raises the failing directions and degrades most of those that already transferred. Taking the better of the two adaptations per direction, 14 of the 20 end closer to chance than they began and 6 end further from it; five of those six involve Montiferru, the smallest and last-added region, and move upward, while the sixth is Manavgat→Muğla moving downward from 0.470 to 0.443. The 14 to 6 split should be read at the precision of limitation (ix) in Section 5.9, since Bejís→Manavgat is counted as compressed on a margin of 0.001.

In the six directions where raw transfer was below chance, the best label-free method recovers at most 34 % of the gap to the within-region reference, so the remaining unrecovered fraction is at least 0.66 everywhere. The half-split control above shows that the larger part of that remainder is already incurred inside the region, so it should not be read as a measure of concept shift. Seven directions show *negative* recovery, meaning adaptation moves the score away from the reference; in six of those raw transfer was already above chance and adaptation destroyed that advantage. Label-free alignment

Table 5: Cross-region transfer matrix, thermal model, TSG population. Target ROC-AUC with 2-cell spatial-block bootstrap 95% CIs (1000 replicates). CORAL is applied after region-wise z-scoring ( $\lambda = 10^{-5}$ ).

| Direction           | Raw                  | Region-wise z-score  | CORAL                |
|---------------------|----------------------|----------------------|----------------------|
| Manavgat→Bejís      | 0.326 [0.305, 0.349] | 0.477 [0.451, 0.502] | 0.511 [0.484, 0.534] |
| Bejís→Manavgat      | 0.444 [0.408, 0.480] | 0.457 [0.420, 0.497] | 0.555 [0.528, 0.583] |
| Manavgat→Muğla      | 0.470 [0.451, 0.488] | 0.431 [0.411, 0.449] | 0.443 [0.423, 0.461] |
| Muğla→Manavgat      | 0.401 [0.378, 0.426] | 0.559 [0.531, 0.587] | 0.560 [0.535, 0.587] |
| Manavgat→Evia       | 0.613 [0.593, 0.631] | 0.542 [0.520, 0.565] | 0.539 [0.518, 0.561] |
| Evia→Manavgat       | 0.686 [0.653, 0.716] | 0.516 [0.489, 0.544] | 0.527 [0.500, 0.553] |
| Bejís→Muğla         | 0.618 [0.601, 0.635] | 0.518 [0.501, 0.535] | 0.507 [0.489, 0.524] |
| Muğla→Bejís         | 0.583 [0.561, 0.607] | 0.535 [0.512, 0.557] | 0.560 [0.538, 0.581] |
| Bejís→Evia          | 0.383 [0.363, 0.402] | 0.532 [0.509, 0.551] | 0.499 [0.479, 0.518] |
| Evia→Bejís          | 0.448 [0.426, 0.470] | 0.549 [0.524, 0.575] | 0.549 [0.525, 0.573] |
| Muğla→Evia          | 0.653 [0.636, 0.671] | 0.561 [0.543, 0.580] | 0.563 [0.545, 0.582] |
| Evia→Muğla          | 0.577 [0.560, 0.593] | 0.501 [0.485, 0.518] | 0.530 [0.515, 0.546] |
| Montiferru→Manavgat | 0.567 [0.539, 0.594] | 0.573 [0.538, 0.609] | 0.606 [0.574, 0.639] |
| Manavgat→Montiferru | 0.533 [0.488, 0.580] | 0.586 [0.540, 0.629] | 0.592 [0.550, 0.631] |
| Montiferru→Bejís    | 0.548 [0.521, 0.578] | 0.574 [0.552, 0.596] | 0.569 [0.548, 0.591] |
| Bejís→Montiferru    | 0.594 [0.560, 0.631] | 0.550 [0.500, 0.601] | 0.574 [0.530, 0.621] |
| Montiferru→Muğla    | 0.619 [0.604, 0.634] | 0.576 [0.562, 0.589] | 0.565 [0.550, 0.579] |
| Muğla→Montiferru    | 0.531 [0.495, 0.568] | 0.587 [0.549, 0.624] | 0.584 [0.547, 0.623] |
| Montiferru→Evia     | 0.586 [0.565, 0.606] | 0.630 [0.611, 0.649] | 0.624 [0.605, 0.641] |
| Evia→Montiferru     | 0.647 [0.608, 0.682] | 0.568 [0.528, 0.609] | 0.581 [0.539, 0.623] |

Table 6: Transfer-gap decomposition (four-AOI set, 12 directions). Within = target’s within-region thermal AUC; best adapted = the better of z-score/CORAL; recovered fraction = (adapted – raw)/(within – raw), signed and unclipped, with paired bootstrap CI (1000 replicates). Montiferru directions are not part of this decomposition (per-pair absolute decompositions exist without fraction CIs). The status column asks whether the *adapted* value clears chance and uses the 2-cell adapted intervals of Table 5. The adapted arms were not recomputed at the coarser blocking of Appendix A(c), which covers the raw arm and the paired delta only.

| Direction      | Within | Raw   | Best adapted<br>(method) | Recovered<br>fraction [CI] | Status                              |
|----------------|--------|-------|--------------------------|----------------------------|-------------------------------------|
| Manavgat→Bejís | 0.918  | 0.326 | 0.511<br>(CORAL)         | +0.31 [+0.28,<br>+0.34]    | recovery,<br>chance not<br>excluded |
| Bejís→Manavgat | 0.870  | 0.444 | 0.555<br>(CORAL)         | +0.26 [+0.18,<br>+0.34]    | recovery<br>above chance            |
| Muğla→Manavgat | 0.870  | 0.401 | 0.560<br>(CORAL)         | +0.34 [+0.28,<br>+0.40]    | recovery<br>above chance            |
| Bejís→Evia     | 0.912  | 0.383 | 0.532<br>(z-score)       | +0.28 [+0.23,<br>+0.32]    | recovery<br>above chance            |
| Evia→Bejís     | 0.918  | 0.448 | 0.549<br>(z-score)       | +0.22 [+0.16,<br>+0.27]    | recovery<br>above chance            |
| Manavgat→Muğla | 0.859  | 0.470 | 0.443<br>(CORAL)         | −0.07 [−0.12,<br>−0.03]    | <b>negative</b><br><b>recovery</b>  |
| Muğla→Bejís    | 0.918  | 0.583 | 0.560<br>(CORAL)         | −0.07 [−0.16,<br>+0.01]    | <b>negative</b><br><b>recovery</b>  |
| Manavgat→Evia  | 0.912  | 0.613 | 0.542<br>(z-score)       | −0.23 [−0.32,<br>−0.14]    | <b>negative</b><br><b>recovery</b>  |
| Evia→Manavgat  | 0.870  | 0.686 | 0.527<br>(CORAL)         | −0.86 [−1.18,<br>−0.60]    | <b>negative</b><br><b>recovery</b>  |
| Evia→Muğla     | 0.859  | 0.577 | 0.530<br>(CORAL)         | −0.17 [−0.22,<br>−0.11]    | <b>negative</b><br><b>recovery</b>  |
| Muğla→Evia     | 0.912  | 0.653 | 0.563<br>(CORAL)         | −0.35 [−0.43,<br>−0.27]    | <b>negative</b><br><b>recovery</b>  |
| Bejís→Muğla    | 0.859  | 0.618 | 0.518<br>(z-score)       | −0.42 [−0.51,<br>−0.34]    | <b>negative</b><br><b>recovery</b>  |

therefore does not act as a repair mechanism.

#### 4.4. *Transferability diagnostics: only conditional similarity orders transfer*

660 Twenty candidate diagnostics from four families were each rank-correlated with the same target quantity, the raw thermal transfer AUC over the twenty ordered directions, under one common pair-based bootstrap.

**Only two diagnostics have intervals excluding zero, and both are conditional.** The stronger is the sign-agreement fraction over interval-supported features, at  $\rho = +0.84$  [ $+0.58, +0.88$ ]. The cosine variant reaches 665  $+0.81$ . No marginal measure was shown to order the matrix. That includes area-of-applicability-style dissimilarity in predictor space, climatic distance and geographic distance. The niche-overlap and regime families were not shown to order it either. The learned domain classifier is at ceiling, separating 670 source from target at  $\text{AUC} \geq 0.96$  for every pair. It always succeeds, which is why it carries no ordering information.

Two limits are stated with the result rather than after it. The first is size. The index's tie structure caps the achievable Spearman at  $+0.861$ , so the observed  $+0.840$  sits essentially on that ceiling. Its exact one-sided 675 permutation p is  $0.0060$ , the smallest this tie structure can produce, against a Bonferroni threshold of  $0.0026$  over the nineteen computed variants. No outcome of this diagnostic could have cleared family-wise correction on ten effective pairs. The second limit is labels. Signed associations need burned labels in both regions. The family that works is therefore not available before 680 deployment, while the family that fails is.

A third limit is about selection, and it is the sharpest of the three. The two rows that clear zero are the *supported-feature* variants. Their feature

Table 7: Transferability diagnostics versus raw thermal transfer, by family. Spearman  $\rho$  against raw transfer AUC with pair-based bootstrap 95 % CIs. Exp. is the sign expected if the diagnostic orders transfer. The member named is the one with the largest absolute correlation in its family, which is not always in the expected direction. All twenty individual diagnostics are in Appendix B, Table B1.

| Family                                          | Diagnostics | Exp. | Largest correlation in family                                                                              | Sign as expected | Any CI excluding 0 |
|-------------------------------------------------|-------------|------|------------------------------------------------------------------------------------------------------------|------------------|--------------------|
| <b>P(y x)</b><br><b>condi-</b><br><b>tional</b> | 6           | +    | <b>+0.84</b><br>[ <b>+0.58</b> ,<br><b>+0.88</b> ]<br>agree-<br>ment<br>fraction,<br>supported<br>features | yes              | <b>yes, 2 of 6</b> |
| P(x y=1)<br>niche<br>overlap                    | 5           | +    | +0.24<br>[−0.45,<br>+0.74]<br>Schoener’s<br>D, 1-D<br>mean                                                 | yes              | no                 |
| P(x)<br>marginal                                | 6           | −    | −0.32<br>[−0.78,<br>+0.33]<br>domain-<br>classifier<br>AUC                                                 | yes              | no                 |
| P(y)<br>regime<br>structure                     | 2           | −    | +0.29<br>[−0.38,<br>+0.74] <sup>log</sup><br>effective-<br>N<br>distance                                   | <b>no</b>        | no                 |
| geographic                                      | 1           | −    | −0.24<br>[−0.84,                                                                                           | yes              | no                 |

subset is chosen by whether two regions' bootstrap intervals happen to be disjoint. That is a data-dependent selection, made on the same data, with  
685 no correction. The unselected counterparts over all nine features are  $\rho = +0.50$   $[-0.17, +0.83]$  for the cosine and  $\rho = +0.18$   $[-0.40, +0.72]$  for the agreement count. Both span zero. The result lives in the selection step, and is reported as such.

**The families are compared on unequal samples, and equalis-**  
690 **ing them does not change the ordering.** The marginal, applicability, climatic and geographic rows sit on twelve directions, because those diagnostics exist only for the four-region subset. The supported-conditional rows sit on sixteen and the rest on twenty. A reader may reasonably ask whether the marginal family's failure is a statement about power rather than about  
695 diagnostics. Every row was therefore recomputed on the common twelve directions. The published values reproduce to  $4.8 \times 10^{-5}$ . The conditional rows still lead, at  $+0.87$   $[+0.65, +0.88]$  and  $+0.85$   $[+0.43, +0.88]$ , and every marginal row still spans zero. The ordering is not an artefact of unequal samples.

#### 700 4.5. *The contrast pair: similarity is not sufficient*

The clearest single view of Table 7 needs no ranking at all. Manavgat and Muğla lie in the same country and the same fire year. They are 307 km apart by the centroid geodesic distance this paper uses as a diagnostic, and their nearest boundaries are 191 km apart. Their burned cells occupy  
705 the most similar environmental envelope of any pair in the matrix, with per-feature Schoener's D of 0.77 to 0.89. Yet five of nine feature-response directions point opposite ways. Elevation is one of them, at a signed AUC of

0.374 [0.289, 0.471] in Manavgat against 0.611 [0.532, 0.690] in Muğla, with disjoint intervals. All four absolute thermal channels are among the others.  
 710 Transfer is below chance in both directions at the point estimate, at 0.470 and 0.401. Both directions also sit inside the nominal area of applicability, at 0.875 and 0.531 of target cells.

Bejís and Montiferru sit at the opposite extreme. Their burned envelopes barely overlap, and they carry the most dissimilar values on every overlap  
 715 measure. They transfer above chance in both directions. That half of the contrast rests on point estimates: neither direction carries a verdict at 5 km blocking, and the marginal applicability audit was never produced for Montiferru. The claim is one of *sufficiency*. Similarity does not guarantee transfer, and dissimilarity does not preclude it. The claim rests on two  
 720 coexisting counterexamples, so it does not depend on the number of pairs available.

#### 4.6. Interventions: pooling and feature removal

**(a) Pooled multi-region training.** At the point estimate, training on the pooled primary populations of the other four regions never beats the best  
 725 single-source transfer for any target, with shortfalls of 0.02 to 0.22, and it stays 0.28 to 0.50 AUC below the within-region ceiling; for two targets it falls below the pairwise mean and below chance, though only Bejís is below chance with interval support, at 0.417 [0.369, 0.467]. Aggregation does not recover what single-source transfer loses, and since it is evaluated entirely  
 730 outside the saturation radius its failure is what separation alone predicts.

**(b) Removing the direction-reversing features.** The two predictors whose signed association reverses between regions **with bootstrap sup-**

port are **elevation\_mean** and **lst\_anomaly\_mean**. The per-region signed associations and the three pair-level reversals that meet that criterion are  
735 in Appendix B, Tables B2 and B3. Retraining without them costs **-0.081** of mean within-region AUC, supported in every region (per-region deltas  $-0.060$ ,  $-0.130$ ,  $-0.074$ ,  $-0.063$ ,  $-0.079$ , every interval entirely below zero), and changes mean transfer by **+0.014** [**-0.017**, **+0.045**], which spans zero.

Two qualifications belong with those numbers rather than after them.  
740 First, the debit is not the thermal block’s. Dropping elevation alone accounts for  $-0.061$  of it, from 0.888 to 0.827. Dropping the LST anomaly alone accounts for  $-0.013$ , from 0.888 to 0.875. Roughly three quarters of the cost is therefore the removal of a **baseline** terrain variable, not of a thermal channel. Second, both figures are post-selection. The two predictors were  
745 chosen because they reverse, using the same data on which the  $-0.081$  and the  $+0.014$  are then estimated. No correction for that selection is applied.

Table 8:

| Configuration                | Mean within-region AUC | Mean transfer AUC |
|------------------------------|------------------------|-------------------|
| full                         | 0.888                  | 0.541             |
| drop <b>elevation_mean</b>   | 0.827                  | 0.546             |
| drop <b>lst_anomaly_mean</b> | 0.875                  | 0.544             |
| drop both                    | 0.807                  | 0.556             |

What this measures is a local cost with no compensating transfer gain. It does not measure an exchange, because the transfer side is a null on both arms. The thermal block’s own contribution to transfer is  $+0.004$ , with  
750 an interval spanning zero (Section 4.3). Removing the reversing predictors returns  $+0.014$ , with an interval spanning zero. Two nulls on the portability

axis are not a price paid.

#### *4.7. Sensitivity analyses*

Six design choices were varied with everything else held fixed. They are  
755 the Evia AOI and its prevalence, the CORAL regularisation constant, the  
blocking scale, the closure date of the predictor window, the quality screening  
of the coarse thermal input, and the contrast between the normalised and  
the absolute dryness channels. None changes a conclusion above. Two bound  
how the results should be read, so they are carried into the main text here.

760 Coarsening the blocks from 1 km to 5 km moves the verdict counts from  
ten positive, seven negative and three uncertain to six, four and ten. Support  
is removed from seven verdicts and added to none. The point estimates are  
unchanged, but that is an identity rather than a result. The blocking scale  
is the resampling unit, and it cannot move an estimate computed once over  
765 all target cells.

Manavgat’s whole downstream chain was then rebuilt from a quality-  
screened MODIS input. That changes the downscaled surface on 22,304 of  
24,150 cells, by up to 10.9 °C. No signed univariate association moves by  
more than +0.0003. This closes the one processing-artefact candidate for  
770 that region’s behaviour. Appendix A reports all six arms, including one that  
tests a claim of Section 1.2 and does not uphold it.

#### *4.8. The same geography, a second fire: reversal with place held constant*

Every result above compares different places. Muğla admits a stricter test,  
because a second fire burned inside the identical AOI on the identical analysis  
775 grid. The 2021 event has a predictor window closing on 28 July. The 2022

event has a matched window closing on 20 June. Region, bounding box, cell definition, feature registry and processing chain are the same in both arms. The design holds place fixed. It does not hold season fixed, because the 2022 fire ignites about six weeks earlier, and it does not hold the population fixed, because the 2022 arm is defined by removing the 2021 scar.

**Elevation reverses with bootstrap support.** In 2021 higher ground burned preferentially, at 0.611 with its interval entirely above 0.5; in 2022 lower ground did, at 0.296 with its interval entirely below. The two intervals are disjoint, 0.532 against 0.355, and the difference is  $-0.317$   $[-0.414, -0.220]$ . The same predictor, the same region, the same grid, and an association that points the opposite way in two fires eleven months apart. The four absolute thermal channels change side as well, from bootstrap-supported lower-values-burn in 2021 to higher-values-burn point estimates in 2022.

Two structural properties must be read alongside these numbers, because they have no analogue in the twenty-direction matrix. The two arms are not disjoint samples: they share 38,789 of the 2022 arm’s 38,790 cells, and elevation, slope and land cover are identical to the digit across all 73,098 grid cells, so only NDVI and the six thermal channels carry new information between them. And the removal is the target’s own positive class, so in the 2022 to 2021 direction not one target positive is present in the source training population while 38,789 of 38,819 target negatives are, and membership of the source training set alone separates the 2021 target’s classes at ROC-AUC 0.9996. We therefore cannot bound what that asymmetry does to the two transfer numbers, only state that a direction with this structure is not comparable to the twenty between-region directions. The elevation reversal is

Table 9: Signed univariate feature-burned AUC, Muğla 2021 versus 2022. Raw AUC against **burned**, never folded to  $\max(\text{AUC}, 1 - \text{AUC})$ ; 10-cell ( $\approx 5$  km) spatial-block bootstrap, 1,000 replicates, seed 42 (Section 3.14). Analysis population 41,730 rows / 2,911 burned (2021) and 38,790 rows / 331 burned (2022). **Positive-carrying 5 km blocks: 70 for the 2021 arm and 11 for the 2022 arm.** Table 2’s note sets sixteen as the floor this design supports at that blocking, so the 2022 intervals here fall below the paper’s own standard and are read as indicative, exactly as the 20-cell row of Table 2 is. The 2022 arm is additionally a single compact scar, so its eleven blocks are contiguous. Read from `paper/step9g_raw/.../mugla_2021__mugla_2022_event_relative/step9g_direction_reversal_table.csv`.

| Feature               | 2021 AUC [95 %<br>CI]           | 2022 AUC [95 %<br>CI]           | Reversal                        |
|-----------------------|---------------------------------|---------------------------------|---------------------------------|
| <b>elevation_mean</b> | <b>0.611 [0.532,<br/>0.690]</b> | <b>0.296 [0.230,<br/>0.355]</b> | <b>bootstrap-<br/>supported</b> |
| current_lst_mean      | 0.325 [0.271,<br>0.382]         | 0.515 [0.434,<br>0.580]         | point only                      |
| current_tvdi_mean     | 0.336 [0.275,<br>0.398]         | 0.594 [0.475,<br>0.674]         | point only                      |
| downscaled_lst_mean   | 0.307 [0.253,<br>0.366]         | 0.508 [0.435,<br>0.571]         | point only                      |
| fused_lst_mean        | 0.325 [0.272,<br>0.383]         | 0.519 [0.436,<br>0.583]         | point only                      |
| ndvi_mean             | 0.662 [0.616,<br>0.704]         | 0.707 [0.624,<br>0.777]         | none                            |
| slope_mean            | 0.637 [0.582,<br>0.686]         | 0.558 [0.468,<br>0.634]         | none                            |
| lst_anomaly_mean      | 0.485 [0.395,<br>0.566]         | 0.380 [0.249,<br>0.502]         | none                            |
| tvdi_difference_mean  | 0.490 [0.396,<br>0.575]         | 0.397 [0.265,<br>0.514]         | none                            |

the finding; the structural asymmetry is a competing explanation this design cannot exclude. The known part of the bias runs the safe way: the removed cells are high, with a median of 563 m, and unburned in 2022, so removing them raises the 2022 elevation AUC and makes the reversal smaller rather  
805 than larger.

#### *4.9. What target labels cost: the recovery curve*

Everything above measures a failure; this prices it. Label-free alignment does not close the residual gap, so the missing resource is information about the target that alignment cannot synthesise, and the direct way to supply it  
810 is target labels. A frozen few-shot diagnostic answers how many, for three regions across all six ordered directions, using one 10-cell ( $\sim 5$  km) spatial block as the unit of labelling effort and reading recovery against a matched target-only ceiling of 0.777 to 0.824.

Thirty-two labelled blocks recover **85 to 89 % of that ceiling in three**  
815 **of the six directions**, two of which started below chance; 51 to 57 % in two more, the directions where the conditional gap is widest; and 30 % in the sixth. That budget is 7 to 20 % of the target's natural-vegetation population, about 800 km<sup>2</sup>, so it is a real answer and not a cheap one. Two properties of the measurement bound it and are stated in the supplement rather than  
820 buried: at the top budget the selection pool is nearly exhausted, so the narrow upper-budget intervals reflect saturation rather than precision, and the labelled blocks are drawn from the event being predicted, which is not a resource available before that event burns. The protocol, the full curve and the limits are in Supplement S1.

825 Two limits belong with the number. At small budgets the same interven-

tion damages the direction that already transfers best without any labels. And six directions in three regions cannot support a general label budget, so none is offered.

## 5. Discussion

### 830 5.1. Principal findings

Three results carry this paper. First, the pre-fire thermal block raises spatially blocked within-region ROC-AUC by +0.056 to +0.153 in every one of five Mediterranean regions, with bootstrap support at 1 km and 5 km blocking, and contributes +0.004  $[-0.028, +0.036]$  across twenty ordered  
835 transfer directions, indistinguishable from zero and with a sign that varies by pair. The static baseline transfers no better, at a mean of 0.537 against 0.541, so the failure is not a property of the dynamic block specifically. Nor is any part of it detectably attributable to region crossing: a within-region half-split returns 0.574 and skill is near chance from about 10 km outward,  
840 so the cross-region mean of 0.541 adds nothing measurable to a loss already incurred inside a region. Separation does not order the matrix, though, and the six below-chance directions are not what extrapolation produces. Removing the two reversing predictors, elevation and the LST anomaly, costs  $-0.081$  of within-region skill, supported in every region and three quarters  
845 attributable to elevation, and changes transfer by +0.014  $[-0.017, +0.045]$ , which also spans zero. Second, of twenty candidate transferability diagnostics only two have intervals excluding zero and both are conditional, while no marginal measure was shown to order the matrix. The marginal family, which includes area-of-applicability dissimilarity, is the only one runnable

850 before deployment. Third, the mechanism is a reversal in the sign of the predictor-burning association, and it persists inside one study area across two fires.

### 5.2. *Why the thermal increment is real but local*

A distinction has to be drawn before this section can say anything useful. 855 The within-region increment and the transfer failure are measured at different separations, and Section 4.3 shows that most of the difference between them is already present inside a single region. What follows therefore reads the increment as local in the literal sense: it holds where training and evaluation cells are interleaved, and it is largely gone by about 10 km, whether or not 860 a region boundary has been crossed. What that does not explain is why the directions differ from each other, and in particular why six of them are anti-predictive.

The within-region increment is not an artefact to be explained away. It replicates in five independent regions, survives coarsening of the spatial 865 blocks to  $\sim 5$  km with its interval intact, persists in the secondary all-valid population, and survives a predictor window closed 7 and 14 days earlier in all five regions. The direction of that last result is region-specific and should not be overstated: the contribution strengthens in Bejís ( $0.058 \rightarrow 0.079$  at 14 days) and Muğla ( $0.115 \rightarrow 0.128$ ), is flat in Montiferru, and **weakens** 870 **monotonically in Evia** ( $0.156 \rightarrow 0.149 \rightarrow 0.135$ ). What holds everywhere is survival, not improvement. The transfer failure is therefore not evidence that the thermal signal is spurious. It is evidence that the fitted relationship is *local*.

The natural objection is that different Mediterranean regions are simply

875 different systems, so that a predictor meaning one thing in one place and  
another elsewhere is not instability but a comparison of two different systems.  
The two Muğla events answer it directly: region, AOI, grid, feature registry  
and processing chain identical, static predictors identical cell by cell, and  
elevation’s association with burning reversing from 0.611 [0.532, 0.690] to  
880 0.296 [0.230, 0.355], a difference of  $-0.317$  [ $-0.414$ ,  $-0.220$ ] (Section 4.8).  
Holding geography fixed does not stabilise the direction of the relationship.

Three caveats attach and none is dismissed. Season and year are con-  
founded, and **that confound cannot be resolved in this study area**:  
the events sit 42 days apart in median burn day-of-year, neither year con-  
885 tains a second event at the other’s phase, and a calendar-matched arm would  
carry nine burned cells against a gate minimum of thirty. The population is  
not held fixed either, the 2022 arm being the 2021 arm with the 2021 scar  
removed. And the 2022 arm rests on one compact scar and eleven positive-  
carrying 5 km blocks. What survives all three is the narrower reading that  
890 *something* reverses the elevation-burning association with geography, grid  
and pipeline held fixed, which is why Section 4.8 is reported as corroboration  
rather than as the load-bearing evidence for the mechanism.

The sharpest supported reversal belongs to elevation, a static predictor,  
so instability is a property of the predictor-to-burning mapping generally  
895 rather than of thermal channels specifically. This is the same conclusion the  
baseline transfer arm reaches from the other direction: the instability is not  
the thermal block’s peculiarity.

### 5.3. *Why the similarity-based diagnostics fail*

A diagnostic built on distance in predictor space asks one question: do  
900 the target’s predictor values look like the training data’s? That is not the  
question that matters when the failure is conditional. A target region can  
sit well inside the training envelope while the relationship between those  
predictors and burning points the other way. Manavgat to Muğla is exactly  
that case, with 0.875 of target cells inside the weighted area of applicability  
905 and transfer below chance.

The same holds for the niche-overlap and regime families. At the point  
estimates, the pair with the highest burned-niche overlap in the matrix  
fails in both directions while the pair with the lowest transfers in both, so  
no monotone function of overlap can order the outcomes. The qualifier is  
910 not decorative: at the 5 km blocking this paper otherwise defends, neither  
Bejís↔Montiferru direction carries a verdict and Manavgat→Muğla loses its  
below-chance support, so this contrast is a statement about point estimates  
and is made as one. The domain classifier is at ceiling for every pair, which  
is informative in itself: the regions are trivially separable in predictor space  
915 whether or not their transfer works, so separability carries no ordering infor-  
mation here.

### 5.4. *What the conditional diagnostic is, and what it is not*

The sign-agreement index reaches  $\rho = +0.84$  [+0.58, +0.88] against ob-  
served transfer. Three limits are reported with it rather than after it. Its  
920 tie structure caps the achievable Spearman at +0.861, so the observed value  
sits essentially on its own ceiling and the ranking it supports is coarse. Its  
exact one-sided permutation p is 0.0060, the smallest that tie structure can

produce, against a Bonferroni threshold of 0.0026 over nineteen computed variants. No outcome could have cleared family-wise correction on ten effective pairs. And it needs burned labels in both regions, which makes it a mechanism diagnosis rather than a pre-deployment screen.

That last limit is the substantive one, and it is why the paper’s practical conclusion is not "use this index instead". It is that the family of diagnostics that can be run before deployment is the family that fails here, and the information that would order transfer is the information a practitioner does not have when the decision is made. The honest statement is that these twenty candidates were **not shown to order transfer**, not that they are shown incapable of it.

### 5.5. *What the two interventions do and do not show*

Both interventions show the same shape. Pooling four regions never beats the best single-source transfer for any target and stays 0.28 to 0.50 below the within-region ceiling, so aggregation does not manufacture the missing conditional information. Removing the reversing predictors costs  $-0.081$  of within-region skill with interval support in every region, and changes mean transfer by  $+0.014$   $[-0.017, +0.045]$ , which is not distinguishable from zero.

That pair of numbers is easy to read as an exchange, and it is not one. The thermal block’s own contribution to transfer is  $+0.004$ , with an interval spanning zero. Removing the reversing predictors returns  $+0.014$ , with an interval spanning zero. Both arms are null on the portability axis. What the interventions measure is a local cost and no compensating transfer gain. That is not a conservation law, and it is not a rate at which local skill can be sold for portability. No such rate is estimated here.

### 5.6. *The regime hypothesis, reported as it happened*

A regime-structure explanation was stated in advance and the data confirmed the null. The regime-distance correlation has the wrong sign at the point estimate, the most regime-similar pair fails in both directions, and the most regime-different pair transfers above chance. One mundane explanation can be set aside: cutting Muğla, by far the largest population, to Manavgat's cell count and separately to its positive count leaves Muğla's transfer behaviour inside the subsampling range in both roles. The error was in the hypothesised grouping, not in the data, and with ten pairs this cannot refute regime typology [2] in general.

### 5.7. *The empirical contrast with Dimarco et al.*

Dimarco et al. [7] transfer successfully across a comparable Mediterranean design and we do not, and the two results are not in conflict. Their predictors are attributes of a place, ours describe the state of a surface in one season, which suggests that the relationship between domain similarity and transfer success is predictor-class dependent.

That reading is one of at least two, and we cannot separate them here. Their response variable is human-driven **ignition**; ours is burned **area**. Ignition likelihood is dominated by human access and activity, stable and near-universal drivers, whereas burned extent is dominated by spread. A predictor-class explanation and a response-variable explanation are therefore confounded in this comparison. Our own data speak against a simple predictor-class reading in any case: the static baseline transfers at a mean of 0.537 here, so within this cohort the place-attribute class does not travel either.

### 5.8. Implications

For practice, transfer skill has to be *measured* rather than inferred from  
975 similarity, and the pre-deployment diagnostics currently used for that inference did not order it here. Where a model must be moved, the resource that closes the gap is target labels: thirty-two labelled 5 km blocks recover 85 to 89 % of the target’s own ceiling in three of six directions, and 30 to 57 % in the rest. That is a real answer and not a cheap one, being 7 to 20 % of  
980 the target’s natural-vegetation population, and at small budgets the same intervention damages the direction that already transfers best.

For method development, the results point away from better unsupervised alignment. Two label-free methods applied carefully moved fourteen of twenty directions towards chance, and every direction they improved involves  
985 the smallest region. A sign reversal is not a distribution mismatch, and no realignment of inputs can repair it.

### 5.9. Limitations

The transfer failure is a finding, not a limitation. The limitations are the boundaries on how far it generalises.

990 (i) **No meteorological covariates** enter the models, so we cannot say how local skill and portability behave for a mixed thermal-plus-weather predictor set.

(ii) **Temporal transfer is measured for one region only**, and even there year and seasonal phase are confounded, a confound that cannot be  
995 resolved in this study area for the reason given in Section 5.2. Its 331 burned cells also leave the thermal reversals unresolved at interval level, and the

pair holds place fixed but not population. Those two arms are additionally the only transfer directions here computed by us rather than read from the pipeline author’s frozen export, with his unmodified code and the same  
1000 pinned environment.

(iii) **All labels derive from a single burned-area product**, MCD64A1 [9], whose omission and commission characteristics [3] bound every model evaluated here. No second burned-area product covers 2021 and 2022 at this resolution; the companion paper reports what an independent active-fire ob-  
1005 servation says about the omission concern.

(iv) **Evia remains the most imbalance-atypical population** even after the AOI extension, at a TSG prevalence of 0.287 against 0.038 to 0.072 in three of the others.

(v) **Each region contributes one fire season**, so regional concept  
1010 shift is confounded with event meteorology, and distinguishing them requires multi-year labels.

(vi) **Cross-region point estimates carry an implementation tolerance** of roughly  $\pm 0.02$  to 0.03 across scikit-learn versions. All reported numbers are fixed to one verified version, but exact reproduction elsewhere  
1015 requires the archived environment.

(vii) **The diagnostic correlations rest on an effective sample of ten region pairs.** Both the successes and the failures of Section 4.4 should be read at that power.

(viii) **Manavgat’s atypical transfer behaviour remains unexplained.**  
1020 It is where the conditional diagnosis bites hardest and where feature removal recovers most. Two candidates have now been tested and neither survives.

The first was its meteorology: the predictor window was not extreme. The second was the quality-screening difference in its coarse thermal input. It was the more worrying of the two, because the change it induces correlates with elevation at +0.615. It does not propagate. The region’s entire downstream chain was rebuilt from a quality-screened MODIS input. That changes the downscaled surface on 22,304 of 24,150 cells, by up to 10.9 °C. No signed univariate association moves by more than +0.0003. Elevation’s stays at 0.374 [0.290, 0.472] in both arms, and the population is unchanged. The result is structural rather than fortunate. Elevation is a DEM variable the screening cannot touch, and fusion falls back on the MODIS-derived surface across only 2.14 percentage points of coverage. With one fire season per region the remaining candidates are not separable in this design.

(ix) **The interval-support counts are less stable than the point estimates behind them.** Several verdicts sit within a thousandth of their reference value, and at 1 km blocking the published split of ten positive, seven negative and three uncertain turns on a lower bound of  $-0.00045$ . The point estimates and the sign pattern are stable; the counts are not. Every sentence in this paper that leans on an exact count of supported directions should be read at that precision.

## 6. Conclusions

Pre-fire thermal dryness adds a real and repeatable increment to burned-area discrimination. Across five Mediterranean regions, ROC-AUC was raised by +0.06 to +0.15 over a static and near-static terrain, fuel and greenness baseline. That baseline’s only time-varying member is the vegetation-index

composite. The gain survived spatial blocking at about 5 km, the coarsest scale this design supports as an interval. It also survived a predictor window closed up to two weeks before the first labelled burning.

That skill is local. Paired per direction, the same predictor block adds  
1050 nothing distinguishable from zero to cross-region transfer. The mean is  
+0.004 over twenty directions and its interval spans zero under all four  
between-direction resampling units we computed. Its contribution changes  
sign from one direction to another. The paired deltas run from  $-0.148$  to  
+0.132, twelve positive and eight negative. At the conservative 5 km block-  
1055 ing, five to six directions are helped and three to four are harmed with interval  
support, with the rest carrying no verdict. The block is also the swing factor  
at the chance line.

Two controls fix the meaning of all this. The baseline arm transfers at a  
mean of 0.537 against the thermal model's 0.541, so the static predictor class  
1060 is not portable here either. And a within-region half-split, applied with no  
refit through the same code path, returns 0.574, with skill near chance from  
about 10 km outward. Most of the loss is incurred before any boundary is  
crossed, and no further loss attributable to region crossing is detectable on  
top of it. The honest summary is therefore not that these predictors fail to  
1065 cross regions, but that a model of this kind does not generalise far beyond  
its training footprint, and that a region boundary adds nothing measurable  
to a loss that has already happened. That is a null about region crossing,  
not a demonstration that distance is the cause, and it leaves the variation  
between directions unexplained: separation does not order the matrix, and  
1070 six directions are anti-predictive with interval support, which extrapolation

alone does not produce.

Removing the two reversing predictors, elevation and the LST anomaly, costs  $-0.081$  of mean within-region skill, with interval support in every region and roughly three quarters of it attributable to elevation, a baseline terrain  
1075 variable; it changes mean transfer by  $+0.014$ , an estimate whose interval spans zero. A local cost is measured. No compensating transfer gain is, on either arm, so no exchange between the two is demonstrated.

The relationship itself is also unstable, and that is a separate finding. The direction of the link between dryness and burning changes from one region  
1080 to another, with bootstrap support for two predictors. This is measured on univariate associations and does not depend on any fitted model, so the distance result above neither establishes it nor removes it. What the two together rule out is the comfortable reading in which regional concept shift explains the whole transfer matrix. Most of its *level* is reached without  
1085 leaving a region. What varies around that level, including six anti-predictive directions, is not explained by separation and is where the instability matters.

None of the similarity diagnostics tested here ordered the transfer matrix: predictor-space distance, domain separability, niche overlap and regime structure all failed. On ten effective pairs those null results mean not shown to  
1090 order transfer, rather than shown not to. Label-blind adaptation compresses transfer towards chance instead of repairing it, and pooled multi-region training does not escape it.

The practical implication is a change in what is checked before a dynamic-state fire model is transferred. The question is not whether the target region  
1095 lies inside the source’s environmental envelope. In our matrix, the pair with

the highest envelope overlap failed in both directions, and the pair with the lowest overlap transferred in both. Both statements are made at the point estimate. The question is whether the signed feature-response directions agree.

1100 That check requires a labelled probe in the target region, and so, on inspection, do the niche and regime diagnostics. Only the marginal family can be computed before any target label exists, and it is the family that failed. How large the probe has to be is not established here, so no label budget is claimed. A signed association with a usable interval may need  
1105 fewer labels than a refitted model, but that was not tested. What is clear is that the supervised recalibration that label-free alignment cannot deliver is not cheap. Section 4.9 prices it for three regions: thirty-two labelled 5 km blocks recovered 85 to 89 % of the target ceiling in three of six directions, 51 to 57 % in two more and 30 % in the sixth, and at its top budget the labelled  
1110 set already holds most of one target region’s burned cells. Where no target labels exist, the transfer performance of such models should be treated as unknown rather than inferred from similarity.

For model builders, local skill and portability should be reported separately rather than assumed to travel together. A predictor block worth a  
1115 large within-region increment may contribute nothing distinguishable from zero across regions, with a sign that varies by pair, and reporting only within-region validation hides that completely. Three extensions are left for future work: temporal transfer, meteorological covariates, and physically normalised dryness variables. Each is a step towards the self-calibrating thermal  
1120 monitoring system that motivated this study.

## Appendix A. Appendix A. Sensitivity analyses

Each arm below varies one design choice and leaves everything else fixed. None changes a conclusion in the main text. They are reported so that a reader can see which choices were tested, and what each was worth. Arm (f) is the exception in one respect: it is not a robustness check but a test of a claim the introduction makes, and the claim is not upheld.

**(a) Evia AOI and prevalence.** The raw transfer arms were repeated with the legacy, high-prevalence Evia box. Every qualitative conclusion is unchanged. Thermal raw transfer AUCs move by up to 0.07, the largest being Evia to Bejís at 0.378 to 0.448, and no direction changes side of the chance line.

**(b) CORAL regularisation.** Over nine  $\lambda$  values from 0 to  $10^{-1}$  on four directions, CORAL transfer AUC moves by at most 0.014 within any direction and 0.008 within the thermal family, so no CORAL-dependent conclusion here is sensitive to  $\lambda$  in that range. The companion paper reports the canonical  $\lambda = 1$ , which the released sweep omits.

**(c) Blocking scale.** Recomputing the transfer quantities at 10-cell ( $\approx 5$  km) blocking from the frozen per-cell predictions widens the intervals and moves the verdict counts, from ten positive, seven negative and three uncertain at 1 km to six, four and ten at 5 km. Coarser blocking therefore removes support from seven verdicts and adds none. Across five seeds the counts run 5 to 6 positive, 3 to 4 negative and 10 to 11 uncertain, so they are not exact.

The point estimates are unchanged. That is an identity rather than a result, and it should not be offered as robustness. The blocking scale is the bootstrap *resampling unit*, and each point estimate is computed once

over all target cells, so no choice of block size could have moved one. The comparison does establish two things. Every verdict that changes, changes towards "no verdict". And no direction crosses the chance line under the widened intervals. The counts are the fragile part of this paper. The sign  
1150 pattern is not thereby shown to be robust. It is simply not tested by this variation.

**(d) Predictor-window closure.** The predictor window was closed 7 and 14 days earlier in all five regions. The thermal contribution stays positive and bootstrap-supported everywhere. The direction of the change is region-  
1155 specific. It strengthens in Bejís, from 0.058 to 0.079 at 14 days, and in Muğla, from 0.115 to 0.128. It is flat in Montiferru. It weakens monotonically in Evia, from 0.156 to 0.149 to 0.135. What holds in every region is survival, not improvement. That is the claim made in Section 5.2.

**(e) Quality screening of the coarse thermal input.** Two of the  
1160 five regions' MODIS inputs are quality-screened and three are not. The split follows export date rather than design, and it induces an elevation-correlated change at the input, at  $r = +0.615$  in Manavgat. Manavgat's entire downstream chain was rebuilt from a quality-screened input. That changes the downscaled surface on 22,304 of 24,150 cells, by up to 10.9 °C. No  
1165 signed univariate association moves by more than +0.0003. Elevation stays at 0.374 [0.290, 0.472] in both arms, the population is unchanged, and the within-region increment moves from [+0.055, +0.079] to [+0.054, +0.077]. The reason is structural. Elevation is a DEM variable the screening cannot touch, and fusion falls back on the MODIS-derived surface across only 2.14  
1170 percentage points of coverage. Details are in `paper/modis_qc_downstream_`

propagation.md.

(f) **Normalised against absolute dryness channels.** Section 1.2 argues that an internally normalised index should be less exposed to absolute-temperature offsets between regions than raw land surface temperature. The thermal block contains both kinds, so the argument can be tested directly as a feature-set contrast.

One thing should be said before the result, because it makes the outcome less surprising than it might otherwise appear. The normalised channels are the two that hold their direction in the same-geography two-event comparison of Section 4.8, where the absolute channels move. That is a statement about one region across two fires. **Across regions it does not hold:** `1st_anomaly_mean` is itself one of the two bootstrap-supported reversing predictors, reversing between Bejís and Evia, which is why Section 4.6b drops it alongside elevation. Being internally normalised protects a channel against the offset between two seasons in one place. It does not, on this evidence, protect it against a change of place.

Three feature sets were run over all twenty directions and all five within-region folds, with the classifier, population, folds and bootstrap held fixed. The harness aborts unless its reference configuration lands on the frozen exports. It did exactly: the maximum absolute difference between the reference arm and the frozen transfer AUCs is 0.000000 across all twenty directions.

**The advantage is not found.** Per direction, the normalised set minus the absolute set has a mean of  $-0.0037$ . It is positive in 12 of 20 directions and runs from  $-0.099$  to  $+0.094$ , so the spread is an order of magnitude larger than the difference. The normalised channels transfer very slightly worse

Table A.10:

| Feature set                             | Mean transfer AUC | Directions > 0.5 | Mean<br>within-region<br>AUC |
|-----------------------------------------|-------------------|------------------|------------------------------|
| baseline only                           | 0.5371            | 16               | 0.7896                       |
| baseline +<br>normalised<br>anomalies   | 0.5442            | 15               | 0.8551                       |
| baseline +<br>absolute surface<br>state | 0.5479            | 13               | 0.8564                       |
| all ten features<br>(reference)         | 0.5414            | 14               | 0.8883                       |

at the mean, and nothing here supports a design rule favouring anomaly-referenced dryness for portability.

Within region the two sub-blocks are indistinguishable, at 0.8551 and 0.8564, and each recovers most of the gap between the baseline’s 0.7896 and the full model’s 0.8883. What separates them in sign stability does not become transferable skill.

This arm strengthens the negative finding rather than softening it. Four differently constituted feature sets were tried. None clears a mean of 0.548 across regions, and all differ sharply within them. The baseline arm also gives an independent confirmation of the control reported in Section 4.3: its mean transfer of 0.5371 was recomputed here from the modelling datasets and matches the 0.537 read from the frozen per-direction export. Details are

in paper/anomaly\_only\_transfer.md.

The comparison is between two sub-blocks of one thermal set on one  
1210 cohort. It does not test normalised dryness indices in general, and it does  
not test a normalisation fitted against a pooled multi-region reference rather  
than each region's own baseline years.

## Appendix B. Appendix B. All twenty transferability diagnostics

Section 4.4 summarises this table by family. The full ranking is given  
1215 here because the negative result is the point: a reader should be able to see  
every candidate that was tried, not only the families' best members, and to  
check that no diagnostic was dropped after it failed.

**Table B1. All transferability diagnostics versus raw thermal  
transfer (20 ordered directions).** Spearman  $\rho$  with pair-based bootstrap  
1220 95 % CIs. Exp. = expected sign. Rows with  $n = 12$  exist only for the four-  
AOI subset, because those diagnostics were never produced for Montiferru.  
The supported-features conditional rows use the 16 directions (8 pairs) with  
at least one CI-supported feature.

*Table note.* A null row means the diagnostic was **not shown to order**  
1225 **transfer** on this design, not that it was shown incapable of ordering it.  
With ten effective region pairs the power is low, and the intervals are wide  
enough to admit moderate true correlations in either direction. The last row  
is retained rather than deleted because it was computed: on two directions  
the statistic has no meaningful value, and reporting that is more honest than  
1230 dropping the variant.

**Reading the two rows that clear zero.** Both are supported-feature

Table B.11:

| Diagnostic                                    | Family                                 | Exp. | n dir | Spearman $\rho$<br>[95% CI] | CI excludes 0 |
|-----------------------------------------------|----------------------------------------|------|-------|-----------------------------|---------------|
| <b>Agreement fraction, supported features</b> | <b><math>P(y x)</math> conditional</b> | +    | 16    | <b>+0.84</b> [+0.58, +0.88] | <b>yes</b>    |
| <b>Cosine, supported features</b>             | <b><math>P(y x)</math> conditional</b> | +    | 16    | <b>+0.81</b> [+0.33, +0.88] | <b>yes</b>    |
| Cosine, all 9 features                        | $P(y x)$ conditional                   | +    | 20    | +0.50 [−0.17, +0.83]        | no            |
| Vector Spearman, all 9                        | $P(y x)$ conditional                   | +    | 20    | +0.27 [−0.36, +0.77]        | no            |
| Agreement count, all 9                        | $P(y x)$ conditional                   | +    | 20    | +0.18 [−0.40, +0.72]        | no            |
| Schoener’s D, 1-D mean                        | $P(x y=1)$ niche                       | +    | 20    | +0.24 [−0.45, +0.74]        | no            |
| Warren’s I, 1-D mean                          | $P(x y=1)$ niche                       | +    | 20    | +0.22 [−0.42, +0.73]        | no            |
| Schoener’s D, PCA-2D                          | $P(x y=1)$ niche                       | +    | 20    | +0.10 [−0.51, +0.68]        | no            |
| Warren’s I, PCA-2D                            | $P(x y=1)$ niche                       | +    | 20    | −0.07 [−0.66, +0.49]        | no            |
| Mahalanobis, burned centroids                 | $P(x y=1)$ niche                       | −    | 20    | −0.23 [−0.75, +0.44]        | no            |
| Domain-classifier AUC                         | $P(x)$ marginal                        | −    | 20    | −0.32 [−0.78, +0.33]        | no            |
| Predictor-space mean dissimilarity            | $P(x)$ marginal                        | −    | 12    | −0.10 [−0.54, +0.43]        | no            |
| Predictor-space p95 dissimilarity             | $P(x)$ marginal                        | −    | 12    | −0.08 [−0.59, +0.49]        | no            |
| Fraction inside weighted AoA                  | $P(x)$ marginal                        | +    | 12    | +0.22 [−0.48, +0.59]        | no            |
| Fraction inside                               | $P(x)$ marginal                        | +    | 12    | +0.08 [−0.89, +0.43]        | no            |

variants, where the predictor subset is chosen by whether two regions' bootstrap intervals happen to be disjoint. That is a data-dependent selection made on the same data, uncorrected. Their all-nine-feature counterparts are  
1235 in the table and both span zero. The result lives in the selection step, and Section 4.4 says so.

**Equal-sample check.** The families sit on unequal samples: marginal, applicability, climatic and geographic rows on twelve directions, the supported-conditional rows on sixteen, the rest on twenty. Recomputing every row on  
1240 the common twelve directions reproduces the published values to  $4.8 \times 10^{-5}$  and leaves the ordering unchanged — the conditional rows still lead at +0.87 [+0.65, +0.88] and +0.85 [+0.43, +0.88], every marginal row still spans zero. Source: `paper/diagnostics_common_subset.md`.

*Appendix B.1. B2. The signed univariate associations the reversal claim*  
1245 *rests on*

Section 4.6b drops two predictors because their signed association with burning reverses between regions with bootstrap support, and Contribution 3 names the reversal as the mechanism. This is the evidence for both. It was computed for the frozen analysis and is reproduced here because the claim  
1250 is otherwise asserted rather than shown.

**Table B2. Signed univariate AUC of each predictor against burned, by region.** Primary natural-vegetation population; 10-cell ( $\sim 5$  km) spatial-block bootstrap, 1000 replicates, seed 42. The AUC is never folded to  $\max(\text{AUC}, 1 - \text{AUC})$ , so a value below 0.5 means lower values  
1255 rank burned and is a direction rather than weakness. **Bold** marks a region whose own interval excludes 0.5.

Table B.12:

| Feature             | Manavgat             | Bejís                | Muğla                | Evia                 | Montiferru           |
|---------------------|----------------------|----------------------|----------------------|----------------------|----------------------|
| <b>elevation_</b>   | <b>0.374</b> [0.289, | <b>0.643</b> [0.558, | <b>0.611</b> [0.532, | 0.541 [0.448,        | 0.584 [0.395,        |
| <b>mean</b>         | 0.471]               | 0.729]               | 0.690]               | 0.626]               | 0.762]               |
| <b>slope_mean</b>   | 0.531 [0.423,        | 0.521 [0.439,        | <b>0.637</b> [0.582, | 0.487 [0.418,        | <b>0.652</b> [0.506, |
|                     | 0.642]               | 0.605]               | 0.686]               | 0.554]               | 0.771]               |
| <b>ndvi_mean</b>    | <b>0.636</b> [0.587, | 0.559 [0.497,        | <b>0.662</b> [0.616, | <b>0.639</b> [0.575, | 0.586 [0.450,        |
|                     | 0.676]               | 0.619]               | 0.704]               | 0.701]               | 0.704]               |
| <b>lst_anomaly_</b> | 0.482 [0.428,        | <b>0.418</b> [0.364, | 0.485 [0.395,        | <b>0.640</b> [0.567, | 0.395 [0.285,        |
| <b>mean</b>         | 0.530]               | 0.480]               | 0.566]               | 0.710]               | 0.535]               |
| <b>current_lst_</b> | 0.538 [0.452,        | 0.477 [0.401,        | <b>0.325</b> [0.271, | <b>0.377</b> [0.301, | 0.370 [0.248,        |
| <b>mean</b>         | 0.621]               | 0.547]               | 0.382]               | 0.456]               | 0.513]               |
| <b>current_</b>     | 0.552 [0.460,        | 0.517 [0.429,        | <b>0.336</b> [0.275, | <b>0.362</b> [0.285, | <b>0.356</b> [0.233, |
| <b>tvdi_mean</b>    | 0.641]               | 0.595]               | 0.398]               | 0.442]               | 0.499]               |
| <b>tvdi_</b>        | 0.449 [0.391,        | 0.512 [0.443,        | 0.490 [0.396,        | 0.519 [0.444,        | <b>0.378</b> [0.282, |
| <b>difference_</b>  | 0.505]               | 0.583]               | 0.575]               | 0.589]               | 0.497]               |
| <b>mean</b>         |                      |                      |                      |                      |                      |
| <b>downscaled_</b>  | 0.552 [0.466,        | 0.484 [0.400,        | <b>0.307</b> [0.253, | <b>0.377</b> [0.297, | 0.365 [0.240,        |
| <b>lst_mean</b>     | 0.637]               | 0.560]               | 0.366]               | 0.459]               | 0.511]               |
| <b>fused_lst_</b>   | 0.540 [0.454,        | 0.481 [0.404,        | <b>0.325</b> [0.272, | <b>0.376</b> [0.300, | 0.370 [0.248,        |
| <b>mean</b>         | 0.622]               | 0.551]               | 0.383]               | 0.456]               | 0.513]               |

**What counts as a reversal.** A pair of regions is called a reversal only when their signed associations point to opposite sides of 0.5 **and each region’s own interval excludes 0.5**. That is stricter than requiring the two regions’ intervals to be disjoint, and the difference matters: for `current_1st_mean` between Manavgat and Muğla the two intervals are disjoint, but Manavgat’s own interval, [0.452, 0.621], includes 0.5, so Manavgat has no established direction to reverse from. That pair is a point reversal, not a supported one.

**Table B3. The cross-region reversals that meet the stricter criterion.** Difference intervals are from the same paired bootstrap.

| Table B.13:                   |          |       |          |       |                         |
|-------------------------------|----------|-------|----------|-------|-------------------------|
| Feature                       | Region A | AUC   | Region B | AUC   | Difference [95 % CI]    |
| <code>elevation_mean</code>   | Manavgat | 0.374 | Bejís    | 0.643 | +0.269 [+0.138, +0.397] |
| <code>elevation_mean</code>   | Manavgat | 0.374 | Muğla    | 0.611 | +0.235 [+0.102, +0.360] |
| <code>1st_anomaly_mean</code> | Bejís    | 0.418 | Evia     | 0.640 | +0.221 [+0.123, +0.313] |

Three pair-level reversals across **two** features, which is why Section 3.12 removes exactly those two. Twenty-nine further pairs reverse at the point estimate only, spread across eight of the nine features, and they are not counted. The conservative criterion costs the paper findings rather than manufacturing them: a difference interval on the pair, which is the instrument Section 4.8 uses, would support more reversals than the three listed here.

## Appendix C. Appendix C. Protocol detail

These protocols are given here in full rather than in Methods, because  
1275 each is a specification a reader needs only when checking the corresponding  
result, and none is needed to follow the argument. Sections 3.2, 3.3, 3.11 and  
3.14 state in summary what each does.

### *Appendix C.1. C.1 Burned-area label, the reconstructed analysis grid, and the admissibility gate*

1280 Labels come from MODIS MCD64A1 Collection 6.1 (MODIS/061/MCD64A1)  
retrieved through Google Earth Engine [10], whose omission and commission  
characteristics [3] bound every model here. The product is exported onto  
the 30 m EPSG:4326 reference grid used by the rest of the pipeline, which  
duplicates each native  $\sim 500$  m observation across the 30 m pixels beneath  
1285 it.

The analysis grid is therefore **reconstructed rather than native**: the  
30 m grid is partitioned into non-overlapping blocks of  $\text{round}(500 / 30) =$   
 $17 \times 17$  pixels, giving a nominal cell edge of 510 m, and each cell is identified  
by its integer block indices `row_500m` and `col_500m`. This approximates the  
1290 MODIS sinusoidal cell rather than reproducing it, being anchored to the  
Landsat/EPSSG:4326 reference grid, and marginal blocks are truncated. The  
cell is square in degrees but not on the ground, about 510 m north to south  
and 390 to 407 m east to west, so its area is 0.199 to 0.208 km<sup>2</sup> against a  
MODIS cell's 0.250 km<sup>2</sup>; every block-size label in this paper is therefore the  
1295 north-south dimension. The companion paper develops what follows from  
that.

A cell's representative burn date is the **mode** of the positive sub-pixel day-of-year values in the block, tested against the region's label window. Because the exported raster carries no out-of-window positives, in this dataset  
1300 a single in-window positive sub-pixel makes the cell burned. The fraction of positive sub-pixels agreeing with the modal date is recorded as **burn\_date\_pixel\_agreement\_fraction** but no agreement threshold is imposed. The label never affects a cell's eligibility for modelling: unburned, all-no-data and out-of-window cells all remain as the negative class.

1305 Two safeguards are recorded rather than assumed. Cells that burned inside a region's own predictor window are removed from its analysis universe rather than counted as unburned, which ran for Muğla, Evia and Montiferru, is recorded as not run for Manavgat and has no recorded status for Bejís. Burning in earlier years is screened for no region in the five-region cohort;  
1310 the only historical exclusion in the study removes the 2021 Muğla scar from the 2022 event-relative experiment of Section 3.14.

Before any modelling each region passes a gate that answers one question: of the cells labelled burned, what fraction is dominated by natural vegetation? A region is admitted as a wildfire candidate when that fraction  
1315 reaches 0.50 and at least 30 burned cells are present, and is rejected as a cropland-dominated control when the cropland fraction reaches 0.50 instead. The gate uses ESA WorldCover classes aggregated to the same reconstructed cells. Its purpose is to separate burned area produced by natural-fuel combustion from burned area produced by post-harvest stubble burning, which  
1320 MCD64A1 does not distinguish. Verdicts are reported in Section 4.1.

## Appendix C.2. C.2 Transferability diagnostics versus transfer

Twenty candidate diagnostics from four families are each rank-correlated against the same target quantity, the raw thermal transfer AUC over the twenty ordered directions, under one common pair-based bootstrap. The families are marginal predictor-distribution measures  $P(x)$ , burned-niche overlap  $P(x|y=1)$ , fire-regime label-pattern structure  $P(y)$ , and conditional feature-response direction  $P(y|x)$ .

The **marginal** family includes area-of-applicability-style dissimilarity in scaled, importance-weighted predictor space [21, 22, 16], climatic and geographic distance, and a learned domain classifier trained to separate source from target cells. The **niche-overlap** family includes Schoener’s D [34] and Warren’s I [41] over the burned cells of each region. The **regime** family includes burned-patch component counts and effective component counts. The **conditional** family is the sign-agreement index: the fraction of predictors whose signed univariate association points the same way in both regions, computed over all predictors and over the subset whose associations are individually interval-supported.

Two properties of this design are stated in advance because they bound what it can show. The candidate set was fixed in a project analysis log before the correlations were computed; that log is not a formal pre-registration and no independent timestamped registration exists. And with five regions the effective sample is ten unordered pairs, so nineteen computed variants are evaluated at that power with no family-wise error control claimed; where a diagnostic clears an interval test but not a Bonferroni threshold, both are reported.

*Appendix C.3. C.3 Same-geography event-to-event comparison (Mugla 2021 versus 2022)*

Muğla is the one region for which a second fire event is analysed on the identical AOI and analysis grid, which allows the direction of feature-label associations to be compared with geography held fixed. The 2022 experiment is registered with windows anchored to its own event: a 58-day predictor window closing the day before ignition and a 49-day label window opening on it, matching the 2021 durations exactly.

**This is not a clean temporal-transfer design and is not presented as one.** The 2022 event ignites about six weeks earlier in the season than the 2021 event, so calendar year and seasonal phase are confounded and no difference can be attributed to the year alone. The registered claim is same-geography event-to-event. The pair is deliberately kept out of the twenty-direction matrix, and that exclusion is enforced in released code rather than merely asserted here.

Two further properties bound the comparison. A historical-burn exclusion applies to the 2022 arm alone, masking every cell that burned in 2021 so the previous year’s scar does not enter the following year’s analysis: it removes 3,073 cells, of which 2,941 belong to the primary population, taking it from 41,730 rows to 38,790, a drop of 7.0 %. The two arms therefore share 38,789 of 38,790 cells and three byte-identical static predictors, and in the 2022-to-2021 direction every target positive lies outside the source training population. Section 4.8 states what follows for how these two numbers may be read.

## CRediT authorship contribution statement

1370     **Yunus Emre Cogurcu:** Conceptualization, Methodology, Formal analysis, Investigation, Writing – original draft, Writing – review and editing, Supervision. **Emrehan Metin:** Software, Data curation, Investigation, Validation, Writing – review and editing.

## Declaration of competing interest

1375     The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

## Funding

1380     This work is an output of project 17506, supported by the Scientific Research Projects (BAP) unit of Çukurova University.

## Data availability

1385     All satellite inputs are public and are obtained through Google Earth Engine. The processing pipeline, its configuration and the frozen numeric outputs are publicly available; the repository, the commit of record and the licence are given in the data and code availability statement of the Methods, together with the three components that fall outside that release. No digital object identifier is minted and no archival deposit exists.





Figure C.1: **Study regions and the gate negative control.** Locations of the five wild-fire regions analysed in this study (blue, circular markers): Manavgat 2021 and Muğla 2021 (Türkiye), Bejís 2022 (Spain), North Evia 2021 (Greece, extended AOI) and Montiferru 2021 (Sardinia, Italy). Rectangles are the true analysis areas of interest; the ring marks each centroid, because at basin scale the rectangles alone are hard to locate. Kozan 2023 (grey, dashed outline, square marker) is *not* a study region: it is the negative control used to calibrate the burned-landcover admissibility gate, and it enters no model, transfer or diagnostic anywhere in this paper. Its rectangle is the true extent of its processed AOI and is larger than the study AOIs for that reason alone; the difference in size carries no methodological meaning and is not a design choice. Coastlines are Natural Earth 1:50 m; country borders are omitted deliberately. The scale bar is exact only at 38°N. Region bounding boxes match Table 1 to  $5 \times 10^{-3}$  degrees.



**Figure C.2: Processing pipeline and evaluation programme.** Flow from source products (left) to the transfer and diagnostic programme (right); every box carries the Methods subsection that specifies it. The five source products (MCD64A1 burned area, Landsat Collection 2 Level-2, MOD11A1, Copernicus DEM GLO-30 and ESA WorldCover) merge onto the reconstructed  $\sim 510$  m analysis grid ( $17 \times 17$  pixels of the 30 m reference grid). The burned-landcover admissibility gate has two outcomes: regions that *pass* proceed to the natural-vegetation primary population, while the rejecting branch (grey, dashed) terminates in the Kozan 2023 negative control, whose burned area is 98% cropland and which is never modelled. The dashed connector deliberately carries no arrowhead onward, so the control cannot be read as entering the pipeline. Feature sets feed two evaluations in parallel: within-region (random forest under spatial-block cross-validation) and cross-region transfer (source-only, 20 ordered directions, with label-blind z-score and CORAL adaptation). Their levels are combined in the transfer-gap decomposition (§3.12). The footer states what applies throughout: leakage hard-exclusion, seed 42, 1000-replicate spatial-block bootstrap (§3.13), and the sensitivity designs (§3.16). Section references are verified against the Methods heading list when the figure is built.




Figure C.3: **The thermal increment replicates in all five regions and survives coarser spatial blocking.** Natural-vegetation primary population; random forest with the primary weighted configuration; all intervals are 95% spatial-block bootstrap intervals (1000 replicates, seed 42). **(a)** Baseline versus thermal out-of-fold ROC-AUC per region at the default 2-cell ( $\sim 1$  km) blocking: the open marker is the static baseline feature set, the filled marker is the baseline plus the six thermal predictors, and the connecting bar is the increment. The right-hand column reports the paired  $\Delta AUC$  and its 95% interval; it is a fixed column placed beyond the plotted AUC range, so no label position depends on a data value. Every interval excludes zero. **(b)** Absolute thermal ROC-AUC against spatial block size (2, 10 and 20 cells,  $\approx 1, 5$  and 10 km): absolute skill falls as blocking coarsens, in every region. **(c)** The thermal increment  $\Delta AUC$  against the same block sizes, with 95% intervals: the increment does *not* melt away, and its interval excludes zero in all ten of the 2- and 10-cell region-block-size combinations. The five 20-cell intervals also lie above zero but rest on 6 to 33 positive-carrying blocks each, so they are plotted as indicative rather than as intervals (Table 2, table note). In (c) the five regions are offset horizontally by 0.15 cell so their whiskers separate; **this offset is cosmetic, since all five regions are evaluated at exactly the same three block sizes** (2, 10, 20 cells), and the offset is smaller than half the spacing between block sizes, so the groups cannot merge. Colour and marker shape both encode region, so the panels remain readable in greyscale.




Figure C.4: **Cross-region transfer, and what label-blind adaptation does to it.** Thermal transfer ROC-AUC for all 20 ordered source–target directions among the five regions, natural-vegetation primary population. **(a)** Raw transfer, no adaptation. **(b)** After region-wise z-score standardisation. **(c)** After CORAL. Rows are the source region, columns the target; the grey diagonal is within-region performance and is *not* a transfer result. The diverging colour scale is centred exactly on chance (0.5, marked by the black rule on the colour bar), so orange is below chance and purple above. Each panel header states that panel’s range: adaptation compresses the spread from 0.326–0.686 (raw) to 0.431 to 0.630 (z-score) and 0.443 to 0.624 (CORAL). It pulls the extremes toward chance rather than lifting transfer. Cells below chance are hatched. The hatch is a redundant cue for print: a diverging palette is symmetric in luminance about its centre, so in greyscale a below-chance cell would otherwise be indistinguishable from an equally distant above-chance one. Hatching follows the *exact* value while the printed label is rounded to two decimals, so two cells print “0.50” with different hatching: CORAL Bejís→Evia is 0.4991 (hatched, below chance) and z-score Evia→Muğla is 0.5013 (un-hatched, above chance). Every cell carries its value, and every label’s colour is chosen by measured WCAG contrast against that cell’s rendered colour (worst case 4.7:1). Confidence intervals for these directions are given in Table 4.




Figure C.5: **Label-blind adaptation compresses transfer toward chance rather than recovering it.** All 12 ordered directions among the four regions carrying the full decomposition (Manavgat, Bejís, Muğla, Evia), natural-vegetation primary population. Each row is one direction. The open circle is raw transfer ROC-AUC and the arrow runs to the best label-blind adapted result for that direction (region-wise z-score or CORAL, whichever is better); the grey vertical tick is the within-region reference for that target, and the dashed vertical rule is chance. **Solid blue arrows are the five directions where adaptation moves transfer toward the within-region reference (recovery); dashed orange arrows are the seven where it moves *away* from it (negative recovery).** Line style and colour both encode this class, because the two hues differ by only 2.30:1 in relative luminance and would nearly merge in greyscale. The pattern is the point: all six directions starting above chance are pushed down, and five of the six starting below chance are pushed up, so adaptation acts on the spread rather than on skill. The exception is Manavgat→Muğla, which starts below chance at 0.470 and is pushed further down to 0.443. The worst case is Evia→Manavgat, where raw transfer of 0.686 is dragged down to 0.527 against a within-region reference of 0.870, a recovered fraction of  $-0.862$ . No arrow reaches its within-region tick.




Figure C.6: **Pooling every other region never beats the best single source.** Leave-one-region-out evaluation: for each held-out target region (rows), a model trained on the pooled remaining four regions (filled orange circle) is compared against the best individual source region for that same target (open blue diamond), with the within-region ceiling as a grey vertical tick and chance as the dashed rule. Natural-vegetation primary population, thermal feature set, no adaptation. The pooled model lies to the left of the best pairwise source in all five targets without exception. Adding source regions does not accumulate transferable signal. In two targets the pooled model falls *below chance* at the point estimate (Manavgat 0.469, Bejís 0.417), and only for Bejís does the interval lie entirely below chance, while a single well-chosen source reaches 0.686 and 0.583 respectively. Every point remains far below the within-region ceiling.




Figure C.7: **Removing the direction-reversing features is a zero-sum trade-off.**

Mean within-region ROC-AUC across the five regions (horizontal) against mean transfer ROC-AUC across the 20 ordered directions (vertical), for four feature configurations: the full thermal set, the set without the LST anomaly, without elevation, and without both. Each configuration has its own marker and is named in the key; the grey path connects them in that order. Moving left along the path, transfer rises monotonically (0.541  $\rightarrow$  0.556) while within-region skill falls monotonically (0.888  $\rightarrow$  0.807). Both monotonicities are verified when the figure is built, so the figure cannot silently contradict the claim. The transfer gain is 0.014 AUC and the within-region cost 0.081 AUC, about six times larger: what is removed is locally informative and non-portable, and removing it buys portability by giving up skill rather than by finding transferable signal.




Figure C.8: **Burned-niche overlap does not determine transfer: the contrast pair.** Per-feature signed univariate AUC, the probability that a predictor takes a higher value on a burned cell than on an unburned one, for the two most informative region pairs, natural-vegetation primary population. Arrows run from chance (0.5) to the observed value, so their direction is the sign of the association and their length its strength; whiskers are 95% spatial-block ( $\approx 5$  km) bootstrap intervals. Within each pair the two regions are distinguished by colour, by line style (solid for the first region named, dashed for the second) and by vertical offset; line style is included because the two hues differ by only 2.30:1 in relative luminance and would nearly merge in greyscale. **A filled arrowhead means that region's 95% interval excludes 0.5, so the direction is statistically supported; an open arrowhead means the interval contains 0.5 and the direction is not supported.** Rows shaded grey are features whose sign disagrees between the two regions. Grey side bars are the per-feature Schoener's  $D$  overlap. **(a)** Manavgat–Muğla, the pair with the *highest* burned-niche overlap ( $\bar{D} = 0.83$ ), transfers below chance in both directions (0.470 and 0.401). **(b)** Bejís–Montiferru, the pair with the *lowest* overlap ( $\bar{D} = 0.48$ ), transfers above chance in both directions (0.594 and 0.548). High niche overlap is therefore not sufficient for transfer. The arrowhead fill also exposes the coverage limit of the supported-agreement index, which is defined only on features whose direction is supported in *both* regions: there are two such features in (a) (NDVI and elevation) and **none in (b)**, where the two regions' supported sets are disjoint and Montiferru's intervals

## References

- [1] Alkan Akıncı, H., Akıncı, H., 2023. Machine learning based forest fire susceptibility assessment of Manavgat district (Antalya), Turkey. *Earth Science Informatics* 16, 397–414. doi:10.1007/s12145-023-00953-5.
- [2] Archibald, S., Lehmann, C.E.R., Gómez-Dans, J.L., Bradstock, R.A., 2013. Defining pyromes and global syndromes of fire regimes. *Proceedings of the National Academy of Sciences* 110, 6442–6447. doi:10.1073/pnas.1211466110.
- [3] Boschetti, L., Roy, D.P., Giglio, L., Huang, H., Zubkova, M., Humber, M.L., 2019. Global validation of the Collection 6 MODIS burned area product. *Remote Sensing of Environment* 235, 111490. doi:10.1016/j.rse.2019.111490.
- [4] Breiman, L., 2001. Random forests. *Machine Learning* 45, 5–32. doi:10.1023/A:1010933404324.
- [5] de Bruin, S., Brus, D.J., Heuvelink, G.B.M., van Ebbenhorst Tengbergen, T., Wadoux, A.M.J.C., 2022. Dealing with clustered samples for assessing map accuracy by cross-validation. *Ecological Informatics* 69, 101665. doi:10.1016/j.ecoinf.2022.101665.
- [6] Chuvieco, E., Cocero, D., Riaño, D., Martín, P., Martínez-Vega, J., de la Riva, J., Pérez, F., 2004. Combining NDVI and surface temperature for the estimation of live fuel moisture content in forest fire danger rating. *Remote Sensing of Environment* 92, 322–331. doi:10.1016/j.rse.2004.01.019.

- [7] Dimarco, N.A., Faraji, I., Wahbi, M., Maatouk, M., Boulaassal, H., Yazidi Aalaoui, O., El Kharki, O., 2026. Generalizing human-driven wildfire ignition models across Mediterranean regions using harmonized remote-sensing and machine-learning data. *Geomatics* 6, 13. doi:10.3390/geomatics6010013.
- [8] Gelabert, P.J., Jiménez-Ruano, A., Ochoa, C., Alcasena, F., Sjöström, J., Marrs, C., Ribeiro, L.M., Palaiologou, P., Bentu Martínez, C., Chuvieco, E., Vega-García, C., Rodrigues, M., 2025. Assessing human-caused wildfire ignition likelihood across Europe. *Natural Hazards and Earth System Sciences* 25, 4713–4729. doi:10.5194/nhess-25-4713-2025.
- [9] Giglio, L., Boschetti, L., Roy, D.P., Humber, M.L., Justice, C.O., 2018. The Collection 6 MODIS burned area mapping algorithm and product. *Remote Sensing of Environment* 217, 72–85. doi:10.1016/j.rse.2018.08.005.
- [10] Gorelick, N., Hancher, M., Dixon, M., Ilyushchenko, S., Thau, D., Moore, R., 2017. Google Earth Engine: Planetary-scale geospatial analysis for everyone. *Remote Sensing of Environment* 202, 18–27. doi:10.1016/j.rse.2017.06.031.
- [11] Huang, Y., Liu, N., Chen, W., Tan, W., Deng, Y., Wang, Z., Chlus, A., Shen, J., Townsend, P.A., 2026. Enhancing transferability of foliar trait retrieval models: A comparative analysis of transfer learning strategies and domain shift characterization. *Remote Sensing of Environment* 337, 115336. doi:10.1016/j.rse.2026.115336.

- 1435 [12] Iban, M.C., Sekertekin, A., 2022. Machine learning based wildfire susceptibility mapping using remotely sensed fire data and GIS: A case study of Adana and Mersin provinces, Turkey. *Ecological Informatics* 69, 101647. doi:10.1016/j.ecoinf.2022.101647.
- [13] Jodhani, K., Bhatiya, L., Sirvi, L., Chanda, A., Thakkar, S., Rathnayake, U., 2026. Artificial intelligence and remote sensing frameworks for wildfire monitoring and risk analysis across multiple ecosystems: A review. *AIMS Environmental Science* 13, 56–98. doi:10.3934/environsci.2026004.
- 1440 [14] Kondylatos, S., Prapas, I., Camps-Valls, G., Papoutsis, I., 2023. Mesogeos: A multi-purpose dataset for data-driven wildfire modeling in the Mediterranean, in: *Advances in Neural Information Processing Systems* 36 (NeurIPS 2023) Datasets and Benchmarks Track, pp. 50661–50676. doi:10.52202/075280–2204.
- 1445 [15] Liu, C., Mostafavi, A., 2025. WildfireGenome: Interpretable machine learning reveals local drivers of wildfire risk and their cross-county variation. *arXiv preprint arXiv:2511.11589*. doi:10.48550/arXiv.2511.11589, arXiv:2511.11589.
- [16] Ludwig, M., Moreno-Martinez, A., Hölzel, N., Pebesma, E., Meyer, H., 2023. Assessing and improving the transferability of current global spatial prediction models. *Global Ecology and Biogeography* 32, 356–368. doi:10.1111/geb.13635.
- 1455 [17] Maffei, C., Alfieri, S.M., Menenti, M., 2018. Relating spatiotemporal

- patterns of forest fires burned area and duration to diurnal land surface temperature anomalies. *Remote Sensing* 10, 1777. doi:10.3390/rs10111777.
- [18] Maffei, C., Lindenbergh, R., Menenti, M., 2021. Combining multispectral and thermal remote sensing to predict forest fire characteristics. *ISPRS Journal of Photogrammetry and Remote Sensing* 181, 400–412. doi:10.1016/j.isprsjprs.2021.09.016.
- [19] Maffei, C., Menenti, M., 2019. Predicting forest fires burned area and rate of spread from pre-fire multispectral satellite measurements. *ISPRS Journal of Photogrammetry and Remote Sensing* 158, 263–278. doi:10.1016/j.isprsjprs.2019.10.013.
- [20] Marino, E., Yáñez, L., Guijarro, M., Madrigal, J., Senra, F., Rodríguez, S., Tomé, J.L., 2024. Transferability of empirical models derived from satellite imagery for live fuel moisture content estimation and fire risk prediction. *Fire* 7, 276. doi:10.3390/fire7080276.
- [21] Meyer, H., Pebesma, E., 2021. Predicting into unknown space? Estimating the area of applicability of spatial prediction models. *Methods in Ecology and Evolution* 12, 1620–1633. doi:10.1111/2041-210X.13650.
- [22] Meyer, H., Pebesma, E., 2022. Machine learning-based global maps of ecological variables and the challenge of assessing them. *Nature Communications* 13, 2208. doi:10.1038/s41467-022-29838-9.
- [23] Meyer, H., Reudenbach, C., Hengl, T., Katurji, M., Nauss, T., 2018. Improving performance of spatio-temporal machine learning models using

forward feature selection and target-oriented validation. *Environmental Modelling & Software* 101, 1–9. doi:10.1016/j.envsoft.2017.12.001.

- [24] Milà, C., Mateu, J., Pebesma, E., Meyer, H., 2022. Nearest neighbour distance matching leave-one-out cross-validation for map validation. *Methods in Ecology and Evolution* 13, 1304–1316. doi:10.1111/2041-210X.13851.
- [25] Moreno-Torres, J.G., Raeder, T., Alaiz-Rodríguez, R., Chawla, N.V., Herrera, F., 2012. A unifying view on dataset shift in classification. *Pattern Recognition* 45, 521–530. doi:10.1016/j.patcog.2011.06.019.
- [26] Muñoz-Sabater, J., Dutra, E., Agustí-Panareda, A., Albergel, C., Arduini, G., Balsamo, G., Boussetta, S., Choulga, M., Harrigan, S., Hersbach, H., Martens, B., Miralles, D.G., Piles, M., Rodríguez-Fernández, N.J., Zsoter, E., Buontempo, C., Thépaut, J.N., 2021. ERA5-Land: a state-of-the-art global reanalysis dataset for land applications. *Earth System Science Data* 13, 4349–4383. doi:10.5194/essd-13-4349-2021.
- [27] Pausas, J.G., Keeley, J.E., 2021. Wildfires and global change. *Frontiers in Ecology and the Environment* 19, 387–395. doi:10.1002/fee.2359.
- [28] Persello, C., Bruzzone, L., 2012. Active learning for domain adaptation in the supervised classification of remote sensing images. *IEEE Transactions on Geoscience and Remote Sensing* 50, 4468–4483. doi:10.1109/TGRS.2012.2192740.
- [29] Ploton, P., Mortier, F., Réjou-Méchain, M., Barbier, N., Picard, N., Rossi, V., Dormann, C., Cornu, G., Viennois, G., Bayol, N., Lyapustin,

- 1505 A., Gourlet-Fleury, S., Pélissier, R., 2020. Spatial validation reveals poor predictive performance of large-scale ecological mapping models. *Nature Communications* 11, 4540. doi:10.1038/s41467-020-18321-y.
- [30] Podschwit, H., Jolly, W., Alvarado, E., Verma, S., Ponce, B., Markos, A., Aliaga-Nestares, V., Rodriguez-Zimmermann, D., 2022. Reliability of cross-regional applications of global fire danger models: A Peruvian case study. *Fire Ecology* 18, 25. doi:10.1186/s42408-022-00150-7.
- 1510 [31] Roberts, D.R., Bahn, V., Ciuti, S., Boyce, M.S., Elith, J., Guillera-Aroita, G., Hauenstein, S., Lahoz-Monfort, J.J., Schröder, B., Thuiller, W., Warton, D.I., Wintle, B.A., Hartig, F., Dormann, C.F., 2017. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography* 40, 913–929. doi:10.1111/ecog.02881.
- [32] Rousseau, J.S., Betts, M.G., 2022. Factors influencing transferability in species distribution models. *Ecography* 2022, e06060. doi:10.1111/ecog.06060.
- 1520 [33] Sandholt, I., Rasmussen, K., Andersen, J., 2002. A simple interpretation of the surface temperature/vegetation index space for assessment of surface moisture status. *Remote Sensing of Environment* 79, 213–224. doi:10.1016/S0034-4257(01)00274-7.
- [34] Schoener, T.W., 1968. The *Anolis* lizards of Bimini: Resource partitioning in a complex fauna. *Ecology* 49, 704–726. doi:10.2307/1935534.
- 1525

- [35] Sun, B., Feng, J., Saenko, K., 2016. Return of frustratingly easy domain adaptation, in: Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence (AAAI-16), AAAI Press, Phoenix, Arizona, USA. pp. 2058–2065. doi:10.1609/aaai.v30i1.10306.
- 1530 [36] Tuia, D., Persello, C., Bruzzone, L., 2016. Domain adaptation for the classification of remote sensing data: An overview of recent advances. IEEE Geoscience and Remote Sensing Magazine 4, 41–57. doi:10.1109/MGRS.2016.2548504.
- [37] Valavi, R., Elith, J., Lahoz-Monfort, J.J., Guillera-Arroita, G., 2019. 1535 blockCV: An R package for generating spatially or environmentally separated folds for k-fold cross-validation of species distribution models. Methods in Ecology and Evolution 10, 225–232. doi:10.1111/2041-210X.13107.
- [38] Vesk, P.A., Morris, W.K., Neal, W.C., Mokany, K., Pollock, L.J., 2021. 1540 Transferability of trait-based species distribution models. Ecography 44, 134–147. doi:10.1111/ecog.05179.
- [39] Vibhandik, P., Sawant, S., Joshi, A., Bidwe, R., 2026. A systematic literature review on forest fire susceptibility mapping using geo-spatial technology and future research directions. Discover Artificial Intelligence 1545 6, 921. doi:10.1007/s44163-026-00921-0.
- [40] Wadoux, A.M.J.C., Heuvelink, G.B.M., de Bruin, S., Brus, D.J., 2021. Spatial cross-validation is not the right way to evaluate map accuracy.

Ecological Modelling 457, 109692. doi:10.1016/j.ecolmodel.2021.109692.

- 1550 [41] Warren, D.L., Glor, R.E., Turelli, M., 2008. Environmental niche equivalency versus conservatism: Quantitative approaches to niche evolution. *Evolution* 62, 2868–2883. doi:10.1111/j.1558-5646.2008.00482.x.
- [42] Xu, Y., Dai, Y., Chang, L., Wang, Q., Dong, Y., 2026. Does your wildfire prediction model actually work, or just score well? arXiv preprint arXiv:2605.18911. doi:10.48550/arXiv.2605.18911, arXiv:2605.18911.
- 1555 [43] Yates, K.L., Bouchet, P.J., Caley, M.J., Mengersen, K., Randin, C.F., Parnell, S., et al., 2018. Outstanding challenges in the transferability of ecological models. *Trends in Ecology & Evolution* 33, 790–802. doi:10.1016/j.tree.2018.08.001.
- 1560 [44] Yebra, M., Dennison, P.E., Chuvieco, E., Riaño, D., Zylstra, P., Hunt, E.R., Danson, F.M., Qi, Y., Jurdao, S., 2013. A global review of remote sensing of live fuel moisture content for fire danger assessment: Moving towards operational products. *Remote Sensing of Environment* 136, 455–468. doi:10.1016/j.rse.2013.05.029.
- 1565 [45] Zanaga, D., Van De Kerchove, R., Daems, D., De Keersmaecker, W., Brockmann, C., Kirches, G., Wevers, J., Cartus, O., Santoro, M., Fritz, S., Lesiv, M., Herold, M., Tsendbazar, N.E., Xu, P., Ramoino, F., Arino, O., 2022. ESA WorldCover 10 m 2021 v200. Zenodo. doi:10.5281/zenodo.7254221.
- 1570