

Highlights

Local skill, unstable portability: sign reversal in pre-fire thermal predictors across Mediterranean wildfire regions, and the diagnostics that do not order it

Yunus Emre Cogurcu, Emrehan Metin

- Pre-fire thermal predictors add +0.06 to +0.15 AUC within all five fire regions
- The same block adds +0.004 AUC across 20 transfer directions, interval spans zero
- Highest niche-overlap pair fails both ways; lowest transfers both ways (points)
- Of 20 diagnostics, only conditional direction agreement tracks transfer (rho 0.84)
- Label-blind adaptation compresses transfer toward chance; only target labels help

Local skill, unstable portability: sign reversal in pre-fire thermal predictors across Mediterranean wildfire regions, and the diagnostics that do not order it

Yunus Emre Cogurcu^{a,*}, Emrehan Metin^a

^a*Çukurova University, Adana, Türkiye*

Abstract

Pre-fire thermal dryness separates a fire year from a normal year. Models built on it are rarely tested outside the region they were fitted in. Five Mediterranean wildfire regions were analysed here, on about 500 m cells, with MCD64A1 labels and spatially blocked validation. Six pre-fire thermal predictors were added to a terrain, fuel and greenness baseline.

The local skill does not travel, and almost all of it is lost before the region changes. Within-region ROC-AUC rose by +0.056 to +0.153 in every region under blocked cross-validation, but only +0.022 [−0.032, +0.077] when a whole burn scar is withheld. Across twenty ordered transfer directions the same block contributed +0.004 [−0.028, +0.036]. That is not distinguishable from zero, and its sign changes from pair to pair. The static baseline transferred no better, at 0.537 against 0.541. Three evaluations on **identical cells** locate the failure. A model that has the held-out burn scar in its training data scores 0.627; withholding that scar gives 0.552; and a model fitted 306 to 2,802 km away gives 0.559. The fire-specific residual is

*Corresponding author.

Email address: ycogurcu@cu.edu.tr (Yunus Emre Cogurcu)

+0.082 $[-0.011, +0.175]$ and the effect of the 2,800 km is -0.003 $[-0.075, +0.069]$. The failure is therefore a property of contiguous spatial holdout, not of separation distance and not of crossing a region boundary. Two thirds of the apparent collapse from a region-wide 0.782 is the evaluation area itself, which is 34 to 87 % burned with every negative fire-adjacent. Two predictors reverse their association between regions: elevation and the LST anomaly. Removing them cost -0.081 of within-region skill, supported in every region and mostly due to elevation. It changed transfer by $+0.014$ $[-0.017, +0.045]$. Both figures are post-selection estimates. A local cost is measured. No compensating transfer gain is.

The failure is invisible to the diagnostics that can be run before deployment. Twenty candidate diagnostics were rank-correlated against observed transfer. Only two had intervals excluding zero, and both were conditional. The stronger one measures agreement in the sign of each predictor’s association (Spearman $\rho = 0.84$ over sixteen directions). No marginal measure was shown to order the matrix, including predictor-space dissimilarity. At the point estimates, the pair with the highest burned-niche overlap failed in both directions, while the lowest transferred in both. Signed associations need labels on both sides. The family that works is therefore the one a practitioner does not have.

The mechanism of the residual is a reversal of sign. Predictors do not merely weaken across regions. Two reverse the direction of their association with burning, with bootstrap support. A distance in predictor space cannot see that. The same reversal appears inside one study area, between two fires eleven months apart on an identical grid. That arm rests

on one fire and eleven positive-carrying blocks, so it corroborates rather than establishes.

Label-free alignment pushed fourteen of twenty directions towards chance rather than repairing them. Transfer skill has to be measured, not inferred from similarity. The price is target labels: thirty-two labelled 5 km blocks recovered 85 to 89 % of the target’s matched ceiling in three of six directions.

Keywords: wildfire, model transferability, area of applicability, land surface temperature, domain adaptation, concept shift

1. Introduction

Wildfire is a defining disturbance of Mediterranean-basin landscapes, and a changing climate is reshaping where and how it burns [27]. The pressure to anticipate *where* fire will occur has produced a large and still-growing literature on fire susceptibility mapping. It has recently been synthesised from
5 erature on fire susceptibility mapping. It has recently been synthesised from both a methodological and an application perspective [39, 13]. The dominant pattern is well established. Geospatial predictors are assembled over a study region and paired with a historical record of burned area. A supervised classifier is then fitted, most often a random forest [4]. The resulting
10 susceptibility surface is published with a cross-validated AUC in the 0.85 to 0.95 range.

This paper is about a property that the pattern does not report. Predictors that describe the *state* of a surface in a particular season are more informative about that season than predictors that describe a place. They
15 are also, as we show, not portable: the relationship they carry is reparameterised locally, existing everywhere but pointing in different directions in

different places. The dynamic block is worth +0.056 to +0.153 ROC-AUC inside every region under blocked cross-validation, +0.022 when a whole burn scar is withheld, and contributes +0.004 between them, an estimate whose
20 interval spans zero, with a sign that is a property of the source-target pair rather than of the block.

The obvious reading of that pair of numbers is a trade-off, and our own earlier framing took it. The data do not support it. The static baseline transfers at a mean of 0.537 against the thermal model’s 0.541, so the predictor
25 class that ought to travel does not travel either, and removing the reversing predictors returns +0.014 of transfer on an interval that also spans zero. Two nulls on the portability axis are not an exchange. What is measured is a local gain that does not export, and a local cost when the reversing predictors are removed, with no compensating transfer gain on either arm.

30 1.1. Portability goes unmeasured

A susceptibility model’s reported skill is almost always an estimate of *within-region* performance. Held-out folds come from the same study area and season, often from the same fire event, and where folds are drawn at random over cells, spatial autocorrelation inflates the estimate further. The
35 problem is documented across ecological modelling [31, 29] and is addressed by spatially blocked cross-validation [37, 23]. Blocking removes the inflation that random folds produce, and it is the right correction for what it corrects. It does not make the estimate honest about a fire the model has not seen. Section 4.3 measures that on identical cells: a model with the held-out burn
40 scar in its training data scores 0.634 on that scar’s area, and withholding the scar gives 0.552. Part of the apparent gap between a region-level figure and

either of those is the evaluation area rather than the model.

Because only that side of the ledger is reported, portability is never entered at all. A predictor block is adopted on the strength of the increment
45 it delivers inside its training footprint, and whether that increment survives a change of region is not asked, even though any regional product built from locally trained models implicitly promises generalisation beyond that footprint. For fire specifically the transfer question is asked far less often than in adjacent fields. Meteorological fire-danger indices are known not
50 to port cleanly between fire environments [30], and the two recent studies that test model transfer systematically, across Mediterranean countries [7] and US counties [15], both report that it largely succeeds between similar regions. Both transfer models whose dominant predictors are *spatially stationary*: terrain, human modification, night-time lights, population density,
55 long-term climatologies. Each describes a place rather than a season, which is the class our thesis expects to travel well. Species distribution modelling has examined transferability far more systematically [43], and there geographic and environmental similarity do *not* reliably predict transfer success [38, 32].

1.2. Pre-fire thermal dryness is the natural test case

60 The dynamic class used here is pre-fire thermal dryness, and it is the class most plausibly *expected* to transfer. Standard susceptibility predictors are terrain, fuel type, greenness, long-term climatology and proximity to roads and settlements, and most are static or near-static over the timescale at which fire danger varies. Such a set explains poorly why a particular summer burned and the preceding one did not. What changes between those
65 summers is the state of the surface: fuel dryness, moisture stress, and how

anomalously hot the land surface has become relative to its own baseline. Satellite thermal observation gives direct access to part of that state through fuel moisture content [44], and pairing surface temperature with a vegetation index was established as a live-fuel-moisture estimator for fire-danger rating by Chuvieco et al. [6]. The Temperature-Vegetation Dryness Index [33] formalises that feature space into an internally normalised measure. In principle it should be less exposed to absolute-temperature offsets between regions than raw land surface temperature. That theoretical portability advantage is tested here, and it is not found: the two internally normalised channels transfer no better than the four absolute ones, at means of 0.544 and 0.548 over the twenty directions (Appendix A(f)).

That pre-fire thermal state carries genuine information about subsequent fire is established [17, 19, 18, 8]. What has not been established is whether the predictive *skill* of a classifier trained on it in one fire region survives strict, label-free application to another. The physics linking moisture stress to combustion is universal, so portability should be most expected for this class, which is what makes it diagnostic: a loss here cannot be dismissed as a peculiarity of a locally defined covariate.

1.3. Why the loss is invisible to the diagnostics in use

A deficit on transfer can arise two ways. Under **covariate shift** the marginal distribution of the predictors differs while the predictor-response relationship is preserved, which is in principle correctable without target labels, by per-region standardisation or covariance alignment such as CORAL [35], an approach with an established remote-sensing literature [36]. Under **concept shift** the conditional relationship itself changes: a predictor pos-

itively associated with burning in one region is negatively associated with it in another. No label-free realignment of input distributions can repair a sign reversal, because the information needed to detect it is exactly the information being withheld. The taxonomy is canonical [25].

This distinction is what makes the failure hard to anticipate. The diagnostics the field relies on to decide whether a model may be applied to new ground, above all area-of-applicability measures built on dissimilarity in predictor space [21, 22, 16], are marginal by construction. A target region can sit well inside the training data’s predictor envelope while the relationship between those predictors and burning points the other way. This paper measures both sides on the same pairs.

1.4. Contributions

Three findings carry this paper. Each is stated at the strength its interval supports, and the nearest prior work is named inside the claim.

Contribution 1. Local skill and portability, measured separately and per direction. The thermal block is worth $+0.056$ to $+0.153$ ROC-AUC inside every one of five regions under blocked cross-validation, with every bootstrap interval above zero and the result holding in both analysis populations. That increment is substantially a property of interleaved holdout: withholding a whole burn scar leaves $+0.022$ $[-0.032, +0.077]$, and across regions $+0.004$. Across the twenty ordered transfer directions it contributes $+0.004$ $[-0.028, +0.036]$. That interval spans zero under all four between-direction resampling units we computed, and the sign varies by pair: paired deltas run from -0.148 to $+0.132$, twelve positive and eight negative. Two controls bound the reading. The baseline arm transfers at

a mean of 0.537, against the thermal model’s 0.541, so the static predictor class is not the portable one either. Three evaluations scored on **identical cells** locate where the skill goes. A model with the held-out burn scar in its training data returns 0.627 on that scar’s area; withholding the scar returns 0.552; and a model fitted 306 to 2,802 km away returns 0.559. The fire-specific residual is +0.082 [−0.011, +0.175]; the effect of the 2,800 km is **−0.003 [−0.075, +0.069]**. **The failure is a property of contiguous spatial holdout**, not of separation distance and not of crossing a region boundary. It is not shown to be a property of the fire event either, because the held-out patch is defined by the labels and its identity cannot be separated from its location. Two thirds of the apparent fall from the region-wide 0.782 is the evaluation area, which is 34 to 87 % burned against 3.8 to 28.7 % for a region, so a scar-level result compared against a region-level reference overstates the collapse. Separation does not order the matrix either, at $\rho = -0.32$ with an interval spanning zero. Six directions are nonetheless below chance with interval support, which no account of merely lost skill explains, and that residual is what Contributions 2 and 3 address. Feature removal measures the local cost. Dropping the two reversing predictors, elevation and the LST anomaly, costs −0.081 of mean within-region AUC, supported in every region. Roughly three quarters of that is elevation, a *baseline* terrain variable. Mean transfer changes by +0.014, whose interval also spans zero. Both are post-selection estimates (Section 4.6b). A local cost is measured. No compensating transfer gain is. The within-region half is not novel, comparable results existing for these landscapes [1, 12]; the paired contrast against portability is. Dimarco et al. [7] transfer a predominantly stationary

predictor set successfully across a comparable Mediterranean design, though their response variable is ignition rather than burned area (Section 5.7).

Contribution 2. The loss is invisible to the diagnostics that
145 **could be run before deployment, and visible to one that cannot.**

Twenty candidate diagnostics from four families are evaluated against observed transfer under one bootstrap framework. Area-of-applicability-style predictor-space dissimilarity [21, 22, 16], climatic and geographic distance, learned domain separability, niche-overlap statistics and burn-pattern regime
150 distances were none of them shown to order the matrix; the domain classifier is at ceiling, with $AUC \geq 0.96$ for every pair. Only two diagnostics have intervals excluding zero. Both measure conditional direction agreement, that is, whether each predictor’s signed association with burning points the same way in both regions (Spearman $\rho = +0.84$ and $+0.81$ over sixteen directions
155 from eight pairs). Both are computed on a data-selected subset of predictors, and their all-nine-feature counterparts span zero, so the result rests on that selection. Equalising the families onto a common twelve directions leaves the ordering unchanged, so it is not an artefact of unequal samples. Signed associations need burned labels on both sides, so this is a mechanism
160 diagnosis rather than a pre-deployment screen, and the marginal family, the only one runnable before deployment, is the one that fails. The same point survives without any ranking, at the point estimates: the pair with the highest burned-niche overlap fails in both directions while the lowest transfers in both. That contradicts an expectation the fire literature carries [7, 15]
165 and agrees with species distribution modelling [38, 32]. Shift decomposition in applied remote sensing is not itself new [11]; the fire application and the

head-to-head of four diagnostic families on the same pairs are.

Contribution 3. The mechanism is a reversal of sign, and it survives holding geography fixed. Predictors do not merely weaken across
170 regions, they reverse the direction of their association with burning, which is
why a distance in predictor space cannot see the failure. The sharpest reversal
is elevation, whose association points opposite ways in Manavgat and in
Bejís and Muğla, with each region’s own interval excluding chance and the
paired difference excluding zero (Appendix B, Table B3). The same reversal
175 appears inside a single study area, between two fires eleven months apart
on an identical grid, where season and population also differ but place does
not (Section 4.8). That arm rests on one fire and eleven positive-carrying 5
km blocks, below the floor this design sets for itself, so it corroborates the
mechanism rather than establishing it.

180 Two consequences follow. Both are reported as supporting results, not
leading ones. First, label-free alignment by standardisation and covariance
alignment [35] does not repair transfer. It compresses transfer towards chance
in fourteen of twenty directions. We believe this is the first application
of covariance alignment to fire susceptibility. Second, the residual gap is
185 conditional, so the resource that closes it is target labels. That price is
measurable. Thirty-two labelled 5 km blocks recover 85 to 89 % of the
target’s matched ceiling in three of six directions, 51 to 57 % in two more,
and 30 % in the sixth.

Alongside these, a leakage-audited, spatially blocked evaluation and trans-
190 fer protocol is released with code, configuration and frozen outputs, so that
most of this result can be re-run rather than taken on trust. One component

is an exception and is named in the data-and-code availability statement:
the reproduction-check code is not in the released repository at the commit
of record. The few-shot export, previously also listed as an exception, was
195 regenerated at commit 6d7a6a71 and its validator now passes on every check.
This matters given evidence that wildfire transfer conclusions are sensitive
to evaluation design and task formulation [42].

A companion paper treats the observational layer beneath this one, where
preprocessing decisions taken before any model is fitted are shown to carry
200 their own uncertainty budget.

Section 2 reviews the relevant literature. Section 3 describes the study
regions, the data, and the modelling, cross-validation, transfer and adapta-
tion protocol. Section 4 reports results. Section 5 discusses them, including
limitations. Section 6 concludes.

205 **2. Related work**

2.1. Fire susceptibility mapping, and what it reports

Fire susceptibility mapping has been synthesised recently from both method-
ological and application perspectives [39, 13]. The dominant pattern is sta-
ble: geospatial predictors are assembled over a study region, paired with a
210 historical burned-area record, and fitted with a supervised classifier, most
often a random forest [4], and the resulting surface is published with a cross-
validated AUC in the 0.85 to 0.95 range. That figure is a within-region
estimate. Comparable within-region results already exist for the landscapes
studied here [1, 12].

215 *2.2. Pre-fire thermal dryness*

Satellite thermal observation gives repeated access to surface state through fuel moisture content [44]. Pairing land surface temperature with a vegetation index was established as a live-fuel-moisture estimator for fire-danger rating by Chuvieco et al. [6], and the Temperature-Vegetation Dryness Index [33] formalises that feature space into an internally normalised measure, which is why it is often expected to travel between regions better than raw temperature. That pre-fire thermal state carries information about subsequent fire is established [17, 19, 18]. Gelabert et al. [8] found dead fine fuel moisture and its anomalies to be the most influential predictor of human-
225 caused ignition likelihood across Europe, testing generalisation by pooled fitting with per-site evaluation. What has not been tested is whether a classifier built on this class survives strict, label-free application to a region it never saw.

2.3. Spatial validation and transferability

230 Random cross-validation over spatially autocorrelated cells inflates skill estimates [31, 29], and spatially blocked designs are the standard remedy [37, 23]. The remedy is not uncontested: blocked schemes have been argued to introduce pessimistic bias of their own [40, 24, 5]. This paper uses spatial blocking throughout and reports every result at three block sizes, so that a
235 reader can see which verdicts depend on the choice.

Blocking corrects the inflation that random folds produce within a region. It does not speak to performance on a fire the model has not seen, which Section 4.3 measures on identical cells. For that, the field’s instrument is the area of applicability and related predictor-space dissimilarity measures

240 [21, 22, 16], which ask whether a target’s predictor values fall inside the training data’s envelope. Species distribution modelling has examined transferability far more systematically than fire has [43], and two studies there report that geographic and environmental similarity do not reliably predict transfer success [38, 32], and the transferability of satellite-derived empirical
245 models is an active question in fire remote sensing itself [20].

2.4. Dataset shift and label-free adaptation

The canonical taxonomy separates covariate shift, where the predictor distribution moves but the predictor-response relationship holds, from concept shift, where the relationship itself changes [25]. Covariate shift is in
250 principle correctable without target labels, by per-region standardisation or by covariance alignment such as CORAL [35], and domain adaptation has an established remote-sensing literature [36, 28]. Concept shift is not: no realignment of input distributions can repair a reversal in the sign of an association, because detecting it requires the labels being withheld. Shift
255 decomposition in applied remote sensing is not itself new [11]. We are not aware of a prior application of covariance alignment to fire susceptibility, fire occurrence or burned-area prediction.

2.5. Cross-region generalisation of fire models

Few studies test fire-model transfer directly. Podschwit et al. [30] re-
260 port, in a Peruvian case study, that meteorologically derived danger indices do not port cleanly between fire environments. WildfireGenome [15] runs a leave-one-county-out matrix across seven US counties and reports strong within-county performance with highly variable off-diagonal transfer, similar

pairs transferring well and dissimilar pairs collapsing; its label, however, is
 265 a principal-component composite of hazard *indicators* rather than observed
 burned area, and it applies no adaptation. Xu et al. [42] argue that wild-
 fire transfer conclusions depend strongly on evaluation design, a caution we
 address by fixing the protocol in a project log before the diagnostics were
 computed and by reporting every sensitivity axis; that log is not a formal pre-
 270 registration, and Section 3.11 states what was fixed and when. Kondylatos
 et al. [14] provide Mesogeos, a 1 km Mediterranean datacube for data-driven
 wildfire modelling.

The nearest neighbour, and the contrast the present paper draws.

Dimarco et al. [7] is the closest Mediterranean analogue: 500 m predic-
 275 tors harmonised across four countries, tree ensembles under spatial cross-
 validation, transfer tested both leave-one-country-out and as a full 4 x 4
 matrix, with every transfer exceeding AUC 0.80 and bioclimatically similar
 countries scoring higher. They apply no domain adaptation.¹

Much of the design is shared — Mediterranean regions, 500 m cells,
 280 MCD64A1-derived targets, tree ensembles, spatially aware validation, an ex-
 plicit transfer matrix — and two things differ. **The predictor class:** their
 model rests on attributes of a place, all spatially stationary, ours on the state
 of a surface in one season. **The response variable:** theirs is human-driven
 ignition, ours is burned area, and those are different quantities with different
 285 dominant controls. Read together the two results suggest that the relation-
 ship between domain similarity and transfer success may be predictor-class

¹Their Results text refers to "LST anomalies", which is inconsistent with their own

dependent, with Vesik et al. [38] and Rousseau and Betts [32] aligning from species distribution modelling while Dimarco et al. and WildfireGenome run against it. We present that as a live disagreement, and Section 5.7 states why our own data do not settle it.

Methods, where no land surface temperature or TVDI variable appears and the only temperature predictor is a static ERA5-Land seasonal climatology [26]. We follow their Methods, which we take to be authoritative, and note the discrepancy so that the comparison is transparent. Every statement about their work was checked against the full text, read 2026-08-08.

3. Data and methods

3.1. Study regions and temporal windows

Five Mediterranean wildfire regions are analysed (Fig. C.1): Manavgat 2021 and Muğla 2021 in Türkiye, Bejís 2022 in Spain, North Evia 2021 in Greece and Montiferru 2021 in Sardinia. Each is a place-based rectangular area of interest in EPSG:4326, defined from place coverage rather than from a fire perimeter, and not tuned on burned prevalence, gate outcome or any model metric. One AOI choice was label-informed and is stated as such. The North Evia box was extended after the legacy box was found to carry an atypically high burned prevalence. The extended geometry was then defined from place anchors, and the legacy variant is kept as a sensitivity arm (Appendix A). The AOIs are deliberately not clipped to fire perimeters, so that unburned cells around each fire form the negative class rather than being excluded by construction. A sixth region, Kozan 2023, is carried as a negative control and is excluded from all modelling by the gate of Section 3.3.

Each region has two non-overlapping windows. The **predictor window** closes the day before the **label window** opens, so no predictor observation can post-date the first labelled burning. Window lengths vary between regions with the event: 57 to 61 days for predictors and 35 to 59 days for labels,
 315 counted inclusively. A four-year baseline of window-symmetric composites precedes each predictor window and supplies the climatological reference for the anomaly channels. Region identifiers, bounding boxes, window dates and baseline years are registered in `core/regions.py` and are reproduced here.

Table 1: Study regions, areas of interest and temporal windows. Bounding boxes are in EPSG:4326, as registered. Window lengths in brackets are inclusive day counts. The baseline years are the four window-symmetric years preceding each predictor window.

Region	Bounding box (lon min, lat min, lon max, lat max)	Predictor window	Label window	Baseline years
Manavgat 2021 (Türkiye)	31.05, 36.72, 31.85, 37.35	2021-06-01 to 2021-07-27 (57 d)	2021-07-28 to 2021-08-31 (35 d)	2017, 2018, 2019, 2020
Bejís 2022 (Spain)	-1.05, 39.68, -0.35, 40.15	2022-06-15 to 2022-08-14 (61 d)	2022-08-15 to 2022-09-30 (47 d)	2018, 2019, 2020, 2021
Muğla 2021 (Türkiye)	27.1, 36.6, 28.9, 37.45	2021-06-01 to 2021-07-28 (58 d)	2021-07-29 to 2021-09-15 (49 d)	2017, 2018, 2019, 2020
North Evia 2021 (Greece)	23.05, 38.55, 23.85, 39.15	2021-06-05 to 2021-08-02 (59 d)	2021-08-03 to 2021-09-30 (59 d)	2017, 2018, 2019, 2020
Montiferru 2021 (Italy)	8.45, 40.05, 8.75, 40.27	2021-05-25 to 2021-07-23 (60 d)	2021-07-24 to 2021-08-31 (39 d)	2017, 2018, 2019, 2020

3.2. Burned-area label and the ~ 500 m analysis grid

320 Labels come from MODIS MCD64A1 Collection 6.1 [9] retrieved through Google Earth Engine [10], whose omission and commission characteristics [3] bound every model here. The analysis grid is **reconstructed rather**

than native: the pipeline’s 30 m EPSG:4326 reference grid is partitioned into 17 x 17 blocks, giving a nominal 510 m cell that approximates rather
325 than reproduces the MODIS cell, and is square in degrees but not on the ground. A cell’s burn date is the mode of its positive sub-pixel day-of-year values, tested against the region’s label window; the label never affects a cell’s eligibility for modelling. Two safeguards are recorded rather than assumed: cells burning inside a region’s own predictor window are removed, which ran
330 for three of five regions, and burning in earlier years is screened for none. Appendix C.1 gives the full specification, including what follows from the grid’s shape.

3.3. Burned-landcover admissibility gate

Before any modelling each region passes a gate asking what fraction of
335 its burned cells is dominated by natural vegetation, using ESA WorldCover classes on the same cells. A region is admitted at 0.50 with at least 30 burned cells, and rejected as a cropland-dominated control at 0.50 cropland. The purpose is to separate burned area produced by natural-fuel combustion from post-harvest stubble burning, which MCD64A1 does not distinguish.
340 Verdicts are in Section 4.1; the full rule is in Appendix C.1.

3.4. Predictor variables

Ten predictors are used, four baseline and six thermal, all summarised per cell over the predictor window, through the chain shown in Fig. C.2. All optical and thermal predictors come from Landsat 8 Collection 2 Level-2
345 (LANDSAT/LC08/C02/T1_L2), quality-screened per pixel from the QA_PIXEL

band with the water bit deliberately preserved; the coarse-resolution thermal input is MODIS/061/MOD11A1.

- **NDVI**, the predictor-window median of Landsat surface reflectance, is the baseline’s one time-varying member.
- 350 • **Elevation** and **slope** come from the Copernicus DEM GLO-30.
- **Land cover** is ESA WorldCover v200 [45], entering as the dominant class code.
- **Current LST**, the predictor-window median of Landsat surface temperature in °C.
- 355 • **LST anomaly**, a z-score of the current-window LST median against the four baseline years, set to no-data where the baseline standard deviation is below 1.0 °C or observations are too few.
- **TVDI** and **TVDI difference**. The Temperature-Vegetation Dryness Index [33] is computed in the LST-NDVI feature space, with wet and dry edges taken as the 2nd and 98th LST percentiles within each of
360 twenty NDVI bins and the index clamped to [0, 1]. The edges are percentiles of the LST values a given scene contains, so they are fitted per AOI and per window, and a TVDI of 0.5 denotes a different physical moisture state in each region. The companion paper reports what that scene dependence does and does not explain. **TVDI difference**
365 **is the raw anomaly of that index against the same four window-symmetric baseline years used for the LST anomaly: `tvdi_difference = current_tvdi - mean(baseline_tvdi)`, in index units rather than**

standard deviations, which is why it is reported alongside the z-scored
channel rather than in place of it.

- **Downscaled LST and fused LST.** A MODIS-to-Landsat downscaling model is trained on the predictor window and applied to the full 30 m grid. The fused product equals observed Landsat LST wherever that is valid, and the downscaled surface only where it is not. Gap-filling therefore never replaces or blends a valid observation. The gap-filled share is 0.11 % to 9.70 % by region. The downscaler’s own inputs include coordinates, which is the one route by which a coordinate-derived surface re-enters a feature set from which Section 3.14 excludes coordinates. Appendix A(g) reports the increment without these two channels.

3.5. *Cell aggregation, validity and analysis populations*

Cell-level values are means over the ~ 510 m cell, computed from valid 30 m pixels only, with the valid fraction recorded. A cell is `valid_for_modeling` when it meets a 30 % valid-pixel floor and no thermal channel is missing. Two populations are then defined. The **primary** population is natural vegetation, cells whose combined tree, shrub and grass fraction reaches 0.50, which excludes cropland from every burnable mask. The **secondary** population is all valid cells, reported throughout as a sensitivity. Every headline result is given for both.

3.6. *Feature sets and classifier*

Two nested feature sets are compared. The **baseline** set is NDVI, elevation, slope and land cover. The **thermal** set is the baseline plus the

six thermal channels. Land cover is one-hot encoded and numeric features are standardised inside the fitting pipeline. The classifier is a random forest [4] with 300 trees, unlimited depth, `min_samples_leaf = 3`, balanced class weights and `random_state = 42`, identical for every region, population, feature set and transfer direction, so that no comparison in this paper is confounded by a model choice.

3.7. *Spatial-block cross-validation and bootstrap uncertainty*

Cross-validation is spatially blocked. Blocks are squares of the analysis grid formed as `row_500m // B by col_500m // B`, and folds are drawn over whole blocks, so no block is split between training and test. Five folds are used. The block size B is reported at 2, 10 and 20 cells, about 1, 5 and 10 km, for the within-region analysis.

Uncertainty is a spatial-block bootstrap: blocks are resampled with replacement, 1000 replicates, seed 42, with 2.5 and 97.5 percentiles as the interval. Because it is blocks that are resampled, what bounds an interval’s reliability is the number of blocks carrying at least one burned cell, and those counts are reported alongside the intervals. Where a verdict rests on too few such blocks it is stated as indicative rather than as an interval.

3.8. *Cross-region transfer protocol*

For each ordered pair of regions the model is fitted on the source population and applied to the target population with **no target label used at any point**: no refitting, no threshold selection, no calibration. Both the baseline and the thermal feature set are transferred, so that the thermal block’s contribution is a paired difference within a direction rather than a

comparison across directions. All twenty ordered directions among the five regions are computed.

3.9. Label-blind domain adaptation

420 Two label-free remedies are tested on every direction, in both feature sets.

Region-wise z-score. Each region’s numeric features are standardised using its own statistics, source statistics from source data and target statistics from target data, never pooled. The classifier is refitted on the z-scored source and applied to the z-scored target. This removes first- and second-
425 order marginal offsets.

CORAL after region-wise z-score. Starting from the z-scored features, the source covariance is aligned to the target covariance by the standard whitening-recolouring map [35], with regularisation $\lambda = 10^{-5}$ and eigenvalues floored at 10^{-12} . Critically **the transform is applied to the source**
430 **only**; the target is left exactly as it is. The classifier is refitted on the aligned source. Neither variant sees a target label, and both are verified label-blind at run time.

λ sensitivity was assessed over nine values from 0 to 10^{-1} on four directions in both feature families, moving transfer AUC by at most 0.014 in
435 any direction and 0.008 within the thermal family. The canonical $\lambda = 1$ lies outside the released sweep and is reported separately in Appendix A(b).

3.10. Transfer-gap decomposition and the concept-shift diagnostic

For each direction the gap between the target’s own within-region skill and the raw transfer result is split in two. One part is what the best label-free
440 adaptation recovers, and the other is what it does not. The recovered fraction

is defined as (adapted – raw) / (within – raw), signed and unclipped, and its interval comes from the same paired bootstrap. The recovered part bounds what covariate-level correction can achieve. The remainder is unrecovered by that correction; Section 4.3 shows it should not be read as a conditional
445 residual, because much of it is incurred inside a single region.

The mechanism is diagnosed by **signed univariate association**. For each numeric predictor the raw ROC-AUC of that predictor against **burned** is computed in each region and never folded to $\max(\text{AUC}, 1 - \text{AUC})$, so a value below 0.5 is read as a direction rather than as weakness. A reversal is
450 called bootstrap-supported only when the two regions’ intervals are disjoint, under the same 10-cell spatial-block bootstrap.

3.11. Transferability diagnostics versus transfer

Twenty candidate diagnostics from four families are computed for every region pair and rank-correlated (Spearman) against observed raw transfer
455 AUC, under one bootstrap framework that resamples unordered region pairs with both of their ordered directions travelling together. The families are marginal predictor-space measures $P(x)$, including area-of-applicability dissimilarity, climatic and geographic distance and a learned domain classifier; burned-niche overlap $P(x|y=1)$; burn-pattern regime distance $P(y)$; and con-
460 ditional direction agreement $P(y|x)$. No family-wise error control is claimed, and the number of variants computed within each family is reported with the result. Appendix C.2 gives the full specification of each diagnostic and of the bootstrap.

3.12. Interventions

465 **Pooled multi-region training.** Leave-one-region-out: the model is trained on the pooled primary populations of the other four regions and evaluated on the held-out region, with folds blocked as before. This asks whether pooling recovers what single-source transfer loses.

Removal of direction-reversing features. The two predictors whose
470 signed association reverses between regions with bootstrap support are **elevation_mean** (Manavgat against Bejís and against Muğla) and **lst_anomaly_mean** (Bejís against Evia). They are dropped and everything is refitted, within-region and across every direction, and the two are also dropped singly so the cost can be attributed. This measures what the removal costs locally and
475 what it returns on transfer on the same footing. Note that the two features are selected by the same reversal analysis against which the result is then read, so both quantities are post-selection estimates and no correction for that selection is applied.

3.13. Controls on the transfer path

480 Four evaluations establish what the transfer arms are measuring. All use the transfer protocol of Section 3.8 unchanged: fit on the source cells, apply to the target cells, no refit, no recalibration, no threshold selection. All were verified to reproduce the frozen cross-region AUCs exactly before being used.

Within-region half-split. The modelled cells of a region are cut in two
485 at the median of a grid axis, a model is fitted on one half and applied to the other. Both axes and both directions are run, giving four splits per region. A split is discarded when either half is single-class, which happens twice in

one region. Per-split source and target positive counts are reported, because they are unequal and bear on the interpretation.

490 **Foreign-region evaluation.** Each held-out scar area is additionally scored with a model fitted on each of the other four regions in turn, and those four results averaged, so that a same-region model and a foreign one are compared on identical cells. This is the arm that separates the effect of withholding the fire from the effect of separation distance.

495 **Leave-one-scar-out.** Burned connected components are found with 8-connectivity on the analysis grid. Every component of at least 50 cells is held out together with all cells within a buffer of it, a model is fitted on the remainder of the same region, and applied to the held-out area. Buffers of 2, 5 and 10 km are run. A scar held out at 2 km inside its own region is
500 close to its training data and yet wholly unseen, which is what makes the comparison with the foreign-region arm informative. The held-out patch is defined by the labels, so this arm cannot separate the identity of a fire from its location.

Appendix A(h) reports all four in full, with per-split and per-scar tables.
505 Intervals for the four-row ladder are Student t over the eight held-out scars, which is that arm’s resampling unit, rather than the spatial-block bootstrap of Section 3.7 used everywhere else in this paper.

3.14. Leakage control

An explicit forbidden-column set is enforced at every model fit. Coordinates (lon, lat, row, col and their normalised forms), every burn-date and
510 label-provenance column, and the agreement fraction are excluded from all feature sets, and the enforcement runs as an assertion rather than a con-

vention. The natural-vegetation mask is used only to define the population,
never as a predictor. Spatial blocking prevents a cell from sharing a fold with
515 its own neighbours.

Reproducibility. All randomness uses seed 42 and the bootstrap uses
1000 replicates throughout. The transfer and adaptation analysis runs in
an environment separate from the upstream pipeline’s. Every within-region
model was therefore refitted there and compared against the frozen up-
520 stream output, and the independently implemented adaptation was com-
pared against the pipeline’s own. The within-region comparisons agree ex-
actly, and the twenty transfer directions agree to within 1.6×10^{-7} . Random-
forest fits are not bit-identical across library versions. All numbers here were
therefore produced under scikit-learn 1.9.0, or verified against it. If that
525 version is not fixed, cross-region point estimates carry an implementation
tolerance of about ± 0.02 to 0.03 . The companion paper reports the version
sensitivity and the reproduction check in full.

Sensitivity analyses. Every headline result is repeated across two anal-
ysis populations, three spatial-block sizes, the CORAL regularisation sweep
530 and both feature sets. Where a conclusion depends on one of those choices,
the dependence is reported rather than resolved by choosing the favourable
setting.

3.15. Same-geography event-to-event comparison (Muğla 2021 versus 2022)

Muğla admits a comparison in which place is held fixed and the event
535 varies: a second fire burned inside the identical AOI, on the identical analysis
grid, eleven months after the first. Signed univariate AUCs are computed
for both arms under the same 10-cell spatial-block bootstrap used elsewhere.

Season, year and population all differ between the arms, since the 2022 arm is defined by removing the 2021 scar, and Section 4.8 reports what that costs. Appendix C.3 gives the window dates, the population construction and the asymmetry audit.

4. Results

4.1. Study regions and the admissibility gate

All five candidate regions pass the burned-landcover gate as wildfire candidates and the negative control fails it as intended. Kozan 2023 returns a natural-vegetation fraction of 0.017 against the 0.50 threshold and is excluded from all modelling. The separation is not marginal: the five admitted regions carry fractions of 0.723 to 0.991, so the threshold falls in an empty interval rather than between neighbouring cases. That is consistent with the gate separating burned area produced by natural-fuel combustion from post-harvest stubble burning, which MCD64A1 does not distinguish, though one negative control cannot establish it.

North Evia is analysed on an extended AOI. Relative to the legacy box, the extended box leaves the burned scar essentially unchanged while cutting overall prevalence from 0.361 to 0.122 and TSG prevalence from 0.676 to 0.287. The effect of that choice on transfer is reported in Appendix A(a).

4.2. Within-region: the thermal increment replicates in five regions

Adding the six thermal predictors to the baseline raises spatially blocked out-of-fold ROC-AUC in every region. The increment's bootstrap interval excludes zero in all five regions at 1 km and at 5 km blocking (Fig. C.3).

Table 2: Region summary: populations and gate outcomes. Counts from each region’s Step 8A dataset statistics; gate fractions from each region’s burned-landcover gate output. TSG = the primary natural-vegetation population. The TSG columns use the canonical modelled population, `burnable_tree_shrub_grass` and `valid_for_modeling == True`, which is the population every model in this paper was fitted and scored on.

Region	Total cells	Valid cells	Burned	Prevalence (all valid)	TSG cells	Burned in TSG	TSG prevalence	Burned natural- veg frac- tion	Gate verdict
Manavgat 2021	24,150	24,087	796	0.033	20,511	784	0.038	0.984	pass
Bejís 2022	15,759	15,759	1,103	0.070	15,190	1,100	0.072	0.991	pass
Muğla 2021	73,098	73,045	3,026	0.041	41,730	2,911	0.070	0.958	pass
North Evia 2021 (ex- tended)	22,925	22,906	2,788	0.122	9,298	2,664	0.287	0.945	pass
Monte Pirri 2021	3,234	3,173	697	0.220	2,544	539	0.212	0.723	pass

Those are the two scales this design supports as intervals. At 10 km the point estimates hold, from +0.048 to +0.154. They rest on 6 to 33 positive-carrying blocks, so they are reported as indicative rather than as intervals. The table note gives the reason.

565 *Table note (resampling units).* The bootstrap resamples spatial blocks, so what bounds an interval’s reliability is the number of blocks that carry at least one burned cell. Those counts fall quickly as blocks coarsen. At 2 cells they are 235 (Manavgat), 302 (Bejís), 843 (Muğla), 716 (Evia) and 192 (Montiferru), out of 5 439, 3 967, 11 316, 2 566 and 743 blocks. At 10
570 cells they are 28, 19, 70, 41 and 16, out of 237, 176, 576, 155 and 35. At 20 cells they are **12, 6, 33, 15 and 6**, out of 60, 48, 167, 50 and 12. No bootstrap replicate was invalid at any block size, with one exception. At 20 cells Bejís had 5 of 1000 replicates single-class. That is the symptom of the six positive-carrying blocks just reported, and a further reason to
575 read the row as indicative. An equal-tailed percentile interval built on six positive-carrying blocks has no meaningful coverage. Montiferru at 20 cells also feeds only 12 groups into a 5-fold grouped split, so its models train on about ten blocks each. **The 20-cell row of this table should be read as indicative rather than as an interval.** The 10-cell row is the coarsest
580 blocking this design supports properly, because every region there has 16 to 70 positive-carrying blocks. The increment holds at that scale in all five regions. Source: `paper/referee2_numbers.md`, block C, counted from the frozen per-cell prediction tables.

This within-region result is not itself novel. It is reported because the
585 transfer arms below are measured against it.

Table 3: Within-region baseline versus thermal performance and block-size robustness. Primary (TSG) population; spatially blocked 5-fold CV (Section 3.7); paired spatial-block bootstrap, 1000 replicates. Block sizes 2/10/20 cells $\approx$ 1/5/10 km.

Region	Block ($\approx$ scale)	Baseline AUC	Thermal AUC	Δ AUC	Δ AUC 95% CI
Manavgat 2021	2 ($\sim$ 1 km)	0.803	0.870	+0.067	[+0.055, +0.079]
	10 ($\sim$ 5 km)	0.748	0.797	+0.050	[+0.023, +0.077]
	20 ($\sim$ 10 km)	0.683	0.731	+0.048	[+0.014, +0.085]
Bejís 2022	2 ($\sim$ 1 km)	0.862	0.918	+0.056	[+0.048, +0.065]
	10 ($\sim$ 5 km)	0.779	0.825	+0.045	[+0.018, +0.069]
	20 ($\sim$ 10 km)	0.739	0.795	+0.057	[+0.031, +0.090]
Muğla 2021	2 ($\sim$ 1 km)	0.743	0.859	+0.116	[+0.106, +0.125]
	10 ($\sim$ 5 km)	0.698	0.777	+0.079	[+0.050, +0.105]
	20 ($\sim$ 10 km)	0.673	0.733	+0.061	[+0.030, +0.094]
North Evia 2021 (ext.)	2 ($\sim$ 1 km)	0.759	0.912	+0.153	[+0.142, +0.166]
	10 ($\sim$ 5 km)	0.716	0.864	+0.148	[+0.119, +0.183]
	20 ($\sim$ 10 km)	0.679	0.833	+0.154	[+0.124, +0.189]
Montiferu 2021	2 ($\sim$ 1 km)	0.781	0.883	+0.101	[+0.080, +0.125]
	10 ($\sim$ 5 km)	0.620	0.720	+0.100	[+0.017, +0.186]
	20 ($\sim$ 10 km)	0.555	0.681	+0.126	[+0.053, +0.228]

4.3. Cross-region transfer, and what label-free adaptation does to it

Raw transfer is heterogeneous and includes anti-predictive directions. Raw target AUC spans 0.326 to 0.686 (Fig. C.4). Twelve of 20 directions are above chance with CI support. Six are *below* chance with CI support: both directions of Manavgat to Bejís and of Manavgat to Muğla, plus Bejís to Evia and Evia to Bejís. Two intervals span 0.5. Even the best raw transfer, Evia to Manavgat at 0.686, stays far below that target’s own within-region thermal performance of 0.870. Across all directions the raw deficit against the within-region reference is 0.184 to 0.592 AUC. That reference is a region-level blocked estimate; the controls below show it is not matched to a transfer evaluation.

Those counts belong to the 2-cell blocking of Table 7. At the more conservative 10-cell blocking the same points give 9 above, 4 below and 7 uncertain, and no direction changes side of the chance line (Section 4.7). Four of the six below-chance directions keep their support there. Bejís to Manavgat and Manavgat to Muğla lose it and carry no verdict. The qualitative statement is unchanged. The counts should not be read as exact.

The static baseline does not transfer either. This is the control for the reading the rest of the paper invites, and it constrains that reading sharply. The same twenty directions were run with the terrain, fuel and greenness baseline alone. The mean target AUC is **0.537**, against **0.541** for the thermal model. Static attributes of a place are the class Dimarco et al. transfer successfully, and are the class this paper’s framing treats as portable. Here that class is itself barely above chance. The paired per-direction contrast is given in Section 4.6b. What matters here is that the

transfer failure below is not specific to the dynamic block. A baseline that does not travel, plus pre-fire thermal state, gives a model that does not travel.

The thermal block’s paired contribution to transfer, with its interval. The two matrices were differenced direction by direction. The mean is **+0.004**. The directions are not independent, because each region appears in eight of the twenty. The interval therefore depends on what is treated as the resampling unit. All four units the design permits give the same answer:

Table 4:

Resampling unit	n	95 % interval on the mean paired contribution
Directions, naive	20	[−0.027, +0.034]
Unordered pairs, cluster bootstrap	10	[−0.028, +0.036]
Unordered pairs, t on pair means	10	[−0.034, +0.042]
Regions, leave-one-out jackknife	5	[−0.037, +0.046]

Every interval spans zero, and the point estimate is a small fraction of each width. The per-region jackknife shows how little the mean is anchored. Holding out Manavgat, Bejís, Muğla, Evia and Montiferru in turn gives +0.0148, +0.0021, +0.0065, **−0.0081** and +0.0060. **Dropping Evia alone reverses the sign of the headline.** None of these four units propagates within-direction sampling variability. They resample between directions only.

Where the skill is lost, on a matched comparison. Four evaluations are reported. The last three are scored on **identical cells**, so they differ only in what the model was trained on, and the first is included to show why an unmatched comparison misleads. Each uses the transfer protocol unchanged:

fit, then apply with no refit, no recalibration and no threshold selection.

630 The held-out unit is a burned connected component of at least 50 cells together with all cells within 2 km of it. That area is 34 to 87 % burned, against 3.8 to 28.7 % for a whole region, and every negative in it is fire-adjacent, so it is a harder discrimination problem than a region.

Table 5: The four evaluations, scored on identical cells. Primary natural-vegetation population. Rows B, C and D are scored on the held-out scar area; row A is the whole region and is shown to make the mismatch visible. Means and Student t intervals are over the eight held-out scars.

Evaluation	Model trained on	Scored on	Mean AUC	95 % CI
A. Blocked cross-validation, 5 km	the region, scar included	the whole region	0.782	[0.746, 0.818]
B. Same blocked model, restricted	the region, scar included	the scar area	0.634	[0.562, 0.706]
C. Leave-one-scar-out	the region, scar withheld	the scar area	0.552	[0.503, 0.601]
D. Foreign region	another region, 306 to 2,802 km	the scar area	0.555	[0.502, 0.608]

All four rows are means over the **same eight scars**, and the intervals
635 are Student t over those eight, which is the resampling unit for this arm rather than the spatial-block bootstrap used elsewhere in the paper. Eight is a small number and the intervals are wide accordingly. Four of the eight scars are in Muğla, so they are not fully independent.

Most of the apparent collapse is the evaluation area, not the fire.

640 A and B use the same model and differ only in which cells it is scored on, and that alone accounts for 0.155 of the 0.230 fall from A to C, about two thirds.

Any comparison of a scar-level result against a region-level reference inherits that, and the region-level reference is what a paper of this kind normally reports.

645 **Neither withholding the fire nor moving 2,800 km has a measurable cost.** B minus C, the fire-specific residual on identical cells, is **+0.082** **[−0.011, +0.175]**, which spans zero. C minus D, the effect of replacing a same-region model with one fitted between 306 and 2,802 km away, is **−0.003** **[−0.075, +0.069]**, which also spans zero and is an order of mag-
650 nitude smaller. The second is the cleaner result: on the same cells, a model fitted in another country does as well as one fitted in the same region with the scar withheld, and it survives restriction to Muğla, the only region where holding out one scar still leaves the source model properly trained, at 0.579 against 0.597.

655 The first difference should not be over-read. With eight scars this design cannot establish a fire-specific residual, only bound it: it is at most about 0.18, against the 0.148 that separates row A from row B on the same model.

What this establishes is that the failure is a property of **contiguous spatial holdout**. It is not a property of separation distance and not of crossing
660 a region boundary. What it does not establish is the fire event as the unit, because the held-out patch is defined by the labels, so its identity cannot be separated from its location. Two further observations point the same way: a model trained on the other half of the same fire is no better at matched sample size, and the within-region half-split, which does see half of the tar-
665 get scar, returns 0.574 against leave-one-scar-out’s 0.552 with overlapping intervals.

Separation does not order the twenty-direction matrix either. The Spearman correlation between centroid separation and transfer AUC is -0.32 with an interval spanning zero, and the two nearest directions, Manavgat and Muğla at 306 km, are among the worst at 0.470 and 0.401. This is a different quantity from Table 8’s geographic-distance row, computed on the twelve-direction subset, which gives -0.24 ; both span zero. The transfer mean of 0.541 is over all twenty directions and is not the value at any one separation: the 2,802 km pair returns 0.326 and 0.444.

The result is robust to how the held-out patch is defined, at 0.543 to 0.565 across minimum component sizes from 25 to 200 cells and under both 4- and 8-connectivity, and flat across buffers of 2, 5 and 10 km.

The increment declines with the holdout, and is not established once the fire is withheld. Contribution 1 is about the paired thermal-minus-baseline difference, so the same evaluations were run on it.

Table 6:

Evaluation	Thermal increment	Interval
Blocked cross-validation, per region	+0.056 to +0.153	every interval above zero
Within-region half-split	+0.027	positive in 13 of 18 splits
Leave-one-scar-out	+0.022	[−0.032, +0.077]
Cross-region, twenty directions	+0.004	[−0.028, +0.036]

The point estimate falls monotonically as the holdout hardens, and the last two intervals span zero. Under leave-one-scar-out the increment is positive in six of eight scars and averages +0.022, against a baseline of 0.529 and

a thermal 0.552 that are both close to chance on that evaluation; restricted
685 to Muğla it is +0.009 $[-0.051, +0.070]$. **The within-region increment of +0.056 to +0.153 is therefore substantially a property of interleaved holdout.** It is not shown to be zero on an unseen fire, and eight scars cannot show that; what this design establishes is that it is not established there.

690 **Six directions are below chance with interval support**, the sharpest at 0.326 $[0.305, 0.349]$. No account of merely lost skill produces a reliably reversed ranking, so those six need a mechanism acting on the direction of the relationship. That is the subject of Sections 4.4 to 4.8. Per-split and per-scar detail is in Appendix A(h).

695 **Label-blind adaptation compresses the matrix toward chance rather than repairing it.** Under region-wise z-scoring the twenty directions span 0.431 to 0.630 and under CORAL 0.443 to 0.624, roughly half the raw spread, with no adapted direction exceeding 0.631 against the unmatched within-region references of 0.859 to 0.918. Adaptation raises the
700 failing directions and degrades most of those that already transferred. Taking the better of the two adaptations per direction, 14 of the 20 end closer to chance than they began and 6 end further from it; five of those six involve Montiferru, the smallest and last-added region, and move upward, while the sixth is Manavgat→Muğla moving downward from 0.470 to 0.443. The 14 to
705 6 split should be read at the precision of limitation (ix) in Section 5.9, since Bejís→Manavgat is counted as compressed on a margin of 0.001.

Against the right reference, adaptation is not failing (Fig. C.5). The three controls above give an achievable reference for a model applied

Table 7: Cross-region transfer matrix, thermal model, TSG population. Target ROC-AUC with 2-cell spatial-block bootstrap 95% CIs (1000 replicates). CORAL is applied after region-wise z-scoring ($\lambda = 10^{-5}$).

Direction	Raw	Region-wise z-score	CORAL
Manavgat→Bejís	0.326 [0.305, 0.349]	0.477 [0.451, 0.502]	0.511 [0.484, 0.534]
Bejís→Manavgat	0.444 [0.408, 0.480]	0.457 [0.420, 0.497]	0.555 [0.528, 0.583]
Manavgat→Muğla	0.470 [0.451, 0.488]	0.431 [0.411, 0.449]	0.443 [0.423, 0.461]
Muğla→Manavgat	0.401 [0.378, 0.426]	0.559 [0.531, 0.587]	0.560 [0.535, 0.587]
Manavgat→Evia	0.613 [0.593, 0.631]	0.542 [0.520, 0.565]	0.539 [0.518, 0.561]
Evia→Manavgat	0.686 [0.653, 0.716]	0.516 [0.489, 0.544]	0.527 [0.500, 0.553]
Bejís→Muğla	0.618 [0.601, 0.635]	0.518 [0.501, 0.535]	0.507 [0.489, 0.524]
Muğla→Bejís	0.583 [0.561, 0.607]	0.535 [0.512, 0.557]	0.560 [0.538, 0.581]
Bejís→Evia	0.383 [0.363, 0.402]	0.532 [0.509, 0.551]	0.499 [0.479, 0.518]
Evia→Bejís	0.448 [0.426, 0.470]	0.549 [0.524, 0.575]	0.549 [0.525, 0.573]
Muğla→Evia	0.653 [0.636, 0.671]	0.561 [0.543, 0.580]	0.563 [0.545, 0.582]
Evia→Muğla	0.577 [0.560, 0.593]	0.501 [0.485, 0.518]	0.530 [0.515, 0.546]
Montiferru→Manavgat	0.567 [0.539, 0.594]	0.573 [0.538, 0.609]	0.606 [0.574, 0.639]
Manavgat→Montiferru	0.533 [0.488, 0.580]	0.586 [0.540, 0.629]	0.592 [0.550, 0.631]
Montiferru→Bejís	0.548 [0.521, 0.578]	0.574 [0.552, 0.596]	0.569 [0.548, 0.591]
Bejís→Montiferru	0.594 [0.560, 0.631]	0.550 [0.500, 0.601]	0.574 [0.530, 0.621]
Montiferru→Muğla	0.619 [0.604, 0.634]	0.576 [0.562, 0.589]	0.565 [0.550, 0.579]
Muğla→Montiferru	0.531 [0.495, 0.568]	0.587 [0.549, 0.624]	0.584 [0.547, 0.623]
Montiferru→Evia	0.586 [0.565, 0.606]	0.630 [0.611, 0.649]	0.624 [0.605, 0.641]
Evia→Montiferru	0.647 [0.608, 0.682]	0.568 [0.528, 0.609]	0.581 [0.539, 0.623]

Table 8: Transfer-gap decomposition (four-AOI set, 12 directions). Within = target’s within-region thermal AUC; best adapted = the better of z-score/CORAL; recovered fraction = (adapted – raw)/(within – raw), signed and unclipped, with paired bootstrap CI (1000 replicates). Montiferru directions are not part of this decomposition (per-pair absolute decompositions exist without fraction CIs). The status column asks whether the *adapted* value clears chance and uses the 2-cell adapted intervals of Table 7. The adapted arms were not recomputed at the coarser blocking of Appendix A(c), which covers the raw arm and the paired delta only.

Direction	Within	Raw	Best adapted (method)	Recovered fraction [CI]	Status
Manavgat→Bejís	0.918	0.326	0.511 (CORAL)	+0.31 [+0.28, +0.34]	recovery, chance not excluded
Bejís→Manavgat	0.870	0.444	0.555 (CORAL)	+0.26 [+0.18, +0.34]	recovery above chance
Muğla→Manavgat	0.870	0.401	0.560 (CORAL)	+0.34 [+0.28, +0.40]	recovery above chance
Bejís→Evia	0.912	0.383	0.532 (z-score)	+0.28 [+0.23, +0.32]	recovery above chance
Evia→Bejís	0.918	0.448	0.549 (z-score)	+0.22 [+0.16, +0.27]	recovery above chance
Manavgat→Muğla	0.859	0.470	0.443 (CORAL)	−0.07 [−0.12, −0.03]	negative recovery
Muğla→Bejís	0.918	0.583	0.560 (CORAL)	−0.07 [−0.16, +0.01]	negative recovery
Manavgat→Evia	0.912	0.613	0.542 (z-score)	−0.23 [−0.32, −0.14]	negative recovery
Evia→Manavgat	0.870	0.686	0.527 (CORAL)	−0.86 [−1.18, −0.60]	negative recovery
Evia→Muğla	0.859	0.577	0.530 (CORAL)	−0.17 [−0.22, −0.11]	negative recovery
Muğla→Evia	0.912	0.653	0.563 (CORAL)	−0.35 [−0.43, −0.27]	negative recovery
Bejís→Muğla	0.859	0.618	0.518 (z-score)	−0.42 [−0.51, −0.34]	negative recovery

to a fire it has not seen: 0.574 for the half-split, 0.552 for the leave-one-
710 scar-out. The best label-free adaptation averages **0.556** across the twenty
directions, against 0.541 raw. It is therefore at that reference, not far below
it. What it does is regress the matrix towards that value: fourteen of twenty
directions move closer to chance, which harms the directions that already
worked and helps the ones that did not. The verdict that follows is narrower
715 than "alignment fails". Alignment cannot exceed what a model can achieve
on an unseen fire, and it does not.

In the six directions where raw transfer was below chance, the best label-
free method recovers at most 34 % of the gap to the within-region reference,
so the remaining unrecovered fraction is at least 0.66 everywhere. The half-
720 split control above shows that the larger part of that remainder is already
incurred inside the region, so it should not be read as a measure of concept
shift. Seven directions show *negative* recovery, meaning adaptation moves
the score away from the reference; in six of those raw transfer was already
above chance and adaptation destroyed that advantage. Label-free alignment
725 therefore does not act as a repair mechanism.

4.4. *Transferability diagnostics: only conditional similarity orders transfer*

Twenty candidate diagnostics from four families were each rank-correlated
with the same target quantity, the raw thermal transfer AUC over the twenty
ordered directions, under one common pair-based bootstrap.

730 **Only two diagnostics have intervals excluding zero, and both
are conditional.** The stronger is the sign-agreement fraction over interval-
supported features, at $\rho = +0.84$ [+0.58, +0.88]. The cosine variant reaches
+0.81. No marginal measure was shown to order the matrix. That includes

Table 9: Transferability diagnostics versus raw thermal transfer, by family. Spearman ρ against raw transfer AUC with pair-based bootstrap 95 % CIs. Exp. is the sign expected if the diagnostic orders transfer. The member named is the one with the largest absolute correlation in its family, which is not always in the expected direction. All twenty individual diagnostics are in Appendix B, Table B1.

Family	Diagnostics	Exp.	Largest correlation in family	Sign as expected	Any CI excluding 0
P(y x) condi- tional	6	+	+0.84 [+0.58 , +0.88] agree- ment fraction, supported features	yes	yes, 2 of 6
P(x y=1) niche overlap	5	+	+0.24 [−0.45, +0.74] Schoener’s D, 1-D mean	yes	no
P(x) marginal	6	−	−0.32 [−0.78, +0.33] domain- classifier AUC	yes	no
P(y) regime structure	2	−	+0.29 [−0.38, +0.74] ^{log} effective- N distance	no	no
geographic	1	−	−0.24 [−0.84,	yes	no

area-of-applicability-style dissimilarity in predictor space, climatic distance
735 and geographic distance. The niche-overlap and regime families were not
shown to order it either. The learned domain classifier is at ceiling, separating
source from target at $AUC \geq 0.96$ for every pair. It always succeeds, which
is why it carries no ordering information.

Two limits are stated with the result rather than after it. The first is
740 size. The index’s tie structure caps the achievable Spearman at +0.861,
so the observed +0.840 sits essentially on that ceiling. Its exact one-sided
permutation p is 0.0060, the smallest this tie structure can produce, against
a Bonferroni threshold of 0.0026 over the nineteen computed variants. No
outcome of this diagnostic could have cleared family-wise correction on ten
745 effective pairs. The second limit is labels. Signed associations need burned
labels in both regions. The family that works is therefore not available before
deployment, while the family that fails is.

A third limit is about selection, and it is the sharpest of the three. The
two rows that clear zero are the *supported-feature* variants. Their feature
750 subset is chosen by whether two regions’ bootstrap intervals happen to be
disjoint. That is a data-dependent selection, made on the same data, with
no correction. The unselected counterparts over all nine features are $\rho =$
 $+0.50 [-0.17, +0.83]$ for the cosine and $\rho = +0.18 [-0.40, +0.72]$ for the
agreement count. Both span zero. The result lives in the selection step, and
755 is reported as such.

The families are compared on unequal samples, and equalis-
ing them does not change the ordering. The marginal, applicability,
climatic and geographic rows sit on twelve directions, because those diag-

nostics exist only for the four-region subset. The supported-conditional rows
 760 sit on sixteen and the rest on twenty. A reader may reasonably ask whether
 the marginal family’s failure is a statement about power rather than about
 diagnostics. Every row was therefore recomputed on the common twelve di-
 rections. The published values reproduce to 4.8×10^{-5} . The conditional
 rows still lead, at $+0.87 [+0.65, +0.88]$ and $+0.85 [+0.43, +0.88]$, and every
 765 marginal row still spans zero. The ordering is not an artefact of unequal
 samples.

4.5. *The contrast pair: similarity is not sufficient*

The clearest single view of Table 9 needs no ranking at all (Fig. C.8).
 Manavgat and Muğla lie in the same country and the same fire year. They
 770 are 306 km apart by the centroid geodesic distance this paper uses as a
 diagnostic, and their nearest boundaries are 191 km apart. Their burned cells
 occupy the most similar environmental envelope of any pair in the matrix,
 with per-feature Schoener’s D of 0.77 to 0.89. Yet five of nine feature-response
 directions point opposite ways. Elevation is one of them, at a signed AUC of
 775 0.374 [0.289, 0.471] in Manavgat against 0.611 [0.532, 0.690] in Muğla, with
 disjoint intervals. All four absolute thermal channels are among the others.
 Transfer is below chance in both directions at the point estimate, at 0.470
 and 0.401. Both directions also sit inside the nominal area of applicability,
 at 0.875 and 0.531 of target cells.

780 Bejís and Montiferru sit at the opposite extreme. Their burned envelopes
 barely overlap, and they carry the most dissimilar values on every overlap
 measure. They transfer above chance in both directions. That half of the
 contrast rests on point estimates: neither direction carries a verdict at 5

km blocking, and the marginal applicability audit was never produced for
 785 Montiferru. The claim is one of *sufficiency*. Similarity does not guarantee
 transfer, and dissimilarity does not preclude it. The claim rests on two
 coexisting counterexamples, so it does not depend on the number of pairs
 available.

4.6. Interventions: pooling and feature removal

790 **(a) Pooled multi-region training** (Fig. C.6). At the point estimate,
 training on the pooled primary populations of the other four regions never
 beats the best single-source transfer for any target, with shortfalls of 0.02 to
 0.22, and it stays 0.28 to 0.50 AUC below the within-region ceiling; for two
 targets it falls below the pairwise mean and below chance, though only Bejís
 795 is below chance with interval support, at 0.417 [0.369, 0.467]. Aggregation
 does not recover what single-source transfer loses.

(b) Removing the direction-reversing features. The two predictors
 whose signed association reverses between regions **with bootstrap sup-**
port are **elevation_mean** and **lst_anomaly_mean**. The per-region signed
 800 associations and the three pair-level reversals that meet that criterion are
 in Appendix B, Tables B2 and B3. Retraining without them costs **−0.081**
 of mean within-region AUC, supported in every region (per-region deltas
 −0.060, −0.130, −0.074, −0.063, −0.079, every interval entirely below zero),
 and changes mean transfer by **+0.014** [**−0.017**, **+0.045**], which spans zero.

805 Two qualifications belong with those numbers rather than after them.
 First, the debit is not the thermal block’s. Dropping elevation alone accounts
 for −0.061 of it, from 0.888 to 0.827. Dropping the LST anomaly alone
 accounts for −0.013, from 0.888 to 0.875. Roughly three quarters of the

cost is therefore the removal of a **baseline** terrain variable, not of a thermal
810 channel. Second, both figures are post-selection. The two predictors were
chosen because they reverse, using the same data on which the -0.081 and
the $+0.014$ are then estimated. No correction for that selection is applied.

Table 10:

Configuration	Mean within-region AUC	Mean transfer AUC
full	0.888	0.541
drop <code>elevation_mean</code>	0.827	0.546
drop <code>lst_anomaly_mean</code>	0.875	0.544
drop both	0.807	0.556

Fig. C.7 shows the four configurations together. What this measures is
a local cost with no compensating transfer gain. It does not measure an
815 exchange, because the transfer side is a null on both arms. The thermal
block’s own contribution to transfer is $+0.004$, with an interval spanning
zero (Section 4.3). Removing the reversing predictors returns $+0.014$, with
an interval spanning zero. Two nulls on the portability axis are not a price
paid.

820 4.7. Sensitivity analyses

Seven design choices were varied with everything else held fixed. They
are the Evia AOI and its prevalence, the CORAL regularisation constant, the
blocking scale, the closure date of the predictor window, the quality screen-
ing of the coarse thermal input, the contrast between the normalised and the
825 absolute dryness channels, and the removal of the coordinate-informed chan-
nels. None changes a conclusion above. Two bound how the results should

be read, so they are carried into the main text here.

Coarsening the blocks from 1 km to 5 km moves the verdict counts from ten positive, seven negative and three uncertain to six, four and ten. Support
830 is removed from seven verdicts and added to none. The point estimates are unchanged, but that is an identity rather than a result. The blocking scale is the resampling unit, and it cannot move an estimate computed once over all target cells.

Manavgat’s whole downstream chain was then rebuilt from a quality-
835 screened MODIS input. That changes the downscaled surface on 22,304 of 24,150 cells, by up to 10.9 °C. No signed univariate association moves by more than +0.0003. This closes the one processing-artefact candidate for that region’s behaviour. Appendix A reports all seven arms, including one that tests a claim of Section 1.2 and does not uphold it.

840 *4.8. The same geography, a second fire: reversal with place held constant*

Every result above compares different places. Muğla admits a stricter test, because a second fire burned inside the identical AOI on the identical analysis grid. The 2021 event has a predictor window closing on 28 July. The 2022 event has a matched window closing on 20 June. Region, bounding box, cell
845 definition, feature registry and processing chain are the same in both arms. The design holds place fixed. It does not hold season fixed, because the 2022 fire ignites about six weeks earlier, and it does not hold the population fixed, because the 2022 arm is defined by removing the 2021 scar.

Elevation reverses with bootstrap support. In 2021 higher ground
850 burned preferentially, at 0.611 with its interval entirely above 0.5; in 2022 lower ground did, at 0.296 with its interval entirely below. The two inter-

Table 11: Signed univariate feature-burned AUC, Muğla 2021 versus 2022. Raw AUC against **burned**, never folded to $\max(\text{AUC}, 1 - \text{AUC})$; 10-cell (≈ 5 km) spatial-block bootstrap, 1,000 replicates, seed 42 (Section 3.15). Analysis population 41,730 rows / 2,911 burned (2021) and 38,790 rows / 331 burned (2022). **Positive-carrying 5 km blocks: 70 for the 2021 arm and 11 for the 2022 arm.** Table 3’s note sets sixteen as the floor this design supports at that blocking, so the 2022 intervals here fall below the paper’s own standard and are read as indicative, exactly as the 20-cell row of Table 3 is. The 2022 arm is additionally a single compact scar, so its eleven blocks are contiguous. Read from `paper/step9g_raw/.../mugla_2021__mugla_2022_event_relative/step9g_direction_reversal_table.csv`.

Feature	2021 AUC [95 % CI]	2022 AUC [95 % CI]	Reversal
elevation_mean	0.611 [0.532, 0.690]	0.296 [0.230, 0.355]	bootstrap- supported
current_lst_mean	0.325 [0.271, 0.382]	0.515 [0.434, 0.580]	point only
current_tvdi_mean	0.336 [0.275, 0.398]	0.594 [0.475, 0.674]	point only
downscaled_lst_mean	0.307 [0.253, 0.366]	0.508 [0.435, 0.571]	point only
fused_lst_mean	0.325 [0.272, 0.383]	0.519 [0.436, 0.583]	point only
ndvi_mean	0.662 [0.616, 0.704]	0.707 [0.624, 0.777]	none
slope_mean	0.637 [0.582, 0.686]	0.558 [0.468, 0.634]	none
lst_anomaly_mean	0.485 [0.395, 0.566]	0.380 [0.249, 0.502]	none
tvdi_difference_mean	0.490 [0.396, 0.575]	0.397 [0.265, 0.514]	none

vals are disjoint, 0.532 against 0.355, and the difference is -0.317 $[-0.414, -0.220]$. The same predictor, the same region, the same grid, and an association that points the opposite way in two fires eleven months apart. The
855 four absolute thermal channels change side as well, from bootstrap-supported lower-values-burn in 2021 to higher-values-burn point estimates in 2022.

Two structural properties must be read alongside these numbers, because they have no analogue in the twenty-direction matrix. The two arms are not disjoint samples: they share 38,789 of the 2022 arm’s 38,790 cells, and
860 elevation, slope and land cover are identical to the digit across all 73,098 grid cells, so only NDVI and the six thermal channels carry new information between them. And the removal is the target’s own positive class, so in the 2022 to 2021 direction not one target positive is present in the source training population while 38,789 of 38,819 target negatives are, and membership of
865 the source training set alone separates the 2021 target’s classes at ROC-AUC 0.9996. We therefore cannot bound what that asymmetry does to a transfer estimate between the two arms, which is why none is reported here, only state that a direction with this structure is not comparable to the twenty between-region directions. The elevation reversal is the finding; the
870 structural asymmetry is a competing explanation this design cannot exclude. The known part of the bias runs the safe way: the removed cells are high, with a median of 563 m, and unburned in 2022, so removing them raises the 2022 elevation AUC and makes the reversal smaller rather than larger.

4.9. What target labels cost: the recovery curve

875 Everything above measures a failure; this prices it. Label-free alignment does not close the residual gap, so the missing resource is information about

the target that alignment cannot synthesise, and the direct way to supply it is target labels. A frozen labelled-budget diagnostic answers how many, for three regions across all six ordered directions, using one 10-cell (~ 5 km) spatial block as the unit of labelling effort and reading recovery against a
880 matched target-only ceiling of 0.777 to 0.824.

Thirty-two labelled blocks recover **85 to 89 % of that ceiling in three of the six directions**, two of which started below chance; 51 to 57 % in two more, the directions where the conditional gap is widest; and 30 % in the
885 sixth. That budget is 7 to 20 % of the target’s natural-vegetation population, about 800 km², so it is a real answer and not a cheap one. Two properties of the measurement bound it and are stated in the supplement rather than buried: at the top budget the selection pool is nearly exhausted, so the narrow upper-budget intervals reflect saturation rather than precision, and
890 the labelled blocks are drawn from the event being predicted, which is not a resource available before that event burns. The protocol, the full curve and the limits are in Supplement S1.

Two limits belong with the number. At small budgets the same intervention damages the direction that already transfers best without any labels.
895 And six directions in three regions cannot support a general label budget, so none is offered.

5. Discussion

5.1. Principal findings

Three results carry this paper. First, the pre-fire thermal block raises
900 spatially blocked within-region ROC-AUC by +0.056 to +0.153 in every one

of five Mediterranean regions, with bootstrap support at 1 km and 5 km blocking, and contributes $+0.004$ $[-0.028, +0.036]$ across twenty ordered transfer directions, indistinguishable from zero and with a sign that varies by pair. The static baseline transfers no better, at a mean of 0.537 against
905 0.541, so the failure is not a property of the dynamic block specifically. Nor is it attributable to region crossing. On identical cells, withholding the burn scar from training costs $+0.082$ $[-0.011, +0.175]$ and moving the training data 306 to 2,802 km away costs -0.003 $[-0.075, +0.069]$. The failure is a property of contiguous spatial holdout. The six below-chance directions
910 are a separate matter, since no account of merely lost skill produces a reliably reversed ranking. Removing the two reversing predictors, elevation and the LST anomaly, costs -0.081 of within-region skill, supported in every region and three quarters attributable to elevation, and changes transfer by $+0.014$ $[-0.017, +0.045]$, which also spans zero. Second, of twenty candidate
915 transferability diagnostics only two have intervals excluding zero and both are conditional, while no marginal measure was shown to order the matrix. The marginal family, which includes area-of-applicability dissimilarity, is the only one runnable before deployment. Third, the mechanism is a reversal in the sign of the predictor-burning association, and it persists inside one study
920 area across two fires.

5.2. *Why the thermal increment is real but local*

A distinction has to be drawn before this section can say anything useful. The within-region increment and the transfer failure are measured at different separations, and Section 4.3 shows that most of the difference between them
925 is already present inside a single region. What follows therefore reads the

increment as local in a specific sense: it holds where the held-out cells are interleaved with training cells, and most of it is gone once they are not, before the fire or the region changes. What that does not explain is why the directions differ from each other, and in particular why six of them are
930 anti-predictive.

The within-region increment is not an artefact to be explained away, but it is measured under a holdout that Section 4.3 shows is generous. Withholding a whole burn scar reduces it to $+0.022$ $[-0.032, +0.077]$, so what follows describes an increment that is established under interleaved validation and
935 not established beyond it. It replicates in five independent regions, survives coarsening of the spatial blocks to ~ 5 km with its interval intact, persists in the secondary all-valid population, and survives a predictor window closed 7 and 14 days earlier in all five regions. The direction of that last result is region-specific and should not be overstated: the contribution strengthens
940 in Bejís ($0.058 \rightarrow 0.079$ at 14 days) and Muğla ($0.115 \rightarrow 0.128$), is flat in Montiferru, and **weakens monotonically in Evia** ($0.156 \rightarrow 0.149 \rightarrow 0.135$). What holds everywhere is survival, not improvement. The transfer failure is therefore not evidence that the thermal signal is spurious. It is evidence that the fitted relationship is *local*.

945 The natural objection is that different Mediterranean regions are simply different systems, so that a predictor meaning one thing in one place and another elsewhere is not instability but a comparison of two different systems. The two Muğla events answer it directly: region, AOI, grid, feature registry and processing chain identical, static predictors identical cell by cell, and
950 elevation’s association with burning reversing from 0.611 $[0.532, 0.690]$ to

0.296 [0.230, 0.355], a difference of -0.317 [-0.414 , -0.220] (Section 4.8). Holding geography fixed does not stabilise the direction of the relationship.

Three caveats attach and none is dismissed. Season and year are confounded, and **that confound cannot be resolved in this study area**:
955 the events sit 42 days apart in median burn day-of-year, neither year contains a second event at the other’s phase, and a calendar-matched arm would carry nine burned cells against a gate minimum of thirty. The population is not held fixed either, the 2022 arm being the 2021 arm with the 2021 scar removed. And the 2022 arm rests on one compact scar and eleven positive-
960 carrying 5 km blocks. What survives all three is the narrower reading that *something* reverses the elevation-burning association with geography, grid and pipeline held fixed, which is why Section 4.8 is reported as corroboration rather than as the load-bearing evidence for the mechanism.

The sharpest supported reversal belongs to elevation, a static predictor,
965 so instability is a property of the predictor-to-burning mapping generally rather than of thermal channels specifically. This is the same conclusion the baseline transfer arm reaches from the other direction: the instability is not the thermal block’s peculiarity.

5.3. *Why the similarity-based diagnostics fail*

970 A diagnostic built on distance in predictor space asks one question: do the target’s predictor values look like the training data’s? That is not the question that matters when the failure is conditional. A target region can sit well inside the training envelope while the relationship between those predictors and burning points the other way. Manavgat to Muğla is exactly
975 that case, with 0.875 of target cells inside the weighted area of applicability

and transfer below chance.

The same holds for the niche-overlap and regime families. At the point estimates, the pair with the highest burned-niche overlap in the matrix fails in both directions while the pair with the lowest transfers in both, so
980 no monotone function of overlap can order the outcomes. The qualifier is not decorative: at the 5 km blocking this paper otherwise defends, neither Bejís↔Montiferru direction carries a verdict and Manavgat→Muğla loses its below-chance support, so this contrast is a statement about point estimates and is made as one. The domain classifier is at ceiling for every pair, which
985 is informative in itself: the regions are trivially separable in predictor space whether or not their transfer works, so separability carries no ordering information here.

5.4. *What the conditional diagnostic is, and what it is not*

The sign-agreement index reaches $\rho = +0.84$ [+0.58, +0.88] against ob-
990 served transfer. Three limits are reported with it rather than after it, and the sharpest is a selection: the two rows that clear zero are supported-feature variants whose subset is chosen on the same data, and their all-nine-feature counterparts span zero. Its tie structure caps the achievable Spearman at +0.861, so the observed value sits essentially on its own ceiling and the rank-
995 ing it supports is coarse. Its exact one-sided permutation p is 0.0060, the smallest that tie structure can produce, against a Bonferroni threshold of 0.0026 over nineteen computed variants. No outcome could have cleared family-wise correction on ten effective pairs. And it needs burned labels in both regions, which makes it a mechanism diagnosis rather than a pre-
1000 deployment screen.

That last limit is the substantive one, and it is why the paper’s practical conclusion is not "use this index instead". It is that the family of diagnostics that can be run before deployment is the family that fails here, and the information that would order transfer is the information a practitioner does
1005 not have when the decision is made. The honest statement is that these twenty candidates were **not shown to order transfer**, not that they are shown incapable of it.

5.5. What the two interventions do and do not show

Both interventions show the same shape. Pooling four regions never beats
1010 the best single-source transfer for any target and stays 0.28 to 0.50 below the within-region ceiling, so aggregation does not manufacture the missing conditional information. Removing the reversing predictors costs -0.081 of within-region skill with interval support in every region, and changes mean transfer by $+0.014$ $[-0.017, +0.045]$, which is not distinguishable from zero.

1015 That pair of numbers is easy to read as an exchange, and it is not one. The thermal block’s own contribution to transfer is $+0.004$, with an interval spanning zero. Removing the reversing predictors returns $+0.014$, with an interval spanning zero. Both arms are null on the portability axis. What the interventions measure is a local cost and no compensating transfer gain.
1020 That is not a conservation law, and it is not a rate at which local skill can be sold for portability. No such rate is estimated here.

5.6. The regime hypothesis, reported as it happened

A regime-structure explanation was stated in advance and the data confirmed the null. The regime-distance correlation has the wrong sign at the

1025 point estimate, the most regime-similar pair fails in both directions, and the
most regime-different pair transfers above chance. One mundane explanation
can be set aside: cutting Muğla, by far the largest population, to Manav-
gat’s cell count and separately to its positive count leaves Muğla’s transfer
behaviour inside the subsampling range in both roles. The error was in the
1030 hypothesised grouping, not in the data, and with ten pairs this cannot refute
regime typology [2] in general.

5.7. *The empirical contrast with Dimarco et al.*

Dimarco et al. [7] transfer successfully across a comparable Mediterranean
design and we do not, and the two results are not in conflict. Their predictors
1035 are attributes of a place, ours describe the state of a surface in one season,
which suggests that the relationship between domain similarity and transfer
success is predictor-class dependent.

That reading is one of at least two, and we cannot separate them here.
Their response variable is human-driven **ignition**; ours is burned **area**. Ig-
1040 nition likelihood is dominated by human access and activity, stable and
near-universal drivers, whereas burned extent is dominated by spread. A
predictor-class explanation and a response-variable explanation are there-
fore confounded in this comparison. Our own data speak against a simple
predictor-class reading in any case: the static baseline transfers at a mean
1045 of 0.537 here, so within this cohort the place-attribute class does not travel
either.

5.8. *Implications*

For practice, transfer skill has to be *measured* rather than inferred from similarity, and the pre-deployment diagnostics currently used for that inference did not order it here. Where a model must be moved, the resource that closes the gap is target labels: thirty-two labelled 5 km blocks recover 85 to 89 % of the target’s matched ceiling in three of six directions, and 30 to 57 % in the rest. That is a real answer and not a cheap one, being 7 to 20 % of the target’s natural-vegetation population, and at small budgets the same intervention damages the direction that already transfers best.

For method development, the results bound what unsupervised alignment can be asked to do. Two label-free methods applied carefully moved fourteen of twenty directions towards chance, and every direction they improved involves the smallest region. But their mean, 0.556, is at the reference a model can reach on an unseen scar at all (Section 4.3), so they are not failing far below an achievable target; they are regressing the matrix onto it, which costs the directions that already worked. What alignment cannot do is exceed that reference, and a sign reversal is not a distribution mismatch that realigning inputs would repair.

5.9. *Limitations*

The transfer failure is a finding, not a limitation. The limitations are the boundaries on how far it generalises.

(i) **No meteorological covariates** enter the models, so we cannot say how local skill and portability behave for a mixed thermal-plus-weather predictor set.

(ii) **The same-geography comparison covers one region only**, and even there year and seasonal phase are confounded, a confound that cannot be resolved in this study area for the reason given in Section 5.2. Its 331 burned cells also leave the thermal reversals unresolved at interval level,
1075 and the pair holds place fixed but not population. Those two arms are additionally the only transfer directions here computed by us rather than read from the pipeline author’s frozen export, with his unmodified code and the same pinned environment.

(iii) **All labels derive from a single burned-area product**, MCD64A1
1080 [9], whose omission and commission characteristics [3] bound every model evaluated here. No second burned-area product covers 2021 and 2022 at this resolution; the companion paper reports what an independent active-fire observation says about the omission concern.

(iv) **Evia remains the most imbalance-atypical population** even
1085 after the AOI extension, at a TSG prevalence of 0.287 against 0.038 to 0.072 in three of the others.

(v) **Each region contributes one fire season**, so regional concept shift is confounded with event meteorology, and distinguishing them requires multi-year labels.

(vi) **Cross-region point estimates carry an implementation tolerance**
1090 of roughly ± 0.02 to 0.03 across scikit-learn versions. All reported numbers are fixed to one verified version, but exact reproduction elsewhere requires the archived environment.

(vii) **The diagnostic correlations rest on an effective sample of**
1095 **ten region pairs**. Both the successes and the failures of Section 4.4 should

be read at that power.

(viii) **Manavgat’s atypical transfer behaviour remains unexplained.**

It is where the conditional diagnosis bites hardest and where feature removal recovers most. Two candidates have now been tested and neither survives.

1100 The first was its meteorology: the predictor window was not extreme. The second was the quality-screening difference in its coarse thermal input. It was the more worrying of the two, because the change it induces correlates with elevation at +0.615. It does not propagate. The region’s entire downstream chain was rebuilt from a quality-screened MODIS input. That changes the
1105 downscaled surface on 22,304 of 24,150 cells, by up to 10.9 °C. No signed univariate association moves by more than +0.0003. Elevation’s is identical in both arms, at 0.374, and the population is unchanged. The result is structural rather than fortunate. Elevation is a DEM variable the screening cannot touch, and fusion falls back on the MODIS-derived surface across
1110 only 2.14 percentage points of coverage. With one fire season per region the remaining candidates are not separable in this design.

(ix) **The interval-support counts are less stable than the point estimates behind them.**

Several verdicts sit within a thousandth of their reference value, and at 1 km blocking the published split of ten positive,
1115 seven negative and three uncertain turns on a lower bound of -0.00045 . The point estimates and the sign pattern are stable; the counts are not. Every sentence in this paper that leans on an exact count of supported directions should be read at that precision.

6. Conclusions

1120 Pre-fire thermal dryness adds a real and repeatable increment to burned-area discrimination. Across five Mediterranean regions, ROC-AUC was raised by +0.06 to +0.15 over a static and near-static terrain, fuel and greenness baseline. That baseline’s only time-varying member is the vegetation-index composite. The gain survived spatial blocking at about 5 km, the coarsest
1125 scale this design supports as an interval. It also survived a predictor window closed up to two weeks before the first labelled burning.

That skill is local. Paired per direction, the same predictor block adds nothing distinguishable from zero to cross-region transfer. The mean is +0.004 over twenty directions and its interval spans zero under all four
1130 between-direction resampling units we computed. Its contribution changes sign from one direction to another. The paired deltas run from -0.148 to $+0.132$, twelve positive and eight negative. At the conservative 5 km blocking, five to six directions are helped and three to four are harmed with interval support, with the rest carrying no verdict. The block is also the swing factor
1135 at the chance line.

Two controls fix the meaning of all this. The baseline arm transfers at a mean of 0.537 against the thermal model’s 0.541, so the static predictor class is not portable here either. And three evaluations on identical cells locate where it goes: 0.627 with the held-out burn scar in the training data, 0.552
1140 with it withheld, and 0.559 for a model fitted 306 to 2,802 km away. The honest summary is therefore not that these predictors fail to cross regions. **The failure is a property of contiguous spatial holdout:** the fire-specific residual is $+0.082$ $[-0.011, +0.175]$ and the 2,800 km costs -0.003

[−0.075, +0.069]. Two thirds of the apparent collapse from a region-wide
1145 0.782 is the evaluation area itself, which is 34 to 87 per cent burned against
3.8 to 28.7 for a region. Six directions are nonetheless anti-predictive with
interval support, which no account of merely lost skill explains. The one-
event-per-region design also means that a fire cannot be separated here from
the season and meteorology that produced it.

1150 Removing the two reversing predictors, elevation and the LST anomaly,
costs −0.081 of mean within-region skill, with interval support in every region
and roughly three quarters of it attributable to elevation, a baseline terrain
variable; it changes mean transfer by +0.014, an estimate whose interval
spans zero. A local cost is measured. No compensating transfer gain is, on
1155 either arm, so no exchange between the two is demonstrated.

The relationship itself is also unstable, and that is a separate finding. The
direction of the link between dryness and burning changes from one region
to another, with bootstrap support for two predictors. This is measured
on univariate associations and does not depend on any fitted model, so the
1160 distance result above neither establishes it nor removes it. What the two
together rule out is the comfortable reading in which regional concept shift
explains the whole transfer matrix. Most of its *level* is reached without
leaving a region. What varies around that level, including six anti-predictive
directions, is not explained by separation and is where the instability matters.

1165 None of the similarity diagnostics tested here ordered the transfer ma-
trix: predictor-space distance, domain separability, niche overlap and regime
structure all failed. On ten effective pairs those null results mean not shown
to order transfer, rather than shown not to. Label-blind adaptation com-

presses transfer towards chance rather than repairing it, though its mean sits
1170 at the reference a model can reach on an unseen scar, so it is regressing the
matrix onto that level rather than failing beneath it. Pooled multi-region
training does not escape it either.

The practical implication is a change in what is checked before a dynamic-
state fire model is transferred. The question is not whether the target region
1175 lies inside the source’s environmental envelope. In our matrix, the pair with
the highest envelope overlap failed in both directions, and the pair with the
lowest overlap transferred in both. Both statements are made at the point
estimate. The question is whether the signed feature-response directions
agree.

1180 That check requires a labelled probe in the target region, and so, on
inspection, do the niche and regime diagnostics. Only the marginal family
can be computed before any target label exists, and it is the family that
failed. How large the probe has to be is not established here, so no label
budget is claimed. A signed association with a usable interval may need
1185 fewer labels than a refitted model, but that was not tested. What is clear
is that the supervised recalibration that label-free alignment cannot deliver
is not cheap. Section 4.9 prices it for three regions: thirty-two labelled 5
km blocks recovered 85 to 89 % of the target’s matched ceiling in three of
six directions, 51 to 57 % in two more and 30 % in the sixth, and at its top
1190 budget the labelled set already holds most of one target region’s burned cells.
Where no target labels exist, the transfer performance of such models should
be treated as unknown rather than inferred from similarity.

For model builders, local skill and portability should be reported sepa-

1195 rarely rather than assumed to travel together. A predictor block worth a
 large within-region increment may contribute nothing distinguishable from
 zero across regions, with a sign that varies by pair, and reporting only within-
 region validation hides that completely. Three extensions are left for fu-
 ture work: temporal transfer, meteorological covariates, and physically nor-
 malised dryness variables. Each is a step towards the self-calibrating thermal
 1200 monitoring system that motivated this study.

Appendix A. Appendix A. Sensitivity analyses

Each arm below varies one design choice and leaves everything else fixed.
 None changes a conclusion in the main text. They are reported so that a
 reader can see which choices were tested, and what each was worth. Arm (f)
 1205 is the exception in one respect: it is not a robustness check but a test of a
 claim the introduction makes, and the claim is not upheld.

(a) **Evia AOI and prevalence.** The raw transfer arms were repeated
 with the legacy, high-prevalence Evia box. Every qualitative conclusion is
 unchanged. Thermal raw transfer AUCs move by up to 0.07, the largest
 1210 being Evia to Bejís at 0.378 to 0.448, and no direction changes side of the
 chance line.

(b) **CORAL regularisation, including the canonical $\lambda = 1$.** Over
 nine λ values from 0 to 10^{-1} on four directions, CORAL transfer AUC moves
 by at most 0.014 within any direction and 0.008 within the thermal family.
 1215 The canonical $\lambda = 1$ of the cited method lies outside that sweep and was
 computed separately on the same four directions. It changes the thermal
 mean from 0.519 to 0.522, a shift of +0.003 with a per-direction range of

−0.005 to +0.016, and **no direction crosses the chance line**. The $\lambda = 1$ values are 0.510, 0.444, 0.559 and 0.575 against 0.508, 0.445, 0.564 and 0.559 at $\lambda = 0.1$. No CORAL-dependent conclusion in this paper is sensitive to λ over the full range from 0 to 1.

(c) Blocking scale. Recomputing the transfer quantities at 10-cell (≈ 5 km) blocking from the frozen per-cell predictions widens the intervals and moves the verdict counts, from ten positive, seven negative and three uncertain at 1 km to six, four and ten at 5 km. Coarser blocking therefore removes support from seven verdicts and adds none. Across five seeds the counts run 5 to 6 positive, 3 to 4 negative and 10 to 11 uncertain, so they are not exact.

The point estimates are unchanged. That is an identity rather than a result, and it should not be offered as robustness. The blocking scale is the bootstrap *resampling unit*, and each point estimate is computed once over all target cells, so no choice of block size could have moved one. The comparison does establish two things. Every verdict that changes, changes towards "no verdict". And no direction crosses the chance line under the widened intervals. The counts are the fragile part of this paper. The sign pattern is not thereby shown to be robust. It is simply not tested by this variation.

(d) Predictor-window closure. The predictor window was closed 7 and 14 days earlier in all five regions. The thermal contribution stays positive and bootstrap-supported everywhere. The direction of the change is region-specific. It strengthens in Bejís, from 0.058 to 0.079 at 14 days, and in Muğla, from 0.115 to 0.128. It is flat in Montiferru. It weakens monotonically in Evia, from 0.156 to 0.149 to 0.135. What holds in every region is survival,

not improvement. That is the claim made in Section 5.2.

(e) **Quality screening of the coarse thermal input.** Two of the
1245 five regions' MODIS inputs are quality-screened and three are not. The
split follows export date rather than design, and it induces an elevation-
correlated change at the input, at $r = +0.615$ in Manavgat. Manavgat's
entire downstream chain was rebuilt from a quality-screened input. That
changes the downscaled surface on 22,304 of 24,150 cells, by up to 10.9 °C.
1250 No signed univariate association moves by more than +0.0003. Elevation
is identical in both arms, at 0.374, the population is unchanged, and the
within-region increment moves from $[+0.055, +0.079]$ to $[+0.054, +0.077]$.
The reason is structural. Elevation is a DEM variable the screening cannot
touch, and fusion falls back on the MODIS-derived surface across only 2.14
1255 percentage points of coverage. Details are in `paper/modis_qc_downstream_`
`propagation.md`.

(f) **Normalised against absolute dryness channels.** Section 1.2
argues that an internally normalised index should be less exposed to absolute-
temperature offsets between regions than raw land surface temperature. The
1260 thermal block contains both kinds, so the argument can be tested directly as
a feature-set contrast.

One thing should be said before the result, because it makes the outcome
less surprising than it might otherwise appear. The normalised channels are
the two that hold their direction in the same-geography two-event compar-
1265 ison of Section 4.8, where the absolute channels move. That is a statement
about one region across two fires. **Across regions it does not hold:** `1st_`
`anomaly_mean` is itself one of the two bootstrap-supported reversing predic-

tors, reversing between Bejís and Evia, which is why Section 4.6b drops it alongside elevation. Being internally normalised protects a channel against
1270 the offset between two seasons in one place. It does not, on this evidence, protect it against a change of place.

Three feature sets were run over all twenty directions and all five within-region folds, with the classifier, population, folds and bootstrap held fixed. The harness aborts unless its reference configuration lands on the frozen ex-
1275 ports. It did exactly: the maximum absolute difference between the reference arm and the frozen transfer AUCs is 0.000000 across all twenty directions.

Table A.12:

Feature set	Mean transfer AUC	Directions > 0.5	Mean within-region AUC
baseline only	0.5371	16	0.7896
baseline + normalised anomalies	0.5442	15	0.8551
baseline + absolute surface state	0.5479	13	0.8564
all ten features (reference)	0.5414	14	0.8883

The advantage is not found. Per direction, the normalised set minus the absolute set has a mean of -0.0037 . It is positive in 12 of 20 directions and runs from -0.099 to $+0.094$, so the spread is an order of magnitude larger

1280 than the difference. The normalised channels transfer very slightly worse
at the mean, and nothing here supports a design rule favouring anomaly-
referenced dryness for portability.

Within region the two sub-blocks are indistinguishable, at 0.8551 and
0.8564, and each recovers most of the gap between the baseline’s 0.7896 and
1285 the full model’s 0.8883. What separates them in sign stability does not
become transferable skill.

This arm strengthens the negative finding rather than softening it. Four
differently constituted feature sets were tried. None clears a mean of 0.548
across regions, and all differ sharply within them. The baseline arm also
1290 gives an independent confirmation of the control reported in Section 4.3: its
mean transfer of 0.5371 was recomputed here from the modelling datasets
and matches the 0.537 read from the frozen per-direction export. Details are
in `paper/anomaly_only_transfer.md`.

The comparison is between two sub-blocks of one thermal set on one
1295 cohort. It does not test normalised dryness indices in general, and it does
not test a normalisation fitted against a pooled multi-region reference rather
than each region’s own baseline years.

(g) **The coordinate-informed channels.** `downscaled_1st_mean` and
`fused_1st_mean` come from a per-region downscaling model whose own in-
1300 puts include coordinates, which Section 3.14 names as the one route by which
a coordinate-derived surface re-enters a feature set that excludes coordinates.
A coordinate-smoothed surface is by construction locally informative and
non-portable, so if the within-region increment depended on it, that incre-
ment would be this paper’s own result in miniature rather than a finding

1305 about thermal dryness. The arm was therefore run here rather than de-
ferred.

Both channels were dropped and everything refitted, over all twenty direc-
tions and all five within-region folds. The reference configuration reproduces
the frozen exports exactly, at 0.000000 against both the transfer and the
1310 within-region references.

Table A.13:

Region	Increment, full set	Increment without the two channels	Retained
Manavgat 2021	+0.067	+0.063	94 %
Bejís 2022	+0.056	+0.046	82 %
Muğla 2021	+0.116	+0.097	84 %
North Evia 2021	+0.153	+0.145	94 %
Montiferru 2021	+0.101	+0.105	103 %
Mean	+0.099	+0.091	92 %

The increment does not depend on them. It stays positive in every
region and retains 82 % to 103 % of its size, rising slightly in one region. Mean
transfer is likewise unchanged, at 0.545 without the two channels against
0.541 with them and 0.537 for the baseline alone. The local skill this paper
1315 reports is therefore not an artefact of a coordinate-smoothed surface.

This arm was computed independently here and reproduces the compan-
ion paper’s figure for the same quantity, 82 % to 103 %, from a separately
written harness.

Appendix A.1. A(h). The four evaluations of Section 4.3, in full

1320 Section 4.3 reports four evaluations of the same models as a ladder. The per-split and per-scar detail is here, so that the ladder can be checked without leaving the manuscript.

Table A1. Within-region half-split, every split. Source and target positive counts are given because they are unequal, which is the principal
1325 limit on this arm: a straight cut does not produce two exchangeable halves. Two splits are unusable because one half of Manavgat contains no burned cells.

Table A.14:

Region	Axis	Direction	Source positives	Target positives	Thermal AUC	Baseline AUC
manavgat 2021	east-west	low to high	700	84	0.695	0.661
manavgat 2021	east-west	high to low	84	700	0.557	0.580
manavgat 2021	north-south	low to high	0	784	—	—
manavgat 2021	north-south	high to low	784	0	—	—
bejis 2022	east-west	low to high	640	460	0.750	0.784
bejis 2022	east-west	high to low	460	640	0.578	0.560
bejis 2022	north-south	low to high	357	743	0.713	0.662
bejis 2022	north-south	high to low	743	357	0.603	0.613
mugla 2021	east-west	low to high	1,449	1,462	0.448	0.378
mugla 2021	east-west	high to low	1,462	1,449	0.636	0.628
mugla 2021	north-south	low to high	774	2,137	0.533	0.503
mugla 2021	north-south	high to low	2,137	774	0.294	0.549
evia 2021 extended	east-west	low to high	2,052	612	0.589	0.506
evia 2021 extended	east-west	high to low	612	2,052	0.645	0.482
evia 2021 extended	north-south	low to high	2,564	100	0.649	0.513
evia 2021 extended	north-south	high to low	100	2,564	0.533	0.495
montiferru 2021	east-west	low to high	410	129	0.475	0.461
montiferru 2021	east-west	high to low	129	410	0.507	0.513
montiferru 2021	north-south	low to high	438	101	0.583	0.511
montiferru 2021	north-south	high to low	101	438	0.536	0.447

Table A2. Leave-one-scar-out at a 2 km buffer, every scar. Only

1330 Muğla contains more than one burned component of at least 50 cells, so it is the only region where holding one out leaves the source model properly trained. Its four arms mean 0.579; the four arms in the other regions mean 0.525, and the pooled figure of 0.552 averages the two.

Table A.15:

Region	Component	Source positives left	Target positives	Target cells	AUC
evia 2021 extended	1	11	2,653	3,059	0.465
manavgat 2021	1	88	696	1,135	0.592
montiferru 2021	1	97	442	758	0.584
montiferru 2021	5	472	67	195	0.458
mugla 2021	6	1,997	914	1,266	0.595
mugla 2021	1	2,173	738	1,244	0.561
mugla 2021	10	2,272	639	940	0.617
mugla 2021	8	2,363	548	954	0.542

1335 The evaluation populations of the two arms are not comparable with each other or with the transfer targets. A held-out scar with its 2 km collar runs at 34 to 87 % burned prevalence; a whole target region runs at 3.8 to 28.7 %. The scar arm therefore asks for discrimination against the nearest and most similar negatives only, while a transfer arm includes the whole easy far field. Section 4.3 states the consequence: the last three rows of the ladder are not distinguishable by this design.

1340 **Appendix B. Appendix B. All twenty transferability diagnostics**

Section 4.4 summarises this table by family. The full ranking is given here because the negative result is the point: a reader should be able to see every candidate that was tried, not only the families' best members, and to check that no diagnostic was dropped after it failed.

1345 **Table B1. All transferability diagnostics versus raw thermal transfer (20 ordered directions).** Spearman ρ with pair-based bootstrap 95 % CIs. Exp. = expected sign. Rows with $n = 12$ exist only for the four-AOI subset, because those diagnostics were never produced for Montiferru. The supported-features conditional rows use the 16 directions (8 pairs) with
1350 at least one CI-supported feature.

Table note. A null row means the diagnostic was **not shown to order transfer** on this design, not that it was shown incapable of ordering it. With ten effective region pairs the power is low, and the intervals are wide enough to admit moderate true correlations in either direction. The last row
1355 is retained rather than deleted because it was computed: on two directions the statistic has no meaningful value, and reporting that is more honest than dropping the variant.

Reading the two rows that clear zero. Both are supported-feature variants, where the predictor subset is chosen by whether two regions' bootstrap intervals happen to be disjoint. That is a data-dependent selection
1360 made on the same data, uncorrected. Their all-nine-feature counterparts are in the table and both span zero. The result lives in the selection step, and Section 4.4 says so.

Equal-sample check. The families sit on unequal samples: marginal,

Table B.16:

Diagnostic	Family	Exp.	n dir	Spearman ρ [95% CI]	CI excludes 0
Agreement fraction, supported features	P(y x) conditional	+	16	+0.84 [+0.58, +0.88]	yes
Cosine, supported features	P(y x) conditional	+	16	+0.81 [+0.33, +0.88]	yes
Cosine, all 9 features	P(y x) conditional	+	20	+0.50 [−0.17, +0.83]	no
Vector Spearman, all 9	P(y x) conditional	+	20	+0.27 [−0.36, +0.77]	no
Agreement count, all 9	P(y x) conditional	+	20	+0.18 [−0.40, +0.72]	no
Schoener’s D, 1-D mean	P(x y=1) niche	+	20	+0.24 [−0.45, +0.74]	no
Warren’s I, 1-D mean	P(x y=1) niche	+	20	+0.22 [−0.42, +0.73]	no
Schoener’s D, PCA-2D	P(x y=1) niche	+	20	+0.10 [−0.51, +0.68]	no
Warren’s I, PCA-2D	P(x y=1) niche	+	20	−0.07 [−0.66, +0.49]	no
Mahalanobis, burned centroids	P(x y=1) niche	−	20	−0.23 [−0.75, +0.44]	no
Domain-classifier AUC	P(x) marginal	−	20	−0.32 [−0.78, +0.33]	no
Predictor-space mean dissimilarity	P(x) marginal	−	12	−0.10 [−0.54, +0.43]	no
Predictor-space p95 dissimilarity	P(x) marginal	−	12	−0.08 [−0.59, +0.49]	no
Fraction inside weighted AoA	P(x) marginal	+	12	+0.22 [−0.48, +0.59]	no
Fraction inside	P(x) marginal	+	12	+0.08 [−0.89, +0.43]	no

1365 applicability, climatic and geographic rows on twelve directions, the supported-
conditional rows on sixteen, the rest on twenty. Recomputing every row on
the common twelve directions reproduces the published values to 4.8×10^{-5}
and leaves the ordering unchanged — the conditional rows still lead at +0.87
[+0.65, +0.88] and +0.85 [+0.43, +0.88], every marginal row still spans zero.
1370 Source: `paper/diagnostics_common_subset.md`.

*Appendix B.1. B2. The signed univariate associations the reversal claim
rests on*

Section 4.6b drops two predictors because their signed association with
burning reverses between regions with bootstrap support, and Contribution
1375 3 names the reversal as the mechanism. This is the evidence for both. It was
computed for the frozen analysis and is reproduced here because the claim
is otherwise asserted rather than shown.

**Table B2. Signed univariate AUC of each predictor against
burned, by region.** Primary natural-vegetation population; 10-cell (~ 5
1380 km) spatial-block bootstrap, 1000 replicates, seed 42. The AUC is never
folded to $\max(\text{AUC}, 1 - \text{AUC})$, so a value below 0.5 means lower values
rank burned and is a direction rather than weakness. **Bold** marks a region
whose own interval excludes 0.5.

What counts as a reversal. A pair of regions is called a reversal
1385 only when their signed associations point to opposite sides of 0.5 **and each
region’s own interval excludes 0.5**. That is stricter than requiring the
two regions’ intervals to be disjoint, and the difference matters: for `current_
1st_mean` between Manavgat and Muğla the two intervals are disjoint, but
Manavgat’s own interval, [0.452, 0.621], includes 0.5, so Manavgat has no

Table B.17:

Feature	Manavgat	Bejís	Muğla	Evia	Montiferru
elevation_	0.374 [0.289,	0.643 [0.558,	0.611 [0.532,	0.541 [0.448,	0.584 [0.395,
mean	0.471]	0.729]	0.690]	0.626]	0.762]
slope_mean	0.531 [0.423,	0.521 [0.439,	0.637 [0.582,	0.487 [0.418,	0.652 [0.506,
	0.642]	0.605]	0.686]	0.554]	0.771]
ndvi_mean	0.636 [0.587,	0.559 [0.497,	0.662 [0.616,	0.639 [0.575,	0.586 [0.450,
	0.676]	0.619]	0.704]	0.701]	0.704]
lst_anomaly_	0.482 [0.428,	0.418 [0.364,	0.485 [0.395,	0.640 [0.567,	0.395 [0.285,
mean	0.530]	0.480]	0.566]	0.710]	0.535]
current_lst_	0.538 [0.452,	0.477 [0.401,	0.325 [0.271,	0.377 [0.301,	0.370 [0.248,
mean	0.621]	0.547]	0.382]	0.456]	0.513]
current_	0.552 [0.460,	0.517 [0.429,	0.336 [0.275,	0.362 [0.285,	0.356 [0.233,
tvdi_mean	0.641]	0.595]	0.398]	0.442]	0.499]
tvdi_	0.449 [0.391,	0.512 [0.443,	0.490 [0.396,	0.519 [0.444,	0.378 [0.282,
difference_	0.505]	0.583]	0.575]	0.589]	0.497]
mean					
downscaled_	0.552 [0.466,	0.484 [0.400,	0.307 [0.253,	0.377 [0.297,	0.365 [0.240,
lst_mean	0.637]	0.560]	0.366]	0.459]	0.511]
fused_lst_	0.540 [0.454,	0.481 [0.404,	0.325 [0.272,	0.376 [0.300,	0.370 [0.248,
mean	0.622]	0.551]	0.383]	0.456]	0.513]

1390 established direction to reverse from. That pair is a point reversal, not a supported one.

Table B3. The cross-region reversals that meet the stricter criterion. Difference intervals are from the same paired bootstrap.

Table B.18:					
Feature	Region A	AUC	Region B	AUC	Difference [95 % CI]
elevation_mean	Manavgat	0.374	Bejís	0.643	+0.269 [+0.138, +0.397]
elevation_mean	Manavgat	0.374	Muğla	0.611	+0.235 [+0.102, +0.360]
1st_anomaly_mean	Bejís	0.418	Evia	0.640	+0.221 [+0.123, +0.313]

1395 Three pair-level reversals across **two** features, which is why Section 3.12 removes exactly those two. Twenty-nine further pairs reverse at the point estimate only, spread across eight of the nine features, and they are not counted. The conservative criterion costs the paper findings rather than manufacturing them: a difference interval on the pair, which is the instrument Section 4.8 uses, would support more reversals than the three listed here.

1400 **Appendix C. Appendix C. Protocol detail**

These protocols are given here in full rather than in Methods, because each is a specification a reader needs only when checking the corresponding result, and none is needed to follow the argument. Sections 3.2, 3.3, 3.11 and 3.15 state in summary what each does.

1405 *Appendix C.1. C.1 Burned-area label, the reconstructed analysis grid, and
the admissibility gate*

Labels come from MODIS MCD64A1 Collection 6.1 (MODIS/061/MCD64A1) retrieved through Google Earth Engine [10], whose omission and commission characteristics [3] bound every model here. The product is exported onto
1410 the 30 m EPSG:4326 reference grid used by the rest of the pipeline, which duplicates each native ~ 500 m observation across the 30 m pixels beneath it.

The analysis grid is therefore **reconstructed rather than native**: the 30 m grid is partitioned into non-overlapping blocks of $\text{round}(500 / 30) =$
1415 17×17 pixels, giving a nominal cell edge of 510 m, and each cell is identified by its integer block indices `row_500m` and `col_500m`. This approximates the MODIS sinusoidal cell rather than reproducing it, being anchored to the Landsat/EPSC:4326 reference grid, and marginal blocks are truncated. The cell is square in degrees but not on the ground, about 510 m north to south
1420 and 390 to 407 m east to west, so its area is 0.199 to 0.208 km² against a MODIS cell's 0.250 km²; every block-size label in this paper is therefore the north-south dimension. The companion paper develops what follows from that.

A cell's representative burn date is the **mode** of the positive sub-pixel
1425 day-of-year values in the block, tested against the region's label window. Because the exported raster carries no out-of-window positives, in this dataset a single in-window positive sub-pixel makes the cell burned. The fraction of positive sub-pixels agreeing with the modal date is recorded as `burn_date_pixel_agreement_fraction` but no agreement threshold is imposed. The

1430 label never affects a cell’s eligibility for modelling: unburned, all-no-data
and out-of-window cells all remain as the negative class.

Two safeguards are recorded rather than assumed. Cells that burned
inside a region’s own predictor window are removed from its analysis universe
rather than counted as unburned, which ran for Muğla, Evia and Montiferru,
1435 is recorded as not run for Manavgat and has no recorded status for Bejís.
Burning in earlier years is screened for no region in the five-region cohort;
the only historical exclusion in the study removes the 2021 Muğla scar from
the 2022 event-relative experiment of Section 3.15.

Before any modelling each region passes a gate that answers one ques-
1440 tion: of the cells labelled burned, what fraction is dominated by natural
vegetation? A region is admitted as a wildfire candidate when that fraction
reaches 0.50 and at least 30 burned cells are present, and is rejected as a
cropland-dominated control when the cropland fraction reaches 0.50 instead.
The gate uses ESA WorldCover classes aggregated to the same reconstructed
1445 cells. Its purpose is to separate burned area produced by natural-fuel com-
bustion from burned area produced by post-harvest stubble burning, which
MCD64A1 does not distinguish. Verdicts are reported in Section 4.1.

Appendix C.2. C.2 Transferability diagnostics versus transfer

Twenty candidate diagnostics from four families are each rank-correlated
1450 against the same target quantity, the raw thermal transfer AUC over the
twenty ordered directions, under one common pair-based bootstrap. The
families are marginal predictor-distribution measures $P(x)$, burned-niche over-
lap $P(x|y=1)$, fire-regime label-pattern structure $P(y)$, and conditional feature-
response direction $P(y|x)$.

1455 The **marginal** family includes area-of-applicability-style dissimilarity in scaled, importance-weighted predictor space [21, 22, 16], climatic and geographic distance, and a learned domain classifier trained to separate source from target cells. The **niche-overlap** family includes Schoener’s D [34] and Warren’s I [41] over the burned cells of each region. The **regime** family in-
1460 cludes burned-patch component counts and effective component counts. The **conditional** family is the sign-agreement index: the fraction of predictors whose signed univariate association points the same way in both regions, computed over all predictors and over the subset whose associations are individually interval-supported.

1465 Two properties of this design are stated in advance because they bound what it can show. The candidate set was fixed in a project analysis log before the correlations were computed; that log is not a formal pre-registration and no independent timestamped registration exists. And with five regions the effective sample is ten unordered pairs, so nineteen computed variants are
1470 evaluated at that power with no family-wise error control claimed; where a diagnostic clears an interval test but not a Bonferroni threshold, both are reported.

Appendix C.3. C.3 Same-geography event-to-event comparison (Muğla 2021 versus 2022)

1475 Muğla is the one region for which a second fire event is analysed on the identical AOI and analysis grid, which allows the direction of feature-label associations to be compared with geography held fixed. The 2022 experiment is registered with windows anchored to its own event: a 58-day predictor window closing the day before ignition and a 49-day label window opening

1480 on it, matching the 2021 durations exactly.

This is not a clean temporal-transfer design and is not presented as one. The 2022 event ignites about six weeks earlier in the season than the 2021 event, so calendar year and seasonal phase are confounded and no difference can be attributed to the year alone. The registered claim is same-
1485 geography event-to-event. The pair is deliberately kept out of the twenty-direction matrix, and that exclusion is enforced in released code rather than merely asserted here.

Two further properties bound the comparison. A historical-burn exclusion applies to the 2022 arm alone, masking every cell that burned in 2021 so the
1490 previous year’s scar does not enter the following year’s analysis: it removes 3,073 cells, of which 2,941 belong to the primary population, taking it from 41,730 rows to 38,790, a drop of 7.0 %. The two arms therefore share 38,789 of 38,790 cells and three byte-identical static predictors, and in the 2022-to-2021 direction every target positive lies outside the source training population.
1495 Section 4.8 states what follows for how these two numbers may be read.

CRediT authorship contribution statement

Yunus Emre Cogurcu: Conceptualization, Methodology, Formal analysis, Investigation, Writing – original draft, Writing – review and editing, Supervision. **Emrehan Metin:** Software, Data curation, Investigation, Val-
1500 idation, Writing – review and editing.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

1505 Funding

This work is an output of project 17506, supported by the Scientific Research Projects (BAP) unit of Çukurova University.

Data availability

1510 All satellite inputs are public and are obtained through Google Earth Engine. The processing pipeline, its configuration and the frozen numeric outputs are publicly available; the repository, the commit of record and the licence are given in the data and code availability statement of the Methods, together with the three components that fall outside that release. No digital object identifier is minted and no archival deposit exists.

1515 References

- [1] Alkan Akıncı, H., Akıncı, H., 2023. Machine learning based forest fire susceptibility assessment of Manavgat district (Antalya), Turkey. *Earth Science Informatics* 16, 397–414. doi:10.1007/s12145-023-00953-5.
- [2] Archibald, S., Lehmann, C.E.R., Gómez-Dans, J.L., Bradstock, R.A.,
1520 2013. Defining pyromes and global syndromes of fire regimes. *Proceedings of the National Academy of Sciences* 110, 6442–6447. doi:10.1073/pnas.1211466110.




Figure C.1: **Study regions and the gate negative control.** Locations of the five wild-fire regions analysed in this study (blue, circular markers): Manavgat 2021 and Muğla 2021 (Türkiye), Bejís 2022 (Spain), North Evia 2021 (Greece, extended AOI) and Montiferru 2021 (Sardinia, Italy). Rectangles are the true analysis areas of interest; the ring marks each centroid, because at basin scale the rectangles alone are hard to locate. Kozan 2023 (grey, dashed outline, square marker) is *not* a study region: it is the negative control used to calibrate the burned-landcover admissibility gate, and it enters no model, transfer or diagnostic anywhere in this paper. Its rectangle is the true extent of its processed AOI and is larger than the study AOIs for that reason alone; the difference in size carries no methodological meaning and is not a design choice. Coastlines are Natural Earth 1:50 m; country borders are omitted deliberately. The scale bar is exact only at 38°N. Region bounding boxes match Table 1 to 5×10^{-3} degrees.




Figure C.2: Processing pipeline and evaluation programme. Flow from source products (left) to the transfer and diagnostic programme (right); every box carries the Methods subsection that specifies it. The five source products (MCD64A1 burned area, Landsat Collection 2 Level-2, MOD11A1, Copernicus DEM GLO-30 and ESA WorldCover) merge onto the reconstructed ~ 510 m analysis grid (17×17 pixels of the 30 m reference grid). The burned-landcover admissibility gate has two outcomes: regions that *pass* proceed to the natural-vegetation primary population, while the rejecting branch (grey, dashed) terminates in the Kozan 2023 negative control, whose burned area is 98% cropland and which is never modelled. The dashed connector deliberately carries no arrowhead onward, so the control cannot be read as entering the pipeline. Feature sets feed two evaluations in parallel: within-region (random forest under spatial-block cross-validation) and cross-region transfer (source-only, 20 ordered directions, with label-blind z-score and CORAL adaptation). Their levels are combined in the transfer-gap decomposition (§3.12). The footer states what applies throughout: leakage hard-exclusion, seed 42, 1000-replicate spatial-block bootstrap (§3.13), and the sensitivity designs (§3.16). Section references are verified against the Methods heading list when the figure is built.




Figure C.3: **The thermal increment replicates in all five regions and survives coarser spatial blocking.** Natural-vegetation primary population; random forest with the primary weighted configuration; all intervals are 95% spatial-block bootstrap intervals (1000 replicates, seed 42). **(a)** Baseline versus thermal out-of-fold ROC-AUC per region at the default 2-cell (~ 1 km) blocking: the open marker is the static baseline feature set, the filled marker is the baseline plus the six thermal predictors, and the connecting bar is the increment. The right-hand column reports the paired ΔAUC and its 95% interval; it is a fixed column placed beyond the plotted AUC range, so no label position depends on a data value. Every interval excludes zero. **(b)** Absolute thermal ROC-AUC against spatial block size (2, 10 and 20 cells, $\approx 1, 5$ and 10 km): absolute skill falls as blocking coarsens, in every region. **(c)** The thermal increment ΔAUC against the same block sizes, with 95% intervals: the increment does *not* melt away, and its interval excludes zero in all ten of the 2- and 10-cell region-block-size combinations. The five 20-cell intervals also lie above zero but rest on 6 to 33 positive-carrying blocks each, so they are plotted as indicative rather than as intervals (Table 3, table note). In (c) the five regions are offset horizontally by 0.15 cell so their whiskers separate; **this offset is cosmetic, since all five regions are evaluated at exactly the same three block sizes** (2, 10, 20 cells), and the offset is smaller than half the spacing between block sizes, so the groups cannot merge. Colour and marker shape both encode region, so the panels remain readable in greyscale.




Figure C.4: **Cross-region transfer, and what label-blind adaptation does to it.** Thermal transfer ROC-AUC for all 20 ordered source–target directions among the five regions, natural-vegetation primary population. **(a)** Raw transfer, no adaptation. **(b)** After region-wise z-score standardisation. **(c)** After CORAL. Rows are the source region, columns the target; the grey diagonal is within-region performance and is *not* a transfer result. The diverging colour scale is centred exactly on chance (0.5, marked by the black rule on the colour bar), so orange is below chance and purple above. Each panel header states that panel's range: adaptation compresses the spread from 0.326–0.686 (raw) to 0.431 to 0.630 (z-score) and 0.443 to 0.624 (CORAL). It pulls the extremes toward chance rather than lifting transfer. Cells below chance are hatched. The hatch is a redundant cue for print: a diverging palette is symmetric in luminance about its centre, so in greyscale a below-chance cell would otherwise be indistinguishable from an equally distant above-chance one. Hatching follows the *exact* value while the printed label is rounded to two decimals, so two cells print “0.50” with different hatching: CORAL Bejís→Evia is 0.4991 (hatched, below chance) and z-score Evia→Muğla is 0.5013 (un-hatched, above chance). Every cell carries its value, and every label's colour is chosen by measured WCAG contrast against that cell's rendered colour (worst case 4.7:1). Confidence intervals for these directions are given in Table 4.




Figure C.5: **Label-blind adaptation regresses transfer onto the level a model reaches on an unseen scar, rather than recovering it.** All 12 ordered directions among the four regions carrying the full decomposition (Manavgat, Bejís, Muğla, Evia), natural-vegetation primary population. Each row is one direction. The open circle is raw transfer ROC-AUC and the arrow runs to the best label-blind adapted result for that direction (region-wise z-score or CORAL, whichever is better); the grey vertical tick is the within-region reference for that target, and the dashed vertical rule is chance. **Solid blue arrows are the five directions where adaptation moves transfer toward the within-region reference (recovery); dashed orange arrows are the seven where it moves *away* from it (negative recovery).** Line style and colour both encode this class, because the two hues differ by only 2.30:1 in relative luminance and would nearly merge in greyscale. The pattern is the point: all six directions starting above chance are pushed down, and five of the six starting below chance are pushed up, so adaptation acts on the spread rather than on skill. The exception is Manavgat→Muğla, which starts below chance at 0.470 and is pushed further down to 0.443. The worst case is Evia→Manavgat, where raw transfer of 0.686 is dragged down to 0.527 against a within-region reference of 0.870, a recovered fraction of -0.862 . No arrow reaches its within-region tick.




Figure C.6: **Pooling every other region never beats the best single source.** Leave-one-region-out evaluation: for each held-out target region (rows), a model trained on the pooled remaining four regions (filled orange circle) is compared against the best individual source region for that same target (open blue diamond), with the within-region ceiling as a grey vertical tick and chance as the dashed rule. Natural-vegetation primary population, thermal feature set, no adaptation. The pooled model lies to the left of the best pairwise source in all five targets without exception. Adding source regions does not accumulate transferable signal. In two targets the pooled model falls *below chance* at the point estimate (Manavgat 0.469, Bejís 0.417), and only for Bejís does the interval lie entirely below chance, while a single well-chosen source reaches 0.686 and 0.583 respectively. Every point remains far below the within-region ceiling.




Figure C.7: **Removing the direction-reversing features costs locally and returns nothing measurable on transfer.** Mean within-region ROC-AUC across the five regions (horizontal) against mean transfer ROC-AUC across the 20 ordered directions (vertical), for four feature configurations: the full thermal set, the set without the LST anomaly, without elevation, and without both. Each configuration has its own marker and is named in the key; the grey path connects them in that order. Moving left along the path, transfer rises monotonically ($0.541 \rightarrow 0.556$) while within-region skill falls monotonically ($0.888 \rightarrow 0.807$). Both monotonicities are verified when the figure is built, so the figure cannot silently contradict the claim. The within-region cost is 0.081 AUC and is supported in every region. The change in transfer is +0.014 AUC on an interval that spans zero, so no exchange between the two is demonstrated: a local cost is measured and no compensating transfer gain is.




Figure C.8: **Burned-niche overlap does not determine transfer: the contrast pair.** Per-feature signed univariate AUC, the probability that a predictor takes a higher value on a burned cell than on an unburned one, for the two most informative region pairs, natural-vegetation primary population. Arrows run from chance (0.5) to the observed value, so their direction is the sign of the association and their length its strength; whiskers are 95% spatial-block (≈ 5 km) bootstrap intervals. Within each pair the two regions are distinguished by colour, by line style (solid for the first region named, dashed for the second) and by vertical offset; line style is included because the two hues differ by only 2.30:1 in relative luminance and would nearly merge in greyscale. **A filled arrowhead means that region's 95% interval excludes 0.5, so the direction is statistically supported; an open arrowhead means the interval contains 0.5 and the direction is not supported.** Rows shaded grey are features whose sign disagrees between the two regions. Grey side bars are the per-feature Schoener's D overlap. **(a)** Manavgat–Muğla, the pair with the *highest* burned-niche overlap ($\bar{D} = 0.83$), transfers below chance in both directions (0.470 and 0.401). **(b)** Bejís–Montiferru, the pair with the *lowest* overlap ($\bar{D} = 0.48$), transfers above chance in both directions (0.594 and 0.548). High niche overlap is therefore not sufficient for transfer. The arrowhead fill also exposes the coverage limit of the supported-agreement index, which is defined only on features whose direction is supported in *both* regions: there are two such features in (a) (NDVI and elevation) and **none in (b)**, where the two regions' supported sets are disjoint and Montiferru's intervals

- [3] Boschetti, L., Roy, D.P., Giglio, L., Huang, H., Zubkova, M., Humber, M.L., 2019. Global validation of the Collection 6 MODIS burned area product. *Remote Sensing of Environment* 235, 111490. doi:10.1016/j.rse.2019.111490.
- [4] Breiman, L., 2001. Random forests. *Machine Learning* 45, 5–32. doi:10.1023/A:1010933404324.
- [5] de Bruin, S., Brus, D.J., Heuvelink, G.B.M., van Ebbenhorst Tengbergen, T., Wadoux, A.M.J.C., 2022. Dealing with clustered samples for assessing map accuracy by cross-validation. *Ecological Informatics* 69, 101665. doi:10.1016/j.ecoinf.2022.101665.
- [6] Chuvieco, E., Cocero, D., Riaño, D., Martín, P., Martínez-Vega, J., de la Riva, J., Pérez, F., 2004. Combining NDVI and surface temperature for the estimation of live fuel moisture content in forest fire danger rating. *Remote Sensing of Environment* 92, 322–331. doi:10.1016/j.rse.2004.01.019.
- [7] Dimarco, N.A., Faraji, I., Wahbi, M., Maatouk, M., Boulaassal, H., Yazidi Aalaoui, O., El Kharki, O., 2026. Generalizing human-driven wildfire ignition models across Mediterranean regions using harmonized remote-sensing and machine-learning data. *Geomatics* 6, 13. doi:10.3390/geomatics6010013.
- [8] Gelabert, P.J., Jiménez-Ruano, A., Ochoa, C., Alcasena, F., Sjöström, J., Marrs, C., Ribeiro, L.M., Palaiologou, P., Bentu Martínez, C., Chuvieco, E., Vega-García, C., Rodrigues, M., 2025. Assessing

human-caused wildfire ignition likelihood across Europe. *Natural Hazards and Earth System Sciences* 25, 4713–4729. doi:10.5194/nhess-25-4713-2025.

- [9] Giglio, L., Boschetti, L., Roy, D.P., Humber, M.L., Justice, C.O., 2018. The Collection 6 MODIS burned area mapping algorithm and product. Remote Sensing of Environment 217, 72–85. doi:10.1016/j.rse.2018.08.005.
- [10] Gorelick, N., Hancher, M., Dixon, M., Ilyushchenko, S., Thau, D., Moore, R., 2017. Google Earth Engine: Planetary-scale geospatial analysis for everyone. Remote Sensing of Environment 202, 18–27. doi:10.1016/j.rse.2017.06.031.
- [11] Huang, Y., Liu, N., Chen, W., Tan, W., Deng, Y., Wang, Z., Chlus, A., Shen, J., Townsend, P.A., 2026. Enhancing transferability of foliar trait retrieval models: A comparative analysis of transfer learning strategies and domain shift characterization. Remote Sensing of Environment 337, 115336. doi:10.1016/j.rse.2026.115336.
- [12] Iban, M.C., Sekertekin, A., 2022. Machine learning based wildfire susceptibility mapping using remotely sensed fire data and GIS: A case study of Adana and Mersin provinces, Turkey. Ecological Informatics 69, 101647. doi:10.1016/j.ecoinf.2022.101647.
- [13] Jodhani, K., Bhatiya, L., Sirvi, L., Chanda, A., Thakkar, S., Rathnayake, U., 2026. Artificial intelligence and remote sensing frameworks for wildfire monitoring and risk analysis across multiple ecosys-

- tems: A review. *AIMS Environmental Science* 13, 56–98. doi:10.3934/
1570 *environsci.2026004*.
- [14] Kondylatos, S., Prapas, I., Camps-Valls, G., Papoutsis, I., 2023. Meso-
geos: A multi-purpose dataset for data-driven wildfire modeling in the
Mediterranean, in: *Advances in Neural Information Processing Systems*
36 (NeurIPS 2023) Datasets and Benchmarks Track, pp. 50661–50676.
1575 doi:10.52202/075280–2204.
- [15] Liu, C., Mostafavi, A., 2025. WildfireGenome: Interpretable machine
learning reveals local drivers of wildfire risk and their cross-county vari-
ation. *arXiv preprint arXiv:2511.11589*. doi:10.48550/arXiv.2511.
11589, arXiv:2511.11589.
- 1580 [16] Ludwig, M., Moreno-Martinez, A., Hölzel, N., Pebesma, E., Meyer, H.,
2023. Assessing and improving the transferability of current global spa-
tial prediction models. *Global Ecology and Biogeography* 32, 356–368.
doi:10.1111/geb.13635.
- [17] Maffei, C., Alfieri, S.M., Menenti, M., 2018. Relating spatiotemporal
1585 patterns of forest fires burned area and duration to diurnal land sur-
face temperature anomalies. *Remote Sensing* 10, 1777. doi:10.3390/
rs10111777.
- [18] Maffei, C., Lindenbergh, R., Menenti, M., 2021. Combining multi-
spectral and thermal remote sensing to predict forest fire characteristics.
1590 *ISPRS Journal of Photogrammetry and Remote Sensing* 181, 400–412.
doi:10.1016/j.isprsjprs.2021.09.016.

- [19] Maffei, C., Menenti, M., 2019. Predicting forest fires burned area and rate of spread from pre-fire multispectral satellite measurements. *ISPRS Journal of Photogrammetry and Remote Sensing* 158, 263–278. doi:10.1016/j.isprsjprs.2019.10.013.
- [20] Marino, E., Yáñez, L., Guijarro, M., Madrigal, J., Senra, F., Rodríguez, S., Tomé, J.L., 2024. Transferability of empirical models derived from satellite imagery for live fuel moisture content estimation and fire risk prediction. *Fire* 7, 276. doi:10.3390/fire7080276.
- [21] Meyer, H., Pebesma, E., 2021. Predicting into unknown space? Estimating the area of applicability of spatial prediction models. *Methods in Ecology and Evolution* 12, 1620–1633. doi:10.1111/2041-210X.13650.
- [22] Meyer, H., Pebesma, E., 2022. Machine learning-based global maps of ecological variables and the challenge of assessing them. *Nature Communications* 13, 2208. doi:10.1038/s41467-022-29838-9.
- [23] Meyer, H., Reudenbach, C., Hengl, T., Katurji, M., Nauss, T., 2018. Improving performance of spatio-temporal machine learning models using forward feature selection and target-oriented validation. *Environmental Modelling & Software* 101, 1–9. doi:10.1016/j.envsoft.2017.12.001.
- [24] Milà, C., Mateu, J., Pebesma, E., Meyer, H., 2022. Nearest neighbour distance matching leave-one-out cross-validation for map validation. *Methods in Ecology and Evolution* 13, 1304–1316. doi:10.1111/2041-210X.13851.

- [25] Moreno-Torres, J.G., Raeder, T., Alaiz-Rodríguez, R., Chawla, N.V.,
1615 Herrera, F., 2012. A unifying view on dataset shift in classification.
Pattern Recognition 45, 521–530. doi:10.1016/j.patcog.2011.06.019.
- [26] Muñoz-Sabater, J., Dutra, E., Agustí-Panareda, A., Albergel, C., Ar-
duini, G., Balsamo, G., Boussetta, S., Choulga, M., Harrigan, S., Hers-
bach, H., Martens, B., Miralles, D.G., Piles, M., Rodríguez-Fernández,
1620 N.J., Zsoter, E., Buontempo, C., Thépaut, J.N., 2021. ERA5-Land:
a state-of-the-art global reanalysis dataset for land applications. Earth
System Science Data 13, 4349–4383. doi:10.5194/essd-13-4349-2021.
- [27] Pausas, J.G., Keeley, J.E., 2021. Wildfires and global change. Frontiers
in Ecology and the Environment 19, 387–395. doi:10.1002/fee.2359.
- [28] Persello, C., Bruzzone, L., 2012. Active learning for domain adaptation
1625 in the supervised classification of remote sensing images. IEEE Transac-
tions on Geoscience and Remote Sensing 50, 4468–4483. doi:10.1109/
TGRS.2012.2192740.
- [29] Ploton, P., Mortier, F., Réjou-Méchain, M., Barbier, N., Picard, N.,
1630 Rossi, V., Dormann, C., Cornu, G., Viennois, G., Bayol, N., Lyapustin,
A., Gourlet-Fleury, S., Pélissier, R., 2020. Spatial validation reveals
poor predictive performance of large-scale ecological mapping models.
Nature Communications 11, 4540. doi:10.1038/s41467-020-18321-y.
- [30] Podschwit, H., Jolly, W., Alvarado, E., Verma, S., Ponce, B., Markos,
1635 A., Aliaga-Nestares, V., Rodriguez-Zimmermann, D., 2022. Reliability

of cross-regional applications of global fire danger models: A Peruvian case study. *Fire Ecology* 18, 25. doi:10.1186/s42408-022-00150-7.

- [31] Roberts, D.R., Bahn, V., Ciuti, S., Boyce, M.S., Elith, J., Guillerá-Arroita, G., Hauenstein, S., Lahoz-Monfort, J.J., Schröder, B., Thuiller, W., Warton, D.I., Wintle, B.A., Hartig, F., Dormann, C.F., 2017. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography* 40, 913–929. doi:10.1111/ecog.02881.
- [32] Rousseau, J.S., Betts, M.G., 2022. Factors influencing transferability in species distribution models. *Ecography* 2022, e06060. doi:10.1111/ecog.06060.
- [33] Sandholt, I., Rasmussen, K., Andersen, J., 2002. A simple interpretation of the surface temperature/vegetation index space for assessment of surface moisture status. *Remote Sensing of Environment* 79, 213–224. doi:10.1016/S0034-4257(01)00274-7.
- [34] Schoener, T.W., 1968. The *Anolis* lizards of Bimini: Resource partitioning in a complex fauna. *Ecology* 49, 704–726. doi:10.2307/1935534.
- [35] Sun, B., Feng, J., Saenko, K., 2016. Return of frustratingly easy domain adaptation, in: *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence (AAAI-16)*, AAAI Press, Phoenix, Arizona, USA. pp. 2058–2065. doi:10.1609/aaai.v30i1.10306.
- [36] Tuia, D., Persello, C., Bruzzone, L., 2016. Domain adaptation for the classification of remote sensing data: An overview of recent advances.

- IEEE Geoscience and Remote Sensing Magazine 4, 41–57. doi:10.1109/MGRS.2016.2548504.
- 1660
- [37] Valavi, R., Elith, J., Lahoz-Monfort, J.J., Guillerá-Arroita, G., 2019. blockCV: An R package for generating spatially or environmentally separated folds for k-fold cross-validation of species distribution models. *Methods in Ecology and Evolution* 10, 225–232. doi:10.1111/2041-210X.13107.
- 1665
- [38] Vesk, P.A., Morris, W.K., Neal, W.C., Mokany, K., Pollock, L.J., 2021. Transferability of trait-based species distribution models. *Ecography* 44, 134–147. doi:10.1111/ecog.05179.
- [39] Vibhandik, P., Sawant, S., Joshi, A., Bidwe, R., 2026. A systematic literature review on forest fire susceptibility mapping using geo-spatial technology and future research directions. *Discover Artificial Intelligence* 6, 921. doi:10.1007/s44163-026-00921-0.
- 1670
- [40] Wadoux, A.M.J.C., Heuvelink, G.B.M., de Bruin, S., Brus, D.J., 2021. Spatial cross-validation is not the right way to evaluate map accuracy. *Ecological Modelling* 457, 109692. doi:10.1016/j.ecolmodel.2021.109692.
- 1675
- [41] Warren, D.L., Glor, R.E., Turelli, M., 2008. Environmental niche equivalency versus conservatism: Quantitative approaches to niche evolution. *Evolution* 62, 2868–2883. doi:10.1111/j.1558-5646.2008.00482.x.
- [42] Xu, Y., Dai, Y., Chang, L., Wang, Q., Dong, Y., 2026. Does your wildfire prediction model actually work, or just score well?
- 1680

arXiv preprint arXiv:2605.18911. doi:10.48550/arXiv.2605.18911,
arXiv:2605.18911.

- [43] Yates, K.L., Bouchet, P.J., Caley, M.J., Mengersen, K., Randin, C.F.,
1685 Parnell, S., et al., 2018. Outstanding challenges in the transferability of
ecological models. *Trends in Ecology & Evolution* 33, 790–802. doi:10.
1016/j.tree.2018.08.001.
- [44] Yebra, M., Dennison, P.E., Chuvieco, E., Riaño, D., Zylstra, P., Hunt,
E.R., Danson, F.M., Qi, Y., Jurdao, S., 2013. A global review of remote
1690 sensing of live fuel moisture content for fire danger assessment: Moving
towards operational products. *Remote Sensing of Environment* 136,
455–468. doi:10.1016/j.rse.2013.05.029.
- [45] Zanaga, D., Van De Kerchove, R., Daems, D., De Keersmaecker, W.,
Brockmann, C., Kirches, G., Wevers, J., Cartus, O., Santoro, M., Fritz,
1695 S., Lesiv, M., Herold, M., Tsendbazar, N.E., Xu, P., Ramoino, F., Arino,
O., 2022. ESA WorldCover 10 m 2021 v200. Zenodo. doi:10.5281/
zenodo.7254221.