

Highlights

Evaluation geometry and the limits of cross-region transfer in pre-fire thermal wildfire prediction

Emrehan Metin, Yunus Emre Cogurcu

- Scoring the same model at the fire, not region-wide, costs 0.143 AUC
- The cause is the negative pool, not prevalence: +0.155 against -0.000
- Under blocked CV, thermal predictors add +0.06 to +0.15 AUC in five regions
- Equalising the evaluation frame lifts mean transfer from 0.541 to 0.616
- Across 20 transfer directions the thermal gain is +0.004 [-0.028, +0.036]

Evaluation geometry and the limits of cross-region transfer in pre-fire thermal wildfire prediction

Emrehan Metin^a, Yunus Emre Cogurcu^{a,*}

^a*Çukurova University, Adana, Türkiye*

Abstract

Pre-fire thermal wildfire models are rarely tested outside their fitting region. Six thermal predictors were added to a terrain, fuel and greenness baseline in five Mediterranean regions, with MCD64A1 labels and spatially blocked validation. The first result is **evaluation geometry**: with the model held fixed, scoring on the burn scar and its 2 km collar rather than region-wide costs **0.143 ROC-AUC** [**+0.077**, **+0.208**], driven by the negative pool, not class balance. Between regions, equalising the study areas to a 10 km collar, which drops no burned cells, lifts mean transfer from 0.541 to 0.616. Local skill does not travel: the thermal block adds +0.045 to +0.148 within regions at 5 km blocking but +0.004 [−0.028, +0.036] across twenty transfer directions, and the static baseline transfers no better, at 0.537 against 0.541. Transfer skill must be measured on the target region before a model is relied on.

Keywords: wildfire, model transferability, model evaluation, land surface temperature, domain adaptation, concept shift

*Corresponding author.

Email address: `ycogurcu@cu.edu.tr` (Yunus Emre Cogurcu)

1. Introduction

Wildfire is a defining disturbance of Mediterranean-basin landscapes, and a changing climate is reshaping where and how it burns (Pausas and Keeley, 2021). The dominant pattern in fire susceptibility mapping (Vibhandik et al., 2026; Jodhani et al., 2026) is settled: geospatial predictors are assembled over a study region, paired with a historical burned-area record and fitted with a supervised classifier, most often a random forest (Breiman, 2001), and the surface is published with a cross-validated AUC of 0.85 to 0.95. **That pattern does not report portability**, and this paper measures it.

10 1.1. *Portability goes unmeasured*

A susceptibility model’s reported skill is almost always a *within-region* estimate: held-out folds come from the same study area and season, often the same fire, and random folds over cells let spatial autocorrelation inflate it further. The problem is documented across ecological modelling (Roberts et al., 2017; Ploton et al., 2020) and addressed by spatially blocked cross-validation (Valavi et al., 2019; Meyer et al., 2018), but blocking does not make the estimate honest about a fire the model has not seen.

Because only that side of the ledger is reported, portability is never entered: a predictor block is adopted on the increment it delivers inside its training footprint, and whether that survives a change of region is not asked, even though any regional product built from locally trained models implicitly promises generalisation beyond it. Meteorological fire-danger indices are known not to port cleanly in a Peruvian case study (Podschwit et al., 2022), and the two studies that test transfer systematically (Dimarco et al., 2026;

25 Liu and Mostafavi, 2025) both report that it largely succeeds between similar regions — and both transfer models whose dominant predictors are *spatially stationary*. Whether a dynamic, season-specific class behaves the same way is this paper’s question.

1.2. Pre-fire thermal dryness is the natural test case

30 Pre-fire thermal dryness is the dynamic class most plausibly *expected* to transfer. Standard predictors are static or near-static over the timescale at which fire danger varies, so they explain poorly why one summer burned and the preceding one did not; what changes is the state of the surface, to which satellite thermal observation gives partial access (Section 2.2). The physics
35 linking moisture stress to combustion is universal, so portability should be most expected here, which makes the class diagnostic: a loss cannot be dismissed as a peculiarity of a locally defined covariate. The expectation is sharpest for the internally normalised channels, which should be least exposed to absolute-temperature offsets between regions; **they transfer no**
40 **better than the absolute ones** (Appendix A(f)). The expectation motivates the design and is not a finding of it — Section 4.4 reports the associations running the other way.

1.3. Contributions

Two findings carry this paper, stated here as claims and established in
45 Section 4, which carries every interval.

Contribution 1. Where a model is scored decides what it appears to know, and the effect is large enough to dissolve findings of our own. That evaluation extent inflates AUC is established in species

distribution modelling (Section 2.4); what is new is a magnitude under a
50 controlled design. Holding model, predictors and fitting fixed and changing
only which cells are scored costs **0.143 ROC-AUC**, the size of the incre-
ments this design measures for a predictor block, with the cause measured as
the negative pool’s composition rather than class balance (Section 4.3). Ap-
plied between regions it withdraws five claims of our own, including the sign
55 reversal we had offered as the mechanism and the one arm that held place
fixed (Section 4.4), and the practical consequence is a reporting standard
(Section 5.6).

**Contribution 2. Local skill does not travel, and correcting the
frame does not rescue it.** The thermal block is worth a substantial within-
60 region increment in all five regions, every bootstrap interval above zero, and
stays positive when the frame is equalised — but much of it is a property of
interleaved holdout, and across twenty transfer directions the paired con-
tribution spans zero on both frames under the primary resampling unit,
though much closer to the boundary once equalised and not under every
65 admissible unit (Section 4.4), with a sign belonging to the pair rather than
the block. Two controls bound this: the static baseline transfers no better
on either frame, so the failure is not the thermal block’s peculiarity, and on
matched frames and blocking equalised transfer still falls **0.155** short of the
within-region reference. The within-region half is not novel (Alkan Akıncı
70 and Akıncı, 2023; Iban and Sekertekin, 2022); the paired contrast against
portability is.

Two consequences follow, as supporting results: label-free alignment by
standardisation and covariance alignment (Sun et al., 2016) does not repair

transfer, in what we believe is its first application to fire susceptibility; and
75 because the residual is conditional, the resource that closes it is target labels,
priced in Appendix A(u).

The leakage-audited, spatially blocked protocol is released with code,
configuration and frozen outputs, so most of this can be re-run rather than
taken on trust, the release being complete as the declarations record — which
80 matters given evidence that wildfire transfer conclusions are sensitive to eval-
uation design (Xu et al., 2026). A companion paper treats the observational
layer beneath this one.

2. Related work

2.1. Fire susceptibility mapping, and what it reports

85 The dominant pattern is described in Section 1; the figure it publishes,
a cross-validated AUC of 0.85 to 0.95, is a within-region estimate. Com-
parable within-region results already exist for the landscapes studied here
(Alkan Akıncı and Akıncı, 2023; Iban and Sekertekin, 2022).

2.2. Pre-fire thermal dryness

90 Satellite thermal observation gives repeated access to surface state
through fuel moisture content (Yebra et al., 2013). Chuvieco et al. (2004)
established the pairing of land surface temperature with a vegetation
index as a live-fuel-moisture estimator for fire-danger rating, and the
Temperature-Vegetation Dryness Index (Sandholt et al., 2002) formalises
95 that feature space into an internally normalised measure, which is why
it is often expected to travel better than raw temperature. That pre-fire

thermal state carries information about subsequent fire is established (Maffei et al., 2018; Maffei and Menenti, 2019; Maffei et al., 2021), and Gelabert et al. (2025) found dead fine fuel moisture and its anomalies the most
 100 influential predictor of human-caused ignition likelihood across Europe, testing generalisation by pooled fitting with per-site evaluation. What has not been tested is whether such a classifier survives strict, label-free application to an unseen region.

2.3. Spatial validation and transferability

105 The canonical taxonomy separates covariate shift, where the predictor distribution moves but the predictor-response relationship holds, from concept shift, where the relationship itself changes (Moreno-Torres et al., 2012). Covariate shift is in principle correctable without target labels, by per-region standardisation or covariance alignment such as CORAL (Sun et al., 2016),
 110 and domain adaptation has an established remote-sensing literature (Tuia et al., 2016; Persello and Bruzzone, 2012). Concept shift is not: no realignment of inputs can repair a reversal in the sign of an association, because detecting one requires the labels being withheld. Shift decomposition in applied remote sensing is not itself new (Huang et al., 2026), but we are aware
 115 of no prior application of covariance alignment to fire susceptibility, fire occurrence or burned-area prediction.

2.4. Cross-region generalisation of fire models

Few studies test fire-model transfer directly. Podschwit et al. (2022) report, in a Peruvian case study, that meteorologically derived danger indices do not port cleanly between fire environments. WildfireGenome (Liu
 120

and Mostafavi, 2025) runs a leave-one-county-out matrix across seven US counties and reports strong within-county performance with highly variable off-diagonal transfer, similar pairs transferring well and dissimilar pairs collapsing; its label is a principal-component composite of hazard *indicators* rather than observed burned area, and it applies no adaptation. Xu et al. (2026) argue that wildfire transfer conclusions depend strongly on evaluation design, a caution we address by fixing the protocol in a project log before the diagnostics were computed and by reporting every sensitivity axis; that log is not a formal pre-registration, and Section 3.11 states what was fixed and when. Kondylatos et al. (2023) provide Mesogeos, a 1 km Mediterranean datacube.

Evaluation extent and AUC. Species distribution modelling settled long ago that the area a model is evaluated over is not a neutral choice. Lobo et al. (2008) make it the fifth and, by their own ranking, most important reason to distrust AUC comparatively: the extent of the modelled area governs how many easy absences enter the calculation, and therefore the score. VanDerWal et al. (2009) show the same lever on the calibration side, and Barve et al. (2011) give the argument its general form as the accessible area. That literature is qualitative about magnitude, because magnitude is problem-specific, so our Contribution 1 is a measurement inside that framework rather than a new phenomenon. The wildfire literature has largely not imported the lesson: region-wide figures are reported as though they described performance at the fire, and we know of no wildfire study that holds the model fixed and varies only the evaluation cells — which is why Section 5.6 treats the 0.143 as a reporting problem rather than a caveat.

The nearest neighbour, and the contrast this paper draws. Dimarco et al. (2026) is the closest Mediterranean analogue: 500 m predictors harmonised across four countries, tree ensembles under spatial cross-validation, transfer tested both leave-one-country-out and as a full 4×4 matrix, every transfer exceeding AUC 0.80 and bioclimatically similar countries scoring higher, with no domain adaptation.¹ Much of the design is shared — Mediterranean regions, 500 m cells, MCD64A1-derived targets, tree ensembles, spatially aware validation, an explicit transfer matrix — and two things differ. **The predictor class:** their model rests on attributes of a place, all spatially stationary, ours on the state of a surface in one season.

The response variable: theirs is human-driven ignition, ours burned area, which have different dominant controls. Read together the two results suggest the relationship between domain similarity and transfer success may be predictor-class dependent, with Vesik et al. (2021) and Rousseau and Betts (2022) aligning from species distribution modelling while Dimarco et al. and WildfireGenome run against it. We present that as a live disagreement; Section 5.5 states why our data do not settle it.

3. Data and methods

3.1. Study regions and temporal windows

Five Mediterranean wildfire regions are analysed (Fig. 1): Manavgat 2021 and Muğla 2021 in Türkiye, Bejís 2022 in Spain, North Evia 2021 in

¹Their Results text refers to “LST anomalies”, inconsistent with their own Methods, where no land surface temperature or TVDI variable appears and the only temperature predictor is a static ERA5-Land seasonal climatology (Muñoz-Sabater et al., 2021). We follow their Methods and note the discrepancy so the comparison is transparent.

Greece and Montiferru 2021 in Sardinia. Each is a place-based rectangle in EPSG:4326, defined from place coverage rather than from a fire perimeter and deliberately not clipped to it, so that unburned cells around each fire form the negative class. **The record does not show every box fixed before any outcome was seen**, so what it does show is stated region by region. In the pipeline’s version history each box has a single committed value, never changed afterwards. Montiferru’s is derived deterministically from the union of four municipal boundaries. Manavgat’s was drawn to exclude the coastal cropland belt and committed together with the gate-only workflow that first evaluated it, the commit that first gave the pipeline any Manavgat geometry; its registry comments disagree on whether the drawing preceded or followed that first gate run. Bejís’s is still labelled the initial candidate in the registry but was committed after its first gate and model results. Muğla’s coordinates appear in a dated preflight record about four minutes before its first gate result and were committed to the registry only afterwards. No box needed adjusting to pass the gate, whose admitted margins are wide (Section 4.1). **One choice was label-informed and is stated as such:** the North Evia box was extended after the legacy box proved atypically high in burned prevalence, the extended geometry then defined from place anchors and the legacy variant kept as a sensitivity arm. Section 4.4 shows this framing decision is consequential and Appendix C.5(ix) treats it as the design lesson of the paper. A sixth region, Kozan 2023, is carried as a negative control and excluded by the gate of Section 3.3.

Each region has two non-overlapping windows. The **predictor window** closes the day before the **label window** opens, so no predictor observation




Figure 1: Study regions and the gate negative control. Locations of the five wildfire regions analysed in this study (blue, circular markers): Manavgat 2021 and Muğla 2021 (Türkiye), Bejís 2022 (Spain), North Evia 2021 (Greece, extended AOI) and Montiferru 2021 (Sardinia, Italy). Rectangles are the true analysis areas of interest; the ring marks each centroid, because at basin scale the rectangles alone are hard to locate. Kozan 2023 (grey, dashed outline, square marker) is *not* a study region: it is the negative control used to calibrate the burned-landcover admissibility gate, and it enters no model, transfer or diagnostic anywhere in this paper. Its rectangle is the true extent of its processed AOI and is larger than the study AOIs for that reason alone; the difference in size carries no methodological meaning and is not a design choice. Coastlines are Natural Earth 1:50 m; country borders are omitted deliberately. The scale bar is exact only at 38°N. Region bounding boxes match Table C1 to 5×10^{-3} degrees.

can post-date the first labelled burning. Lengths vary with the event: 57 to 61 days for predictors, 35 to 59 for labels, counted inclusively. A four-year baseline of window-symmetric composites precedes each predictor window and supplies the climatological reference for the anomaly channels. Region
195 identifiers, bounding boxes, window dates and baseline years are registered in `core/regions.py` and reproduced in Appendix C, Table C1. The processing chain and the evaluation programme it feeds are drawn in Fig. 2.

3.2. Burned-area label and the ~ 500 m analysis grid

200 Labels come from MODIS MCD64A1 Collection 6.1 (Giglio et al., 2018) retrieved through Google Earth Engine (Gorelick et al., 2017), whose omission and commission characteristics (Boschetti et al., 2019) bound every model here. The analysis grid is **reconstructed rather than native**: the pipeline’s 30 m EPSG:4326 reference grid is partitioned into 17 x 17 blocks,
205 giving a nominal 510 m cell that approximates rather than reproduces the MODIS cell and is square in degrees but not on the ground. A cell’s burn date is the mode of its positive sub-pixel day-of-year values, tested against the label window; the label never affects eligibility for modelling. Two safeguards are recorded rather than assumed: cells burning before the label window
210 opens are removed, which ran for three of five regions and excluded 49 cells in Muğla, 16 in North Evia and 61 in Montiferru, with none arising in Manavgat or Bejís, and burning in earlier years is screened for none. Appendix C.1 gives the full specification, including what follows from the grid’s shape.




Figure 2: Processing pipeline and evaluation programme. Flow from source products (left) to the transfer and diagnostic programme (right); every box carries the Methods subsection that specifies it. The five source products (MCD64A1 burned area, Landsat Collection 2 Level-2, MOD11A1, Copernicus DEM GLO-30 and ESA WorldCover) merge onto the reconstructed ~ 510 m analysis grid (17×17 pixels of the 30 m reference grid). The burned-landcover admissibility gate has two outcomes: regions that *pass* proceed to the natural-vegetation primary population, while the rejecting branch (grey, dashed) terminates in the Kozan 2023 negative control, whose burned area is 98% cropland and which is never modelled. The dashed connector deliberately carries no arrowhead onward, so the control cannot be read as entering the pipeline. Feature sets feed two evaluations in parallel: within-region (random forest under spatial-block cross-validation) and cross-region transfer (source-only, 20 ordered directions, with label-blind z-score and CORAL adaptation). Their levels are combined in the transfer-gap decomposition (§3.12). The footer states what applies throughout: leakage hard-exclusion, seed 42, 1000-replicate spatial-block bootstrap (§3.13), and the sensitivity designs (§3.16). Section references are verified against the Methods heading list when the figure is built.

3.3. Burned-landcover admissibility gate

215 Before any modelling each region passes a gate asking what fraction of
its burned cells is dominated by natural vegetation, using ESA WorldCover
classes (Zanaga et al., 2022) on the same cells. A region is admitted at 0.50
with at least 30 burned cells, and rejected as a cropland-dominated control
at 0.50 cropland. Verdicts and the purpose of the gate are in Section 4.1, the
220 full rule in Appendix C.1.

3.4. Predictor variables

Two feature sets are used. The **baseline** is terrain, fuel and green-
ness: elevation, slope, dominant land cover and a predictor-window median
vegetation-index composite, all static or near-static over the timescale at
225 which fire danger varies. The **thermal** set adds six pre-fire channels: current
land surface temperature, its anomaly against a four-year window-symmetric
baseline at the same cell, the Temperature-Vegetation Dryness Index and its
difference against that baseline, and two coordinate-informed products, a
downscaled and a fused surface temperature. The two differenced channels
230 are the ones constructed to isolate the dynamic anomaly. TVDI’s wet and
dry edges are percentiles of the values a given area and window happen to
contain, so **it is not portable as a physical quantity independently of
any concept shift** (Appendices C.4, C.5).

3.5. Cell aggregation, validity and analysis populations

235 Cell-level values are means over the ~ 510 m cell from valid 30 m pixels
only, with the valid fraction recorded. A cell is `valid_for_modeling` when
it is analysis-eligible, meaning not excluded for pre-label burning, and its

predictors are valid: joint finite NDVI, elevation and slope support over at
 least 30 % of the cell, finite means for those three channels, and at least
 240 one valid land-cover pixel. **Thermal completeness is not part of the
 definition:** depending on the region, 6 % to 58 % of valid cells carry at
 least one missing thermal channel, and missing values are imputed inside the
 fitting pipeline (Section 3.6) rather than by excluding the cell. The **primary**
 population is natural vegetation, cells whose combined tree, shrub and grass
 245 fraction reaches 0.50, which excludes cropland from every burnable mask.
 The **secondary** population is all valid cells; the frozen export carries it for
 the within-region arm in two regions, reported as a sensitivity in Appendix
 A(v), and the transfer matrix is defined on the primary population only.

3.6. Classifier

250 The two feature sets of Section 3.4 are nested, the thermal set being the
 baseline plus the six thermal channels. Land cover is one-hot encoded. Miss-
 ing numeric values are median-imputed and the categorical channel most-
 frequent-imputed, with the imputers fitted inside each training fold, and on
 the source alone in transfer, so no held-out or target statistic enters a fit.
 255 Features are not otherwise standardised; standardisation appears only as
 the adaptation intervention of Section 3.9. The classifier is a random forest
 (Breiman, 2001) with 300 trees, unlimited depth, `min_samples_leaf = 3`,
 balanced class weights and `random_state = 42`, identical for every region,
 population, feature set and transfer direction, so that no comparison here is
 260 confounded by a model choice.

3.7. *Spatial-block cross-validation and bootstrap uncertainty*

Cross-validation is spatially blocked. Blocks are squares of the analysis grid formed as `row_500m // B` by `col_500m // B`, and five folds are drawn over whole blocks, so no block is split between training and test. B is reported
265 at 2, 10 and 20 cells, about 1, 5 and 10 km.

Uncertainty is a spatial-block bootstrap: blocks are resampled with replacement, 1000 replicates, seed 42, with 2.5 and 97.5 percentiles as the interval. Because it is blocks that are resampled, what bounds an interval’s reliability is the number of blocks carrying at least one burned cell, and those
270 counts are reported alongside the intervals. Where a verdict rests on too few such blocks it is stated as indicative rather than as an interval.

3.8. *Cross-region transfer protocol*

For each ordered pair the model is fitted on the source population and applied to the target with **no target label used at any point**: no refitting,
275 no threshold selection, no calibration. Both feature sets are transferred, so the thermal block’s contribution is a paired difference within a direction rather than a comparison across directions. All twenty ordered directions are computed.

3.9. *Label-blind domain adaptation*

280 Two label-free remedies are tested on every direction, in both feature sets.

Region-wise z-score. Each region’s numeric features are standardised using its own statistics, source statistics from source data and target statistics from target data, never pooled. The classifier is refitted on the z-scored

source and applied to the z-scored target. This removes first- and second-
 285 order marginal offsets.

CORAL after region-wise z-score. The source covariance is aligned to the target’s by the standard whitening-recolouring map (Sun et al., 2016), $\lambda = 10^{-5}$. Critically **the transform is applied to the source only**; the target is left as it is, and the classifier is refitted on the aligned source. Neither
 290 variant sees a target label, and both are verified label-blind at run time. λ sensitivity was assessed over nine values on four of the twenty directions, moving transfer AUC by at most 0.014, and no value of λ was selected on performance; the $\lambda = 1$ of the original CORAL formulation lies outside that sweep, while the value used throughout remains $\lambda = 10^{-5}$ (Appendix A(b)).

3.10. *Transfer-gap decomposition and the concept-shift criterion*

The transfer gap is decomposed as (adapted – raw) / (within – raw), signed and unclipped, with its interval from the same paired bootstrap; Section 4.3 shows the remainder should not be read as a conditional residual, because much of it is incurred inside a single region (Appendices A(j), C.7).

300 The mechanism is diagnosed by **signed univariate association**: the raw ROC-AUC of each numeric predictor against **burned** is computed per region and never folded to $\max(\text{AUC}, 1 - \text{AUC})$, so a value below 0.5 is a direction rather than weakness. A reversal is called bootstrap-supported only when the two regions’ point estimates fall on opposite sides of 0.5 **and each**
 305 **region’s own interval excludes 0.5**. That is stricter than requiring the intervals to be disjoint, and a feature failing the second condition is recorded as a point reversal only.

3.11. Interventions

Pooled multi-region training. Leave-one-region-out: the model is
310 trained on the pooled primary populations of the other four regions and
evaluated on the held-out region, with folds blocked as before. This asks
whether pooling recovers what single-source transfer loses.

Removal of direction-reversing features. The two predictors re-
versing with bootstrap support **on the frames as drawn** are **elevation_**
315 **mean** and **lst_anomaly_mean**. Section 4.4 withdraws that support under
an equalised frame, so the selection rule is frame-dependent and this arm is
reported as a measurement under the original protocol, not as a consequence
of an established reversal. Both are dropped and everything refitted, within-
region and across every direction, and each is dropped singly so the cost can
320 be attributed. The two features are selected by the same reversal analysis
the result is then read against, so **both quantities are post-selection**
estimates with no correction applied.

3.12. Controls on the transfer path

Four evaluations establish what the transfer arms are measuring, all using
325 the transfer protocol of Section 3.8 unchanged: a **within-region half-split**
(modelled cells cut at the median of a grid axis, both axes and directions,
a split discarded when either half is single-class); **leave-one-scar-out** (each
burned component of at least 50 cells, dilated by a 2 km buffer, withheld
from training and used as the target); a **foreign-region evaluation** of the
330 same held-out scar areas; and the **same blocked model restricted** to those
areas, which isolates the evaluation region from the training regime. Buffers

of 2, 5 and 10 km were run, 2 km primary. The per-split and per-scar positive counts are unequal and bear on the interpretation (Appendix A(i)).

3.13. Leakage control and reproducibility

335 An explicit forbidden-column set is enforced at every model fit as an as-
sertion rather than a convention: coordinates and their normalised forms,
every burn-date and label-provenance column, and the agreement fraction
are excluded from all feature sets. The natural-vegetation mask defines the
population and is never a predictor. All randomness uses seed 42 and the
340 bootstrap 1000 replicates, with one qualification: the diagnostic bootstraps
of the released appendices use per-measure offsets from that seed rather than
the seed itself, so that independent measures do not share a resampling draw.
Two of the twenty transfer verdicts are not stable across seeds and are iden-
tified in Section 4.5. Because the transfer analysis runs in an environment
345 separate from the upstream pipeline’s, every within-region model was refit-
ted against the frozen upstream output: the within-region comparisons agree
exactly and the twenty transfer directions to within 1.6×10^{-7} , with the tol-
erance that applies if the library version is not pinned in Appendix C.5(vi).
Headline results are repeated across two populations, three block sizes, the
350 CORAL sweep, both feature sets and four classifier capacities, and where a
conclusion depends on one of those choices **the dependence is reported
rather than resolved by choosing the favourable setting** (Appendices
A, C.6).

3.14. Same-geography event-to-event comparison (*Muğla 2021 versus 2022*)

355 Muğla admits a comparison in which place is held fixed and the event varies: a second fire burned inside the identical AOI, on the identical grid, eleven months after the first. Signed univariate AUCs are computed for both arms under the same 10-cell bootstrap used elsewhere. **Season, year and population all differ**, since the 2022 arm is defined by removing the 2021
360 scar (Appendices A(m), C.3).

4. Results

4.1. Study regions and the admissibility gate

All five candidate regions pass the burned-landcover gate and the negative control fails it as intended: Kozan 2023 returns a natural-vegetation
365 fraction of 0.017 against the 0.50 threshold and is excluded. The separation is not marginal, the five admitted regions carrying 0.723 to 0.991, so the threshold falls in an empty interval rather than between neighbouring cases. That is consistent with the gate separating natural-fuel combustion from post-harvest stubble burning, which MCD64A1 does not distinguish,
370 **though one negative control cannot establish it**. North Evia is analysed on an extended AOI, which leaves the burned scar essentially unchanged while cutting TSG prevalence from 0.676 to 0.287; the effect on transfer is in Appendix A(a).

4.2. Within-region: the thermal increment replicates in five regions

375 Adding the six thermal predictors to the baseline raises spatially blocked out-of-fold ROC-AUC in every region, the increment's bootstrap interval

excluding zero in all five at 1 km and at 5 km blocking (Fig. 3) — the two scales this design supports as intervals. At 10 km the point estimates hold, from +0.048 to +0.154, but rest on 6 to 33 positive-carrying blocks and are
380 indicative.

Table 1: Within-region baseline versus thermal performance and block-size robustness. Primary (TSG) population; spatially blocked 5-fold CV (Section 3.7); paired spatial-block bootstrap, 1000 replicates. Block sizes 2/10/20 cells $\approx$ 1/5/10 km. Baseline and Thermal columns are ROC-AUC; the 95 % CI belongs to Δ AUC.

Region	Block	Baseline	Thermal	Δ AUC	95% CI
Manavgat 2021	2	0.803	0.870	+0.067	[+0.055, +0.079]
	10	0.748	0.797	+0.050	[+0.023, +0.077]
	20	0.683	0.731	+0.048	[+0.014, +0.085]
Bejís 2022	2	0.862	0.918	+0.056	[+0.048, +0.065]
	10	0.779	0.825	+0.045	[+0.018, +0.069]
	20	0.739	0.795	+0.057	[+0.031, +0.090]
Muğla 2021	2	0.743	0.859	+0.116	[+0.106, +0.125]
	10	0.698	0.777	+0.079	[+0.050, +0.105]
	20	0.673	0.733	+0.061	[+0.030, +0.094]
North Evia 2021 (ext.)	2	0.759	0.912	+0.153	[+0.142, +0.166]
	10	0.716	0.864	+0.148	[+0.119, +0.183]
	20	0.679	0.833	+0.154	[+0.124, +0.189]
Montiferru 2021	2	0.781	0.883	+0.101	[+0.080, +0.125]
	10	0.620	0.720	+0.100	[+0.017, +0.186]
	20	0.555	0.681	+0.126	[+0.053, +0.228]

Table note (resampling units). The bootstrap resamples spatial blocks, so an interval’s reliability is bounded by the number of blocks carrying at least one burned cell, and those counts fall from 192 to 843 at 2 cells to **6 to 33** at 20 cells. An interval built on six such blocks has no meaningful
385 coverage, so **the 20-cell row should be read as indicative rather than as an interval.** The 10-cell row is the coarsest blocking this design supports




Figure 3: **The thermal increment replicates in all five regions and survives coarser spatial blocking.** Natural-vegetation primary population; random forest with the primary weighted configuration; all intervals are 95% spatial-block bootstrap intervals (1000 replicates, seed 42). **(a)** Baseline versus thermal out-of-fold ROC-AUC per region at the default 2-cell (~ 1 km) blocking: the open marker is the static baseline feature set, the filled marker is the baseline plus the six thermal predictors, and the connecting bar is the increment. The right-hand column reports the paired ΔAUC and its 95% interval; it is a fixed column placed beyond the plotted AUC range, so no label position depends on a data value. Every interval excludes zero. **(b)** Absolute thermal ROC-AUC against spatial block size (2, 10 and 20 cells, ≈ 1 , 5 and 10 km): absolute skill falls as blocking coarsens, in every region. **(c)** The thermal increment ΔAUC against the same block sizes, with 95% intervals: the increment does *not* melt away, and its interval excludes zero in all ten of the 2- and 10-cell region-block-size combinations. The five 20-cell intervals also lie above zero but rest on 6 to 33 positive-carrying blocks each, so they are plotted as indicative rather than as intervals (Table 1, table note). In (c) the five regions are offset horizontally by 0.15 cell so their whiskers separate; **this offset is cosmetic, since all five regions are evaluated at exactly the same three block sizes** (2, 10, 20 cells), and the offset is smaller than half the spacing between block sizes, so the groups cannot merge. Colour and marker shape both encode region, so the panels remain readable in greyscale.

properly, every region there having 16 to 70 positive-carrying blocks, and the increment holds at that scale in all five regions (Appendix A(c)).

This within-region result is not itself novel; it is reported because the
390 transfer arms below are measured against it.

4.3. *What the evaluation frame is worth, within one region*

This section isolates the effect Section 4.4 applies between regions. It comes first because it is measured on one model in one region with no transfer involved, so nothing the transfer arms do can explain it (Appendix A(i)).

395 **Where the skill is lost, on a matched comparison.** Four evaluations are reported; the last three are scored on **identical cells**, so they differ only in what the model was trained on, and the first shows why an unmatched comparison misleads. The held-out unit is a burned connected component of at least 50 cells with all cells within 2 km of it. What makes
400 that harder than a whole region is the composition of its negatives, not its burned fraction, and that is measured rather than argued. Scoring the same predictions on a random sample of region cells drawn at the scar area's own burned fraction gives **0.782 against the region-wide 0.782**, so matching prevalence changes nothing, at $-0.000 [-0.003, +0.002]$; scoring them on the
405 scar area gives 0.627, a fall of **+0.155 [+0.093, +0.217]**. That swaps both pools at once, so each was swapped singly: replacing only the negatives costs **0.147 [0.104, 0.191]**, the whole of it, and replacing only the positives costs **0.002 [-0.052, +0.048]**, null on average though the per-scar effect runs from -0.12 to $+0.12$. **The effect is the negative pool** — every negative
410 in a scar collar is fire-adjacent, sharing the terrain, land cover and synoptic conditions of the positives, whereas a region's negatives include its easy

far field — and ROC-AUC is in any case invariant to class balance at fixed class-conditional distributions (`pool_decomposition.json`).

Table 2: The four evaluations, scored on identical cells. Primary natural-vegetation population. Rows B, C and D are scored on the held-out scar area; row A is the whole region and is shown to make the mismatch visible. Means and Student t intervals are over the eight held-out scars.

Evaluation	Model trained on	Scored on	Mean AUC	95 % CI
A. Blocked cross-validation, 5 km	the region, scar included	the whole region	0.776	[0.738, 0.814]
B. Same blocked model, restricted	the region, scar included	the scar area	0.634	[0.552, 0.716]
C. Leave-one-scar-out	the region, scar withheld	the scar area	0.552	[0.501, 0.602]
D. Foreign region	another region, 306 to 2,802 km	the scar area	0.555	[0.495, 0.616]

All four rows are means over the **same eight scars**, and Bejís has no
415 row C because it is its region’s only component of any size, which is why
the controls above are computed on nine and report 0.782 and 0.627 against
this table’s 0.776 and 0.634. The intervals are Student t over the scars, this
arm’s resampling unit rather than the spatial-block bootstrap used elsewhere.
Four scars are in Muğla and two in Montiferru, and **row A is a region-**
420 **level quantity repeated identically across a region’s scars, so its**
interval is pseudo-replicated and should not be read as coverage.
That repetition propagates into $A - B$ and $A - C$; clustering by region
gives **+0.155 [+0.082, +0.228]** and **+0.251 [+0.093, +0.409]**, both still
excluding zero but the second 2.3 times wider, so **the scar-level intervals**
425 **on these two differences overstate precision** (Appendix A(i)).

The same model, scored two ways on the same region, differs by

0.143 AUC. Rows A and B are one model and one set of out-of-fold predictions; the only change is which cells they are scored on. That change alone costs **0.143** [**+0.077**, **+0.208**] of the 0.225 [**+0.157**, **+0.293**] fall from A to
430 C, about two thirds — a ratio of two estimates, compatible with anything from roughly a third to nine tenths, so the point is the size of the numerator rather than the precision of the fraction. It is the size of the predictor-block increments this design itself measures, so a region-wide blocked figure is an upper bound on what the same model achieves where the fire actually is.

Neither withholding the fire nor moving 2,800 km has a measurable cost. B minus C, the fire-specific residual on identical cells, is **+0.082** [**−0.011**, **+0.175**]; C minus D, the effect of replacing a same-region model with one fitted 306 to 2,802 km away, is **−0.003** [**−0.075**, **+0.069**]. Both span zero and both arms sit close to chance, so **with eight scars, three of**
440 **them starved of positives, this design cannot establish a fire-specific residual, only bound it at about 0.18.** It is not shown to be zero on an unseen fire; what this design establishes is that it is not established there. Nor does it establish the fire *event* as the unit, since the held-out patch is defined by the labels and its identity cannot be separated from its location.
445 Appendix A(z) reports the spread behind row D and the sweeps showing the patch definition does not drive the result.

The increment declines with the holdout, and is not established once the fire is withheld. Contribution 1 is about the paired thermal-minus-baseline difference, so the same evaluations were run on it: **+0.056** to
450 **+0.153** under blocked cross-validation with every interval above zero, **+0.027** on a within-region half-split, **+0.022** [**−0.032**, **+0.077**] under leave-one-

scar-out, and +0.004 $[-0.028, +0.036]$ across regions. The point estimate falls monotonically as the holdout hardens and the last two intervals span zero, so **the within-region increment is substantially a property of interleaved holdout.**

4.4. The same effect between regions, applied to our own matrix

Section 4.3's effect applies with equal force between regions. This test is reported before the transfer matrix because it changes what Section 4.5 and Appendix A(s) can claim; it withdraws nothing in Section 4.3, which is what it is built from. Sources are `aoi_frame_auc_frozen_mugla.csv` and `aoi_frame_transfer_frozen_mugla.csv`, which supersede the pre-correction files for the reason given below, with `collar_frame_bootstrap.csv` and `diagnostics_collar_frame.csv`; code is under `paper/code/` and the elaboration in Appendix A(w).

The five areas of interest are not comparable frames. Each is a rectangle drawn around a fire, and they differ by an order of magnitude in how much unburnt far field they enclose: the share of modelled cells beyond 10 km of any burned cell runs from **2.1 %** in Montiferru to **63.1 %** in Bejis (Appendix B, Table B6). That far field is not neutral — in Manavgat the median elevation of modelled cells rises from 472 m within 5 km of the fire to 1,273 m at 20 to 50 km, against 512 m for the burned cells themselves.

Under an equalised frame the sign reversals do not survive. Restricting every region to cells within 10 km of any burned cell removes only far-field negatives — every burned cell is at distance zero and is retained at any radius — so the protection against a flattering radius is the sweep, not the retention of positives. On that frame all five regions agree in sign

on elevation, LST and TVDI, and both bootstrap-supported elevation reversals of Table B3 disappear. Under this paper’s own criterion from Section 3.10, **no between-region reversal remains bootstrap-supported:**
480 the two features straddling 0.5 at the point estimate are the internally differenced channels, and in each at least one region’s interval covers 0.5 — Evia’s by 0.003. The weaker difference instrument that Table B3’s note commits this paper to does support four opposite-sided pairs on `lst_anomaly_mean`, with three caveats that keep it weak. **The honest statement is that no**
485 **between-region reversal meets the strict criterion under the collar, and that the LST anomaly differs between regions on the weaker one.** Two scope limits belong with it, both in Appendix A(l) and A(w): the interval criterion was evaluated at the 10 km collar only, though at the point estimate the two radii agree exactly; and the collar reduces two regions’ cell
490 counts substantially, so part of the loss of support is a loss of power, against which features supported in both regions *rise* from 1.20 to 3.40 per direction.

The sign they now agree on is not the one the dryness framing predicts: at the point estimates a hotter pre-fire surface burned *less* in every region, and on mutual adjustment temperature survives where greenness does
495 not, so the absolute channels behave here as static land-surface descriptors rather than as dryness. Within distance to the nearest burned cell the sign crosses 0.5 in Manavgat and is attenuated to near-null in Montiferru, so in those two it is largely the residual spatial gradient the collar was introduced to remove. The stratifications, the correlations and the reciprocal adjustment
500 are in Appendix A(k).

The same test removes a result of our own. The diagnostic that

best ordered transfer in our matrix is built from exactly the signed AUCs shown above to be frame artefacts, and was correlated against transfer on the same unequal frames. Recomputed under the collar it is unanimous —
505 1.0 in all eighteen directions, variance exactly zero — so it has no variance left to correlate, and its cosine variant falls from $\rho = +0.81$ to -0.06 . The correction does not weaken that result; it dissolves it (`diagnostics_collar_frame.csv`, `collar_increment_and_cosine.csv`).

Two further arms move with the frame. The same-geography arm
510 is the most extreme case in the cohort — a fixed study area is not a fixed evaluation frame, and its 2022 arm has 93.2 % of cells beyond 10 km of any burned cell against 55.3 % for 2021 — so under the collar its two arms fall on the same side of 0.5 on elevation, the feature whose reversal we had reported. The claim is scoped: that arm is reconstructed from the 2021
515 predictors, valid for the year-invariant channels and **not** the seasonal ones, so **only elevation and slope carry a verdict, on both the reversal is removed, and on the thermal channels the arm is silent**; it rests on eleven positive-carrying blocks against a floor of sixteen (Appendices A(m), A(o)). The within-region increment, by contrast, **survives**, positive in all
520 five regions at a mean of $+0.077$ against $+0.086$ as drawn.

The transfer matrix moves as well. Restricting source and target to the same collar:

A data-provenance defect in this arm was found and corrected: one region’s predictor file at the canonical path had come to differ from the
525 one the frozen tables were computed on, so every arm here was re-run against the frozen export. It moves forty of a hundred per-direction values by up to

Table 3: Cross-region transfer under equalised evaluation frames. Primary natural-vegetation population, thermal model, twenty ordered directions per row. Above/below chance are point counts; the supported counts use a 10-cell (≈ 5 km) spatial-block bootstrap on the target, 1000 replicates, seed 42; **Table B9 reports the same matrix under 2-cell (≈ 1 km) blocking**, which is why its supported counts are the larger 12 and 6 (Section 4.5). Per-direction bounds are in `aoi_frame_transfer_frozen_mugla`.

CSV.

Source frame	Target frame	Mean target AUC	Above chance	Below chance	Supported above / below	Paired thermal delta
full	full (Table B9)	0.541	14 of 20	6	9 / 4	+0.004
full	10 km	0.576	17 of 20	3	11 / 1	+0.003
10 km	full	0.570	17 of 20	3	11 / 3	+0.014
10 km	10 km	0.616	19 of 20	1	15 / 1	+0.023
5 km	5 km	0.608	19 of 20	1	12 / 0	+0.014

0.022 and leaves **every headline quantity above unchanged to within 0.0012**, with the signed AUCs of that region’s two channels moving by at most 0.008 and no verdict changing (Appendix A(w)). The table reports the
530 corrected values, and reproduces the frozen matrix at 0.541 against Table B9’s 0.541 and 14 of 20 exactly. **The baseline control must be restated on this frame:** the static baseline transfers at 0.593 against 0.616, a paired difference of +0.023 rather than +0.004, so the control holds in kind but the gap is about six times larger once frames are comparable, and about four
535 times at the 5 km collar.

What this settles, and what it leaves standing. Five quantities reported below are properties of the frames rather than of the predictor-burning relationship, and are identified as such where they appear: the count of six anti-predictive directions, which becomes one; the sign reversal of ele-
540 vation, LST and TVDI as a mechanism; the sign-agreement diagnostic; the same-geography arm; and the paired thermal contribution, +0.004 as drawn against +0.023 equalised. That last quantity carries the portability null, so it is given an interval on the frame this section argues for: **+0.023 [−0.004, +0.048]** under the pair-cluster resampling of Appendix A(o). It still spans

545 zero, so the null survives the correction — but only just, and under the al-
 ternative admissible unit, clustering by target region, it does not (Appendix
 A(w)). What survives is the central negative result, and its size must be
 stated on a matched comparison: setting 0.616 against a within-region refer-
 ence of about 0.87 would compare a collar number with a full-rectangle one
 550 at 1 km blocking. Recomputed on the same frame at 5 km blocking that
 reference is 0.772, so the shortfall is **+0.155** [**+0.094**, **+0.217**], a Student
 t interval over the five target regions rather than the spatial-block bootstrap
 used elsewhere, and numerically almost identical to Section 4.3’s negative-
 pool effect over nine scars, which is a different quantity on a different unit.
 555 It is real, and 0.155 rather than the 0.25 an unmatched comparison implies.
 Appendix C.5(ix) records the frame as a limitation of this cohort rather than
 of the method.

4.5. *Cross-region transfer, and what label-free adaptation does to it*

Everything in this section is computed on the frames as drawn
 560 and should be read against Section 4.4, which has shown the frames
 are not comparable and that the within-region reference used here is itself
 frame-dependent. The as-drawn matrix is reported because it is what the
 pre-registered protocol yields; per-direction values are in Appendix B, Table
 B9, and the supporting arms in Appendix A(ix).

565 **Raw transfer is heterogeneous and includes anti-predictive di-
 rections.** Target AUC spans 0.326 to 0.686, and Fig. 4 gives the matrix
 direction by direction. The support counts depend on the blocking scale, so
 both are reported: at 2-cell (≈ 1 km) blocking twelve of twenty directions
 are above chance with interval support and six below; at the 10-cell (≈ 5

570 km) blocking this design defends, nine are above and four below with seven uncertain, the counts Table 3 uses. Of the below-chance directions only Man-avgat to Bejís survives frame equalisation, and two verdicts are not stable across bootstrap seeds (`transfer_ci_blocksize.csv`). Even the best raw transfer sits far below that target’s own within-region 0.870, the raw deficit
575 running 0.184 to 0.592 against a reference Section 4.3 shows is not matched to a transfer evaluation.




Figure 4: **Cross-region transfer, and what label-blind adaptation does to it.** Thermal transfer ROC-AUC for all 20 ordered source–target directions among the five regions, natural-vegetation primary population. **(a)** Raw transfer, no adaptation. **(b)** After region-wise z-score standardisation. **(c)** After CORAL. Rows are the source region, columns the target; the grey diagonal is within-region performance and is *not* a transfer result. The diverging colour scale is centred exactly on chance (0.5, marked by the black rule on the colour bar), so orange is below chance and purple above. Each panel header states that panel’s range: adaptation compresses the spread from 0.326–0.686 (raw) to 0.431 to 0.630 (z-score) and 0.443 to 0.624 (CORAL). It pulls the extremes toward chance rather than lifting transfer. Cells below chance are hatched. The hatch is a redundant cue for print: a diverging palette is symmetric in luminance about its centre, so in greyscale a below-chance cell would otherwise be indistinguishable from an equally distant above-chance one. Hatching follows the *exact* value while the printed label is rounded to two decimals, so two cells print “0.50” with different hatching: CORAL Bejís→Evia is 0.4991 (hatched, below chance) and z-score Evia→Muğla is 0.5013 (un-hatched, above chance). Every cell carries its value, and every label’s colour is chosen by measured WCAG contrast against that cell’s rendered colour (worst case 4.7:1). Confidence intervals for these directions are given in Table B9.

In precision terms it is worse than the ROC figures suggest. A susceptibility surface is used as a ranked area budget, so precision-recall is the operational quantity. PR-AUC averages **0.156 against a no-skill baseline**
 580 **of 0.136**, and **six of twenty directions fall below their own baseline**
at the point estimate, five of them with intervals entirely below it; the exception is Bejis to Manavgat, whose interval covers its baseline. Only one direction exceeds twice its baseline. These are frame-as-drawn quantities and the PR arm was not recomputed on the collar (Appendix B, Table B4).

585 **The static baseline does not transfer either**, at a mean of **0.537** against **0.541** for the thermal model, so the failure is not specific to the dynamic block: a baseline that does not travel, plus pre-fire thermal state, gives a model that does not travel. Section 4.4 restates this control on the equalised frame, where it is 0.593 against 0.616.

590 **The thermal block’s paired contribution is +0.004, and the mean is not the informative statistic.** The individual contributions span **−0.148 to +0.133**, twelve positive and eight negative, so the spread is thirty times the mean and the sign belongs to the pair rather than to the block; a mean near zero records cancellation, not consistent absence of effect.
 595 The directions are not independent, since each region appears in eight of the twenty, so the interval depends on the resampling unit — **all four units the design permits give the same answer**, from $[-0.027, +0.034]$ to $[-0.037, +0.046]$, and none propagates within-direction sampling variability. The leave-one-region-out jackknife shows how little the mean is anchored:
 600 **dropping Evia alone reverses its sign.**

Label-blind adaptation compresses the matrix toward chance

rather than repairing it (Fig. 5). Under region-wise z-scoring the twenty directions span 0.431 to 0.630 and under CORAL 0.443 to 0.624, roughly half the raw spread, and fourteen of twenty move closer to chance — which helps the directions that failed and harms those that worked. The committed-in-advance CORAL arm averages **0.552**; taking whichever method scores better per direction gives 0.556, but that selection uses the target labels the protocol forbids, so it is **an oracle upper bound rather than an achievable result**. Even the oracle only reaches the reference a model can reach on an unseen scar (Section 4.3), so alignment is regressing the matrix onto that reference rather than failing beneath it. What it cannot do is exceed the reference, and a sign reversal is not a distribution mismatch that realigning inputs would repair. Appendix A(ix) reports the recovery fractions, including the seven directions with *negative* recovery.

4.6. Further arms

Six further arms bear on the findings above without changing them: the contrast pair (A(s)), the two interventions (A(n)), the sensitivity summary (A(v)), the same-geography two-event arm (A(m)), the distance curve (A(t)) and the label-budget curve (A(u)). Each is stated there with its own limits. Three are plotted here because the shape of the result is the argument: pooling every other region never beats the best single source for any target (Fig. 7); removing the direction-reversing features costs within-region skill and returns nothing measurable on transfer (Fig. 6). The two bearing directly. The two bearing directly on Sections 5.4 and 5.7 are the interventions, where a local cost of -0.081 is measured against a transfer return of $+0.014$




Figure 5: **Label-blind adaptation regresses transfer onto the level a model reaches on an unseen scar, rather than recovering it.** All 12 ordered directions among the four regions carrying the full decomposition (Manavgat, Bejís, Muğla, Evia), natural-vegetation primary population. Each row is one direction. The open circle is raw transfer ROC-AUC and the arrow runs to the best label-blind adapted result for that direction (region-wise z-score or CORAL, whichever is better); the grey vertical tick is the within-region reference for that target, and the dashed vertical rule is chance. **Solid blue arrows are the five directions where adaptation moves transfer toward the within-region reference (recovery); dashed orange arrows are the seven where it moves away from it (negative recovery).** Line style and colour both encode this class, because the two hues differ by only 2.30:1 in relative luminance and would nearly merge in greyscale. The pattern is the point: all six directions starting above chance are pushed down, and five of the six starting below chance are pushed up, so adaptation acts on the spread rather than on skill. The exception is Manavgat→Muğla, which starts below chance at 0.470 and is pushed further down to 0.443. The worst case is Evia→Manavgat, where raw transfer of 0.686 is dragged down to 0.527 against a within-region reference of 0.870, a recovered fraction of -0.862 . No arrow reaches its within-region tick.

$[-0.017, +0.045]$ whose interval spans zero, and the label budget, where thirty-two labelled 5 km blocks recover 85 to 89 % of the target's matched ceiling in three of six directions and 30 to 57 % in the rest — 7 to 20 % of the target's population, drawn from the event being predicted, which is not

630 a resource available before that event burns.




Figure 6: Removing the direction-reversing features costs locally and returns nothing measurable on transfer. Mean within-region ROC-AUC across the five regions (horizontal) against mean transfer ROC-AUC across the 20 ordered directions (vertical), for four feature configurations: the full thermal set, the set without the LST anomaly, without elevation, and without both. Each configuration has its own marker and is named in the key; the grey path connects them in that order. Moving left along the path, transfer rises monotonically ($0.541 \rightarrow 0.556$) while within-region skill falls monotonically ($0.888 \rightarrow 0.807$). Both monotonicities are verified when the figure is built, so the figure cannot silently contradict the claim. The within-region cost is 0.081 AUC and is supported in every region. The change in transfer is +0.014 AUC on an interval that spans zero, so no exchange between the two is demonstrated: a local cost is measured and no compensating transfer gain is.




Figure 7: **Pooling every other region never beats the best single source.** Leave-one-region-out evaluation: for each held-out target region (rows), a model trained on the pooled remaining four regions (filled orange circle) is compared against the best individual source region for that same target (open blue diamond), with the within-region ceiling as a grey vertical tick and chance as the dashed rule. Natural-vegetation primary population, thermal feature set, no adaptation. The pooled model lies to the left of the best pairwise source in all five targets without exception. Adding source regions does not accumulate transferable signal. In two targets the pooled model falls *below chance* at the point estimate (Manavgat 0.469, Bejís 0.417), and only for Bejís does the interval lie entirely below chance, while a single well-chosen source reaches 0.686 and 0.583 respectively. Every point remains far below the within-region ceiling.

5. Discussion

5.1. Reading the two findings together

Section 1.3 states the two findings and Section 4 establishes them; this section argues from them. The relation between them is what makes the paper cohere: the first is not a caveat attached to the second, it is the
635 instrument that sets its size. Applied to our own matrix it withdrew five claims, and what it left standing is a shortfall in transferred skill, not a reversed relationship.

5.2. Why the thermal increment is real but local

640 The within-region increment and the transfer failure are measured at different separations, and Section 4.3 shows most of the difference is already present inside a single region. The increment is local in a specific sense: it holds where held-out cells are interleaved with training cells and most of it is gone once they are not, before the fire or the region changes. It is not an
645 artefact to be explained away, surviving every sensitivity arm of Appendix A(a)–(h), but it is established under interleaved validation and not beyond it — and even there the window-closure arm **weakens monotonically in Evia**, so what holds everywhere is survival, not improvement.

The natural objection is that Mediterranean regions are simply different
650 systems, so a predictor meaning one thing in one place and another elsewhere is a comparison of two systems rather than instability. **We designed the two-Muğla-events arm to answer that objection and it does not answer it:** Section 4.4 shows it is the most extreme frame artefact in the cohort. Three further confounds were never resolved in any case: season and

655 year, which **cannot be resolved in this study area**, an unfixed population, and a positive-block count below this design’s own floor (Appendix C.5(ii)). With one fire per region everywhere else and that arm withdrawn, **this cohort provides no evidence that the transfer shortfall is regional rather than event-specific**. Appendix A(t) adds a length scale for
660 the within-region decay but cannot turn it into an attribution either.

One conclusion survives from the other direction: the static baseline transfers no better than the dynamic block, so whatever the shortfall is, it is not the thermal block’s peculiarity.

5.3. *What the two interventions do and do not show*

665 Both interventions show the same shape. Pooling four regions never beats the best single-source transfer for any target and stays well below the within-region ceiling (Appendix A(n)), so it does not manufacture the missing conditional information. Removing the reversing predictors costs -0.081 of within-region skill with interval support in every region, and returns $+0.014$ $[-0.017,$
670 $+0.045]$ in mean transfer.

That pair of numbers is easy to read as an exchange, and it is not one. Both arms are null on the portability axis — the thermal block’s own contribution is $+0.004$, removal returns $+0.014$, both intervals spanning zero — so they measure a local cost and no compensating gain. That is not a conservation law and not a rate at which local skill can be sold for portability;
675 no such rate is estimated here.

5.4. *The regime hypothesis, reported as it happened*

A regime-structure explanation was stated in advance and the data confirmed the null: the regime-distance correlation has the wrong sign at the point estimate, the most regime-similar pair fails in both directions, and the most regime-different pair transfers above chance. One mundane explanation can be set aside: subsampling Muğla, much the largest population, to Manavgat’s cell count and separately to its positive count leaves its transfer behaviour inside the subsampling range in both roles. The error was in the hypothesised grouping, not in the data, and with ten pairs this cannot refute regime typology (Archibald et al., 2013) in general.

5.5. *The empirical contrast with Dimarco et al.*

Dimarco et al. (2026) transfer successfully across a comparable Mediterranean design and we do not, and the two results are not in conflict. Their predictors are attributes of a place, ours the state of a surface in one season, which suggests the relation between domain similarity and transfer success is predictor-class dependent.

That reading is one of at least two, and we cannot separate them here: their response variable is human-driven **ignition**, dominated by access and activity, while ours is burned **area**, dominated by spread, so a predictor-class explanation and a response-variable explanation are confounded in this comparison. Our own data speak against a simple predictor-class reading in any case: the static baseline transfers at a mean of 0.537 here, so within this cohort the place-attribute class does not travel either.

700 5.6. *Implications*

In precision-recall terms, which is how a susceptibility surface is used, transferred models average a PR-AUC of 0.156 against a no-skill baseline of 0.136 **on the frames as drawn**, six of twenty falling below their own baseline (Section 4.5). Whatever the ROC figures suggest, a model moved
705 to a region it was not fitted in does not usefully rank burned cells there.

The paper supports one concrete change in reporting: alongside a spatially blocked within-region figure, report skill on a held-out burn scar and its surroundings. On these five regions the two differ by about 0.14 ROC-AUC on the same model, the size of the effect such papers usually claim, so
710 **a blocked figure alone should be read as an upper bound**. Transfer skill likewise has to be *measured* on the target region rather than assumed from a within-region figure, and where a model must be moved the resource that closes the gap is target labels (Appendix A(u)): a real answer, but not a cheap one.

715 The results also bound what unsupervised alignment can be asked to do. Even the oracle selection, which uses the target labels the protocol forbids, only reaches the reference a model can reach on an unseen scar (Section 4.5). Alignment is not failing far below an achievable target; it is regressing the matrix onto it, at the cost of the directions that already worked. A sign
720 reversal is not a distribution mismatch that realigning inputs repairs.

5.7. *Limitations*

Ten limitations are stated in full in Appendix C.5; four bind the conclusions above. **The frames are not comparable and this cohort cannot fully repair it** (Section 4.4): the collar radius is itself a choice — 5 km and

725 10 km give 0.608 against 0.616 — and the frames were fixed upstream, so
 we can restrict them but not extend them. Any future cohort should fix the
 frame by an explicit accessible-area rule (Barve et al., 2011) before any pre-
 dictor is computed; that is the main design lesson of this paper. **Each region**
contributes one fire season, so regional concept shift is confounded with
 730 event meteorology and the shortfall cannot be attributed to region rather
 than event (Section 5.2). **The diagnostic correlations rest on an effec-**
tive sample of ten region pairs, so Section 4.6’s successes and failures
 both read at that power. And **the interval-support counts are less sta-**
ble than the point estimates behind them: the published split turns
 735 on a lower bound of -0.00045 . The remaining six, from the absent meteo-
 rological covariates to the single classifier family (Appendix A(h)), bear on
 scope rather than on the conclusions above.

6. Conclusions

Where a fire model is scored decides what it appears to know. Holding
 740 model, predictors and fitting fixed and changing only which cells are scored,
 moving from a whole study region to the burn scar and its 2 km collar
 costs **0.143 ROC-AUC**, with a control measuring the cause as the negative
 pool’s composition rather than class balance. That is the size of the predictor-
 block increments this design itself measures, so a region-wide validation figure
 745 should be read as an upper bound on what a model achieves where fire
 actually occurs, and reported as one.

That correction is not only other people’s problem. Applied to our own
 five-region matrix it withdraws five of our claims. The study areas enclose

very unequal far fields, and equalising them to a 10 km collar lifts mean
750 transfer from 0.541 to 0.616, reduces the directions below chance, and leaves
no sign reversal supported between regions under our own criterion. The
same correction applies to the one arm that held place fixed, two fires in one
study area eleven months apart, which we had wrongly exempted because
its geography was constant while its evaluation frame was not: its elevation
755 reversal does not survive equalisation either, and that arm can speak only
for the year-invariant channels. What had looked like a reversed predictor-
burning relationship was mostly a statement about how five rectangles were
drawn.

What survives is the central result. Pre-fire thermal dryness adds a real
760 and repeatable within-region increment in all five regions, surviving spatial
blocking at about 5 km and a predictor window closed up to two weeks before
the first labelled burning. It does not travel: on matched frames and blocking,
equalised transfer of 0.616 falls 0.155 short of the within-region reference;
the paired cross-region contribution spans zero on both frames under the
765 resampling unit we treat as primary, though not under every admissible one
(Section 4.4), with a sign that varies by pair; and the static baseline transfers
no better on either frame, so this is not a peculiarity of thermal predictors.
Label-free alignment compresses most directions towards chance rather than
repairing them, and supervised recalibration works but is not cheap.

770 Two limits belong with all of this. The agreed thermal direction is the
opposite of the one the dryness framing predicts — hotter pre-fire surfaces
burned less in every region, and holding greenness does not remove it while
holding temperature reverses greenness in two regions — so on this cohort

the absolute channels behave as static land-surface descriptors rather than
775 as a dryness index. And with one fire per region, and the single arm that
held place fixed now withdrawn, this design measures a transfer shortfall but
cannot attribute it to region rather than to event.

Two practical consequences follow. Report local skill and portability sep-
arately, since a predictor block worth a large within-region increment may
780 contribute nothing distinguishable from zero across regions. And fix the
evaluation frame by an explicit accessible-area rule before any predictor is
computed, because a frame chosen for convenience can manufacture both
a transfer failure and a mechanism to explain it. Three extensions are left
for future work: temporal transfer, meteorological covariates, and physically
785 normalised dryness variables.

Appendix A. Appendix A. Where the rest of the appendices are

Two appendices are printed here: **Appendix B**, the diagnostic tables
below, and **Appendix C.5**, the eleven limitations. They carry the per-
direction numbers and the limitations against which the claims are checked,
790 so the paper can be assessed without leaving it.

Appendix A and the protocol sections C.1 to C.4, C.6 and C.7
are released with the paper rather than printed in it, at paper/
supplementary_appendices.md in the repository named in the declarations.
They keep their section names there, so a pointer such as Appendix A(c) or
795 Appendix C.2 in the text resolves in that document unchanged. Nothing in
them is evidence a claim depends on that appears nowhere else; every number
in them also sits in a frozen artefact named in the text. The lettering of this

appendix is kept for that reason: renaming would have broken 148 references
between the paper, the supplementary appendices and the supplementary
800 material.

Appendix B. Appendix B. Supporting tables

These are the per-region and per-direction numbers the Results sections
quote, given in full so that every claim can be checked against the values it
rests on rather than against a summary of them.

805 **Table B2. Signed univariate AUC of each predictor against
burned, by region.** Primary natural-vegetation population; 10-cell (~ 5
km) spatial-block bootstrap, 1000 replicates, seed 42. The AUC is never
folded to $\max(\text{AUC}, 1 - \text{AUC})$, so a value below 0.5 means lower values
rank burned and is a direction rather than weakness. **Bold** marks a region
810 whose own interval excludes 0.5.

Feature	Manavgat	Bejís	Muğla	Evia	Montiferru
elevation_	0.374 [0.289,	0.643 [0.558,	0.611 [0.532,	0.541 [0.448,	0.584 [0.395,
mean	0.471]	0.729]	0.690]	0.626]	0.762]
slope_mean	0.531 [0.423,	0.521 [0.439,	0.637 [0.582,	0.487 [0.418,	0.652 [0.506,
	0.642]	0.605]	0.686]	0.554]	0.771]
ndvi_mean	0.636 [0.587,	0.559 [0.497,	0.662 [0.616,	0.639 [0.575,	0.586 [0.450,
	0.676]	0.619]	0.704]	0.701]	0.704]
lst_anomaly_	0.482 [0.428,	0.418 [0.364,	0.485 [0.395,	0.640 [0.567,	0.395 [0.285,
mean	0.530]	0.480]	0.566]	0.710]	0.535]
current_lst_	0.538 [0.452,	0.477 [0.401,	0.325 [0.271,	0.377 [0.301,	0.370 [0.248,
mean	0.621]	0.547]	0.382]	0.456]	0.513]
current_	0.552 [0.460,	0.517 [0.429,	0.336 [0.275,	0.362 [0.285,	0.356 [0.233,
tvdi_mean	0.641]	0.595]	0.398]	0.442]	0.499]
tvdi_	0.449 [0.391,	0.512 [0.443,	0.490 [0.396,	0.519 [0.444,	0.378 [0.282,
difference_	0.505]	0.583]	0.575]	0.589]	0.497]
mean					
downscaled_	0.552 [0.466,	0.484 [0.400,	0.307 [0.253,	0.377 [0.297,	0.365 [0.240,
lst_mean	0.637]	0.560]	0.366]	0.459]	0.511]
fused_lst_	0.540 [0.454,	0.481 [0.404,	0.325 [0.272,	0.376 [0.300,	0.370 [0.248,
mean	0.622]	0.551]	0.383]	0.456]	0.513]

What counts as a reversal. A pair of regions is called a reversal only when their signed associations point to opposite sides of 0.5 **and each region's own interval excludes 0.5**. That is stricter than requiring the two regions' intervals to be disjoint, and the difference matters: for **current_** **lst_mean** between Manavgat and Muğla the two intervals are disjoint, but Manavgat's own interval, [0.452, 0.621], includes 0.5, so Manavgat has no established direction to reverse from. That pair is a point reversal, not a supported one.

Table B3. The cross-region reversals that meet the stricter criterion. Difference intervals are from the same paired bootstrap.

Feature	Region A	AUC	Region B	AUC	Difference [95 % CI]
<code>elevation_mean</code>	Manavgat	0.374	Bejís	0.643	+0.269 [+0.138, +0.397]
<code>elevation_mean</code>	Manavgat	0.374	Muğla	0.611	+0.235 [+0.102, +0.360]
<code>lst_anomaly_mean</code>	Bejís	0.418	Evia	0.640	+0.221 [+0.123, +0.313]

Three pair-level reversals across **two** features, which is why Section 3.11 removes exactly those two. Twenty-nine further pairs reverse at the point
825 estimate only, spread across eight of the nine features, and they are not counted. The conservative criterion costs the paper findings rather than manufacturing them: a difference interval on the pair, which is the instrument Appendix A(m) uses, would support more reversals than the three listed here.

Appendix B.1. B4. The transfer matrix in precision-recall terms

830 ROC-AUC is reported throughout the main text for comparability with the susceptibility literature. A susceptibility surface is used as a ranked area budget, so precision-recall is the operational quantity, and at target prevalences of 3.8 to 28.7 % the two can differ sharply. Read from the same frozen step9b exports as Table B9.

835 **Table B4. Thermal transfer, PR-AUC against the no-skill baseline.** The baseline is the target’s burned prevalence. Lift is PR-AUC divided by that baseline; a lift of 1 is no better than random ranking. Ordered by lift.

Direction	ROC-AUC	PR-AUC	No-skill	Lift
Evia $\rightarrow$ Manavgat	0.686	0.094	0.038	2.45
Bejís $\rightarrow$ Montiferru	0.594	0.289	0.212	1.36
Evia $\rightarrow$ Montiferru	0.647	0.283	0.212	1.34
Bejís $\rightarrow$ Muğla	0.619	0.093	0.070	1.33
Muğla $\rightarrow$ Evia	0.653	0.379	0.287	1.32
Montiferru $\rightarrow$ Bejís	0.548	0.093	0.072	1.28
Montiferru $\rightarrow$ Muğla	0.619	0.089	0.070	1.27
Evia $\rightarrow$ Muğla	0.577	0.086	0.070	1.23
Muğla $\rightarrow$ Bejís	0.583	0.088	0.072	1.22
Manavgat $\rightarrow$ Evia	0.613	0.343	0.287	1.20
Manavgat $\rightarrow$ Montiferru	0.533	0.243	0.212	1.15
Montiferru $\rightarrow$ Manavgat	0.567	0.043	0.038	1.11
Montiferru $\rightarrow$ Evia	0.586	0.317	0.287	1.11
Muğla $\rightarrow$ Montiferru	0.531	0.214	0.212	1.01
Bejís $\rightarrow$ Manavgat	0.444	0.034	0.038	0.90
Manavgat $\rightarrow$ Muğla	0.470	0.063	0.070	0.90
Evia $\rightarrow$ Bejís	0.448	0.059	0.072	0.82
Bejís $\rightarrow$ Evia	0.383	0.226	0.287	0.79
Muğla $\rightarrow$ Manavgat	0.401	0.029	0.038	0.75
Manavgat $\rightarrow$ Bejís	0.326	0.049	0.072	0.67
Mean	0.541	0.156	0.136	1.16

840 Six directions fall below their own no-skill baseline, and only one exceeds twice it. The six are the same six that are below chance on ROC-AUC, which is what a reversed ranking predicts in either metric. Section 4.4 shows that this count is largely a property of the evaluation frames rather than of

a reversed predictor-burning relationship.

845 *Appendix B.2. B5. The same-geography event pair, in full*

Appendix A(m) reports this arm and Section 4.4 withdraws its **elevation** reversal as a frame artefact; Appendix A(o) explains why the thermal channels of this arm carry no verdict either way. The per-feature values are kept here because the arm is what motivated the frame test, and because its
850 structural asymmetries have no analogue in the twenty-direction matrix.

Table B5. Signed univariate feature-burned AUC, Muğla 2021 versus 2022. Raw AUC against burned, never folded to $\max(\text{AUC}, 1 - \text{AUC})$; 10-cell (≈ 5 km) spatial-block bootstrap, 1,000 replicates, seed 42 (Section 3.14). Analysis population 41,730 rows / 2,911 burned (2021) and
855 38,790 rows / 331 burned (2022). **Positive-carrying 5 km blocks: 70 for the 2021 arm and 11 for the 2022 arm.** Table 1’s note sets sixteen as the floor this design supports at that blocking, so the 2022 intervals here fall below the paper’s own standard and are read as indicative, exactly as the 20-cell row of Table 1 is. The 2022 arm is additionally a single
860 compact scar, so its eleven blocks are contiguous. Read from `paper/step9g_raw/.../mugla_2021__mugla_2022_event_relative/step9g_direction_reversal_table.csv`.

Feature	2021 AUC [95 % CI]	2022 AUC [95 % CI]	Reversal
elevation_mean	0.611 [0.532, 0.690]	0.296 [0.230, 0.355]	bootstrap- supported
current_lst_mean	0.325 [0.271, 0.382]	0.515 [0.434, 0.580]	point only
current_tvdi_mean	0.336 [0.275, 0.398]	0.594 [0.475, 0.674]	point only
downscaled_lst_mean	0.307 [0.253, 0.366]	0.508 [0.435, 0.571]	point only
fused_lst_mean	0.325 [0.272, 0.383]	0.519 [0.436, 0.583]	point only
ndvi_mean	0.662 [0.616, 0.704]	0.707 [0.624, 0.777]	none
slope_mean	0.637 [0.582, 0.686]	0.558 [0.468, 0.634]	none
lst_anomaly_mean	0.485 [0.395, 0.566]	0.380 [0.249, 0.502]	none
tvdi_difference_mean	0.490 [0.396, 0.575]	0.397 [0.265, 0.514]	none

Appendix B.3. B6. The evaluation frames, and the signed associations they produce

Section 4.4 states these results; the per-region values are here.

Table B6. Evaluation-frame geometry of the five study regions.

Primary natural-vegetation population. Distance is Euclidean to the nearest burned cell on the 500 m grid, at 0.45 km per cell. Source `aoi_frame_auc.`

870 `csv`; recomputable by `paper/code/verify_aoi_frame.py`.

Region	cells	burned	median distance to burned	share beyond 10 km
Manavgat	20,511	784	13.4 km	60.1 %
Bejís	15,190	1,100	13.5 km	63.1 %
Muğla	41,730	2,911	11.3 km	55.3 %
Evia	9,298	2,664	8.0 km	43.7 %
Montiferru	2,544	539	2.7 km	2.1 %

Table B7. Signed univariate AUC, frame as drawn against a 10 km collar. Point estimates; the intervals that decide the reversal question are given in the text below and in `collar_frame_bootstrap.csv` (10-cell
875 blocks, 1000 replicates, seed 42). Signed and never folded to $\max(\text{AUC}, 1 - \text{AUC})$, so a value below 0.5 is a direction, not weakness. The collar drops no burned cells in any region.

Signed univariate AUC	Manavgat	Bejís	Muğla	Evia	Montiferru	straddles 0.5
elevation, full frame (Table B2)	0.374	0.643	0.611	0.541	0.584	yes
elevation, 10 km collar	0.561	0.614	0.606	0.648	0.581	no
current LST, full frame	0.538	0.477	0.325	0.377	0.370	yes
current LST, 10 km collar	0.386	0.405	0.332	0.286	0.376	no
current TVDI, full frame	0.552	0.517	0.336	0.362	0.356	yes
current TVDI, 10 km collar	0.392	0.454	0.342	0.250	0.361	no

Table B8. Region summary: populations and gate outcomes.

880 Counts from each region’s Step 8A dataset statistics; gate fractions from
each region’s burned-landcover gate output. TSG = the primary natural-
vegetation population. The TSG columns use the canonical modelled pop-
ulation, `burnable_tree_shrub_grass` and `valid_for_modeling == True`,
which is the population every model in this paper was fitted and scored on.

885

Region	Total cells	Valid cells	Burned	Prevalence (all valid)	TSG cells	Burned in TSG	TSG prevalence	Burned natural-veg fraction	Gate verdict
Manavgat 2021	24,150	24,087	796	0.033	20,511	784	0.038	0.984	pass
Bejís 2022	15,759	15,759	1,103	0.070	15,190	1,100	0.072	0.991	pass
Muğla 2021	73,098	73,045	3,026	0.041	41,730	2,911	0.070	0.958	pass
North Evia 2021 (extended)	22,925	22,906	2,788	0.122	9,298	2,664	0.287	0.945	pass
Montiferru 2021	3,234	3,173	697	0.220	2,544	539	0.212	0.723	pass

Table B9. Cross-region transfer matrix, thermal model, TSG population. Target ROC-AUC with 2-cell spatial-block bootstrap 95% CIs (1000 replicates). CORAL is applied after region-wise z-scoring ($\lambda = 10^{-5}$).

Direction	Raw	Region-wise z-score	CORAL
Manavgat→Bejís	0.326 [0.305, 0.349]	0.477 [0.451, 0.502]	0.511 [0.484, 0.534]
Bejís→Manavgat	0.444 [0.408, 0.480]	0.457 [0.420, 0.497]	0.555 [0.528, 0.583]
Manavgat→Muğla	0.470 [0.451, 0.488]	0.431 [0.411, 0.449]	0.443 [0.423, 0.461]
Muğla→Manavgat	0.401 [0.378, 0.426]	0.559 [0.531, 0.587]	0.560 [0.535, 0.587]
Manavgat→Evia	0.613 [0.593, 0.631]	0.542 [0.520, 0.565]	0.539 [0.518, 0.561]
Evia→Manavgat	0.686 [0.653, 0.716]	0.516 [0.489, 0.544]	0.527 [0.500, 0.553]
Bejís→Muğla	0.618 [0.601, 0.635]	0.518 [0.501, 0.535]	0.507 [0.489, 0.524]
Muğla→Bejís	0.583 [0.561, 0.607]	0.535 [0.512, 0.557]	0.560 [0.538, 0.581]
Bejís→Evia	0.383 [0.363, 0.402]	0.532 [0.509, 0.551]	0.499 [0.479, 0.518]
890 Evia→Bejís	0.448 [0.426, 0.470]	0.549 [0.524, 0.575]	0.549 [0.525, 0.573]
Muğla→Evia	0.653 [0.636, 0.671]	0.561 [0.543, 0.580]	0.563 [0.545, 0.582]
Evia→Muğla	0.577 [0.560, 0.593]	0.501 [0.485, 0.518]	0.530 [0.515, 0.546]
Montiferru→Manavgat	0.567 [0.539, 0.594]	0.573 [0.538, 0.609]	0.606 [0.574, 0.639]
Manavgat→Montiferru	0.533 [0.488, 0.580]	0.586 [0.540, 0.629]	0.592 [0.550, 0.631]
Montiferru→Bejís	0.548 [0.521, 0.578]	0.574 [0.552, 0.596]	0.569 [0.548, 0.591]
Bejís→Montiferru	0.594 [0.560, 0.631]	0.550 [0.500, 0.601]	0.574 [0.530, 0.621]
Montiferru→Muğla	0.619 [0.604, 0.634]	0.576 [0.562, 0.589]	0.565 [0.550, 0.579]
Muğla→Montiferru	0.531 [0.495, 0.568]	0.587 [0.549, 0.624]	0.584 [0.547, 0.623]
Montiferru→Evia	0.586 [0.565, 0.606]	0.630 [0.611, 0.649]	0.624 [0.605, 0.641]
Evia→Montiferru	0.647 [0.608, 0.682]	0.568 [0.528, 0.609]	0.581 [0.539, 0.623]

Appendix C. Appendix C. Limitations

Section 5.7 states the four limitations that bind the conclusions; all ten are here. The rest of this appendix, which specifies the protocol, is released with the supplementary appendices rather than printed here.

895 *Appendix C.1. C.5 Limitations, in full*

Section 5.7 states the four that bind the conclusions; all ten are here.

(i) **No meteorological covariates** enter the models, so we cannot say how local skill and portability behave for a mixed thermal-plus-weather predictor set.

900 (ii) **The same-geography comparison covers one region only**, and even there year and seasonal phase are confounded. That confound cannot be resolved in this study area, and the reason is specific: the two events sit 42 days apart in median burn day-of-year, neither year contains a second event at the other’s phase, and a calendar-matched arm would carry nine
905 burned cells against this design’s gate minimum of thirty. The positive-block count for the 2022 arm is separately below the floor this design sets itself, at eleven against sixteen (Appendix A(m)). Its 331 burned cells also leave the thermal reversals unresolved at interval level, and the pair holds place fixed but not population. Those two arms are additionally the only transfer
910 directions here computed by us rather than read from the pipeline author’s frozen export, with his unmodified code and the same pinned environment.

(iii) **All labels derive from a single burned-area product**, MCD64A1 (Giglio et al., 2018), whose omission and commission characteristics (Boschetti et al., 2019) bound every model evaluated here. No second

915 burned-area product covers 2021 and 2022 at this resolution; the companion paper reports what an independent active-fire observation says about the omission concern.

(iv) **Evia remains the most imbalance-atypical population** even after the AOI extension, at a TSG prevalence of 0.287 against 0.038 to 0.072
920 in three of the others.

(v) **Each region contributes one fire season**, so regional concept shift is confounded with event meteorology, and distinguishing them requires multi-year labels.

(vi) **Cross-region point estimates carry an implementation tolerance** of roughly ± 0.02 to 0.03 across scikit-learn versions. All reported
925 numbers are fixed to one verified version, but exact reproduction elsewhere requires the archived environment.

(vii) **Manavgat’s atypical transfer behaviour remains unexplained.** It is where the conditional diagnosis bites hardest and where
930 feature removal recovers most. Three candidates have now been tested and none survives: its meteorology, which was not extreme; the quality screening of its coarse thermal input, which propagates widely but moves no signed association by more than +0.0003 (Appendix A(v), Appendix A(e)); and the evaluation frame, which explains its elevation figure but not its transfer
935 behaviour (Section 4.4). With one fire season per region the remaining candidates are not separable in this design.

(viii) **The interval-support counts are less stable than the point estimates behind them.** Several verdicts sit within a thousandth of their reference value, and at 1 km blocking the published split of ten positive, seven

940 negative and three uncertain turns on a lower bound of -0.00045 . The point estimates and the sign pattern are stable; the counts are not. Every sentence in this paper that leans on an exact count of supported directions should be read at that precision. (ix) **The five areas of interest are not comparable frames, and this cohort cannot fully repair it** (Section 4.4).

945 We report the equalised arm alongside the frame-as-drawn arm rather than replacing one with the other, because the collar radius is itself a choice and 5 km and 10 km do not agree exactly (0.608 against 0.616). The deeper limitation is that the frames were fixed upstream of this work, in `repo/`, so we can restrict them but not extend them; a region whose rectangle is already

950 fire-scale, Montiferru, cannot be given a far field for symmetry. Nor can their independence from outcomes be fully documented: only Montiferru’s box is derived by a rule, and for Manavgat, Bejís and Muğla the version history cannot show the box fixed before the first gate result (Section 3.1). Any future cohort should fix the frame by an explicit accessible-area rule (Barve et al., 2011) before any predictor is computed, and we treat that as the main

955 design lesson of this paper.

(x) **One classifier family.** The headline numbers use a random forest with unlimited depth, the configuration most able to encode local structure and least able to extrapolate. Appendix A(h) shows the transfer result is not

960 an artefact of that choice: three further estimators, including a penalised linear one, all land between 0.510 and 0.556 and all place fourteen of twenty directions above chance. Those are point estimates without intervals, so the ordering among them is not claimed as a result. Other model families were not tried, and a different inductive bias might behave differently, but within

965 this family the negative result is a property of the predictors rather than of
an unregularised estimator.

Software and data availability

Analysis code and frozen outputs. Name: `thermal-twin`. De-
velopers: Y. E. Cogurcu and E. Metin; contact: `ycogurcu@cu.edu.tr`.
970 First available: 2026. Hardware: a standard desktop computer. Software
required: Python 3.12 with NumPy 2.4.4, pandas 3.0.2 and scikit-learn
1.9.0, the versions every reported number was produced with; scikit-learn in
particular carries a cross-region tolerance of about ± 0.02 to 0.03 between
versions (Appendix C.5(vi)). Program language: Python. Size: about
975 21 MB. Availability: <https://github.com/dryuemco/thermal-twin>,
which holds the scripts that regenerate each reported artefact, the frozen
numeric outputs behind every table and figure, and the supplementary
appendices (Appendix A and protocol sections C.1 to C.4, C.6 and C.7, at
`paper/supplementary_appendices.md`). Cost: free.

980 **Upstream processing pipeline.** Name: `satellite-thermal-
digital-twin`. Developer: E. Metin. First available: 2026. Hard-
ware: a standard desktop computer; the satellite exports run on
Google Earth Engine and need an Earth Engine account. Pro-
gram language: Python. Size: about 20 MB. Availability: [https:](https://github.com/emrehann17/satellite-thermal-digital-twin)
985 [//github.com/emrehann17/satellite-thermal-digital-twin](https://github.com/emrehann17/satellite-thermal-digital-twin), MIT
licence; the commit of record for every number reported here is `6381f4c`.
Three components once outside the release are now inside it: the
reproduction-check driver and its validation package, on which Section 3.14

rests; the few-shot run, reachable through the tag `few-shot-run-19d825b`;
990 and the ERA5-Land diagnostic, whose manifest records commit `a07ea33`.
Cost: free.

Data. All satellite inputs are public and were retrieved through Google Earth Engine: burned-area labels from MODIS MCD64A1 Collection 6.1, land cover from ESA WorldCover, and the thermal, optical and terrain inputs
995 described in Section 3.4; no proprietary or restricted data were used. The modelling dataset each number rests on is identified in the pipeline by SHA-256, which is how the one provenance incident in this project was settled: a rebuild had replaced one region’s file at its canonical path, and the recorded hash identifies the frozen original unambiguously. No digital object identifier
1000 is minted and no archival deposit exists.

Declarations

Funding. This work was supported by the Çukurova University Scientific Research Projects Coordination Unit (Bilimsel Araştırma Projeleri Koordinasyon Birimi) under the Career Starter Project (Kariyer Başlangıç
1005 Projesi) scheme, project code `FKB-2025-17608` (“Termal Dijital İkiz Tabanlı Sürü İHA Sistemi ile Orman Yangınlarının Erken Tespiti ve Önlenmesi”).

Acknowledgments. The authors gratefully acknowledge the Çukurova University Scientific Research Projects Coordination Unit for financial support of this research, and the Department of Computer Engineering
1010 at Çukurova University for providing the laboratory environment and institutional support that made this work possible.

Competing interests. The authors declare no competing interests.

Ethics approval. Not applicable. This study involved no human participants, animal subjects, or personally identifiable data.

1015 **Author contributions.** Stated in CRediT terms. **Emrehan Metin:** Software, Data curation, Investigation, Validation, Writing – review and editing. **Yunus Emre Cogurcu:** Conceptualization, Methodology, Formal analysis, Investigation, Writing – original draft, Writing – review and editing, Supervision.

1020 **Declaration of generative AI and AI-assisted technologies in the manuscript preparation process**

During the preparation of this work the authors used Claude (Anthropic) in order to write and run analysis and verification code against the frozen pipeline outputs, cross-check reported numbers against those outputs, and
1025 draft and edit manuscript text. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the publication.

References

Alkan Akıncı, H., Akıncı, H., 2023. Machine learning based forest fire susceptibility assessment of Manavgat district (Antalya), Turkey. Earth Science
1030 Informatics 16, 397–414. doi:10.1007/s12145-023-00953-5.

Archibald, S., Lehmann, C.E.R., Gómez-Dans, J.L., Bradstock, R.A., 2013. Defining pyromes and global syndromes of fire regimes. Proceedings of the National Academy of Sciences 110, 6442–6447. doi:10.1073/pnas.
1035 1211466110.

- Barve, N., Barve, V., Jiménez-Valverde, A., Lira-Noriega, A., Maher, S.P., Peterson, A.T., Soberón, J., Villalobos, F., 2011. The crucial role of the accessible area in ecological niche modeling and species distribution modeling. *Ecological Modelling* 222, 1810–1819. doi:10.1016/j.ecolmodel.2011.02.011.
- 1040 Boschetti, L., Roy, D.P., Giglio, L., Huang, H., Zubkova, M., Humber, M.L., 2019. Global validation of the Collection 6 MODIS burned area product. *Remote Sensing of Environment* 235, 111490. doi:10.1016/j.rse.2019.111490.
- 1045 Breiman, L., 2001. Random forests. *Machine Learning* 45, 5–32. doi:10.1023/A:1010933404324.
- Chuvieco, E., Cocero, D., Riaño, D., Martín, P., Martínez-Vega, J., de la Riva, J., Pérez, F., 2004. Combining NDVI and surface temperature for the estimation of live fuel moisture content in forest fire danger rating. *Remote Sensing of Environment* 92, 322–331. doi:10.1016/j.rse.2004.01.019.
- 1050 Dimarco, N.A., Faraji, I., Wahbi, M., Maatouk, M., Boulaassal, H., Yazidi Aalaoui, O., El Kharki, O., 2026. Generalizing human-driven wildfire ignition models across Mediterranean regions using harmonized remote-sensing and machine-learning data. *Geomatics* 6, 13. doi:10.3390/geomatics6010013.
- 1055 Gelabert, P.J., Jiménez-Ruano, A., Ochoa, C., Alcasena, F., Sjöström, J., Marrs, C., Ribeiro, L.M., Palaiologou, P., Bentu Martínez, C., Chuvieco, E., Vega-García, C., Rodrigues, M., 2025. Assessing human-caused wildfire

- ignition likelihood across Europe. *Natural Hazards and Earth System Sciences* 25, 4713–4729. doi:10.5194/nhess-25-4713-2025.
- Giglio, L., Boschetti, L., Roy, D.P., Humber, M.L., Justice, C.O., 2018. The Collection 6 MODIS burned area mapping algorithm and product. *Remote Sensing of Environment* 217, 72–85. doi:10.1016/j.rse.2018.08.005.
- Gorelick, N., Hancher, M., Dixon, M., Ilyushchenko, S., Thau, D., Moore, R., 2017. Google Earth Engine: Planetary-scale geospatial analysis for everyone. *Remote Sensing of Environment* 202, 18–27. doi:10.1016/j.rse.2017.06.031.
- Huang, Y., Liu, N., Chen, W., Tan, W., Deng, Y., Wang, Z., Chlus, A., Shen, J., Townsend, P.A., 2026. Enhancing transferability of foliar trait retrieval models: A comparative analysis of transfer learning strategies and domain shift characterization. *Remote Sensing of Environment* 337, 115336. doi:10.1016/j.rse.2026.115336.
- Iban, M.C., Sekertekin, A., 2022. Machine learning based wildfire susceptibility mapping using remotely sensed fire data and GIS: A case study of Adana and Mersin provinces, Turkey. *Ecological Informatics* 69, 101647. doi:10.1016/j.ecoinf.2022.101647.
- Jodhani, K., Bhatiya, L., Sirvi, L., Chanda, A., Thakkar, S., Rathnayake, U., 2026. Artificial intelligence and remote sensing frameworks for wildfire monitoring and risk analysis across multiple ecosystems: A review. *AIMS Environmental Science* 13, 56–98. doi:10.3934/environsci.2026004.

- Kondylatos, S., Prapas, I., Camps-Valls, G., Papoutsis, I., 2023. Meso-geos: A multi-purpose dataset for data-driven wildfire modeling in the Mediterranean, in: Advances in Neural Information Processing Systems 36 (NeurIPS 2023) Datasets and Benchmarks Track, pp. 50661–50676. doi:10.52202/075280–2204.
- 1085
- Liu, C., Mostafavi, A., 2025. WildfireGenome: Interpretable machine learning reveals local drivers of wildfire risk and their cross-county variation. arXiv preprint arXiv:2511.11589. doi:10.48550/arXiv.2511.11589, arXiv:2511.11589.
- 1090
- Lobo, J.M., Jiménez-Valverde, A., Real, R., 2008. AUC: a misleading measure of the performance of predictive distribution models. Global Ecology and Biogeography 17, 145–151. doi:10.1111/j.1466-8238.2007.00358.x.
- 1095
- Maffei, C., Alfieri, S.M., Menenti, M., 2018. Relating spatiotemporal patterns of forest fires burned area and duration to diurnal land surface temperature anomalies. Remote Sensing 10, 1777. doi:10.3390/rs10111777.
- Maffei, C., Lindenbergh, R., Menenti, M., 2021. Combining multi-spectral and thermal remote sensing to predict forest fire characteristics. ISPRS Journal of Photogrammetry and Remote Sensing 181, 400–412. doi:10.1016/j.isprsjprs.2021.09.016.
- 1100
- Maffei, C., Menenti, M., 2019. Predicting forest fires burned area and rate of spread from pre-fire multispectral satellite measurements. ISPRS Journal

of Photogrammetry and Remote Sensing 158, 263–278. doi:10.1016/j.isprsjprs.2019.10.013.

- 1105 Meyer, H., Reudenbach, C., Hengl, T., Katurji, M., Nauss, T., 2018. Improving performance of spatio-temporal machine learning models using forward feature selection and target-oriented validation. *Environmental Modelling & Software* 101, 1–9. doi:10.1016/j.envsoft.2017.12.001.

- Moreno-Torres, J.G., Raeder, T., Alaiz-Rodríguez, R., Chawla, N.V., Herrera, F., 2012. A unifying view on dataset shift in classification. *Pattern Recognition* 45, 521–530. doi:10.1016/j.patcog.2011.06.019.

- Muñoz-Sabater, J., Dutra, E., Agustí-Panareda, A., Albergel, C., Arduini, G., Balsamo, G., Boussetta, S., Choulga, M., Harrigan, S., Hersbach, H., Martens, B., Miralles, D.G., Piles, M., Rodríguez-Fernández, N.J., Zsoter, E., Buontempo, C., Thépaut, J.N., 2021. ERA5-Land: a state-of-the-art global reanalysis dataset for land applications. *Earth System Science Data* 13, 4349–4383. doi:10.5194/essd-13-4349-2021.

- Pausas, J.G., Keeley, J.E., 2021. Wildfires and global change. *Frontiers in Ecology and the Environment* 19, 387–395. doi:10.1002/fee.2359.

- 1120 Persello, C., Bruzzone, L., 2012. Active learning for domain adaptation in the supervised classification of remote sensing images. *IEEE Transactions on Geoscience and Remote Sensing* 50, 4468–4483. doi:10.1109/TGRS.2012.2192740.

- Ploton, P., Mortier, F., Réjou-Méchain, M., Barbier, N., Picard, N., Rossi, V., Dormann, C., Cornu, G., Viennois, G., Bayol, N., Lyapustin, A.,

- Gourlet-Fleury, S., Péliissier, R., 2020. Spatial validation reveals poor predictive performance of large-scale ecological mapping models. *Nature Communications* 11, 4540. doi:10.1038/s41467-020-18321-y.
- Podschwit, H., Jolly, W., Alvarado, E., Verma, S., Ponce, B., Markos, A.,
1130 Aliaga-Nestares, V., Rodriguez-Zimmermann, D., 2022. Reliability of cross-regional applications of global fire danger models: A Peruvian case study. *Fire Ecology* 18, 25. doi:10.1186/s42408-022-00150-7.
- Roberts, D.R., Bahn, V., Ciuti, S., Boyce, M.S., Elith, J., Guillera-Arroita, G., Hauenstein, S., Lahoz-Monfort, J.J., Schröder, B., Thuiller, W.,
1135 Warton, D.I., Wintle, B.A., Hartig, F., Dormann, C.F., 2017. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography* 40, 913–929. doi:10.1111/ecog.02881.
- Rousseau, J.S., Betts, M.G., 2022. Factors influencing transferability in species distribution models. *Ecography* 2022, e06060. doi:10.1111/ecog.
1140 06060.
- Sandholt, I., Rasmussen, K., Andersen, J., 2002. A simple interpretation of the surface temperature/vegetation index space for assessment of surface moisture status. *Remote Sensing of Environment* 79, 213–224. doi:10.1016/S0034-4257(01)00274-7.
- 1145 Sun, B., Feng, J., Saenko, K., 2016. Return of frustratingly easy domain adaptation, in: *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence (AAAI-16)*, AAAI Press, Phoenix, Arizona, USA. pp. 2058–2065. doi:10.1609/aaai.v30i1.10306.

- Tuia, D., Persello, C., Bruzzone, L., 2016. Domain adaptation for the clas-
1150 sification of remote sensing data: An overview of recent advances. *IEEE
Geoscience and Remote Sensing Magazine* 4, 41–57. doi:10.1109/MGRS.
2016.2548504.
- Valavi, R., Elith, J., Lahoz-Monfort, J.J., Guillerá-Arroita, G., 2019.
blockCV: An R package for generating spatially or environmentally sepa-
1155 rated folds for k-fold cross-validation of species distribution models. *Meth-
ods in Ecology and Evolution* 10, 225–232. doi:10.1111/2041-210X.
13107.
- VanDerWal, J., Shoo, L.P., Graham, C., Williams, S.E., 2009. Selecting
pseudo-absence data for presence-only distribution modeling: How far
1160 should you stray from what you know? *Ecological Modelling* 220, 589–594.
doi:10.1016/j.ecolmodel.2008.11.010.
- Vesk, P.A., Morris, W.K., Neal, W.C., Mokany, K., Pollock, L.J., 2021.
Transferability of trait-based species distribution models. *Ecography* 44,
134–147. doi:10.1111/ecog.05179.
- 1165 Vibhandik, P., Sawant, S., Joshi, A., Bidwe, R., 2026. A systematic litera-
ture review on forest fire susceptibility mapping using geo-spatial technol-
ogy and future research directions. *Discover Artificial Intelligence* 6, 921.
doi:10.1007/s44163-026-00921-0.
- Xu, Y., Dai, Y., Chang, L., Wang, Q., Dong, Y., 2026. Does your wild-
1170 fire prediction model actually work, or just score well? *arXiv preprint
arXiv:2605.18911*. doi:10.48550/arXiv.2605.18911, arXiv:2605.18911.

- Yebra, M., Dennison, P.E., Chuvieco, E., Riaño, D., Zylstra, P., Hunt, E.R.,
Danson, F.M., Qi, Y., Jurdao, S., 2013. A global review of remote sens-
ing of live fuel moisture content for fire danger assessment: Moving to-
wards operational products. *Remote Sensing of Environment* 136, 455–468.
1175 doi:10.1016/j.rse.2013.05.029.
- Zanaga, D., Van De Kerchove, R., Daems, D., De Keersmaecker, W., Brock-
mann, C., Kirches, G., Wevers, J., Cartus, O., Santoro, M., Fritz, S., Lesiv,
M., Herold, M., Tsendbazar, N.E., Xu, P., Ramoino, F., Arino, O., 2022.
1180 ESA WorldCover 10 m 2021 v200. Zenodo. doi:10.5281/zenodo.7254221.